

# Building Deep Learning Models Using PyTorch

---

## INTRODUCTION TO PYTORCH



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**PyTorch is a deep learning framework**

**More tightly integrated with Python than TensorFlow**


**Can use Python libraries, debugger**

**Unlike TensorFlow, supports dynamic computation graphs**

**Like other deep learning frameworks uses a forward and backward pass for training**

# Prerequisites and Course Outline

---



# Prereqs - Basic ML

## **Building Machine Learning Models in Python with scikit-learn**

- Using the scikit-learn library for ML

## **How to Think About Machine Learning Algorithms**

- Introduction to machine learning



# Software and Skills

**Be very comfortable programming in Python 3 using Jupyter notebooks**

**Be comfortable with ML concepts**

**Understand the basics of neural networks for deep learning**



# Course Outline

## **Introducing PyTorch**

- Machine learning using NNs
- Deep learning framework, alternative to TensorFlow

## **Building Simple Neural Networks**

- Autograd in PyTorch for back-propagation
- Defining neural networks and connecting them

## **Building Image Classification Models**

- Convolutional Neural Networks in PyTorch
- Using pre-trained models for image classification

## **Building Text Classification Models**

- Recurrent Neural Networks in PyTorch

# Introducing Neural Networks

---

# Reviews: Positive or Negative?





# ML-based Classifier

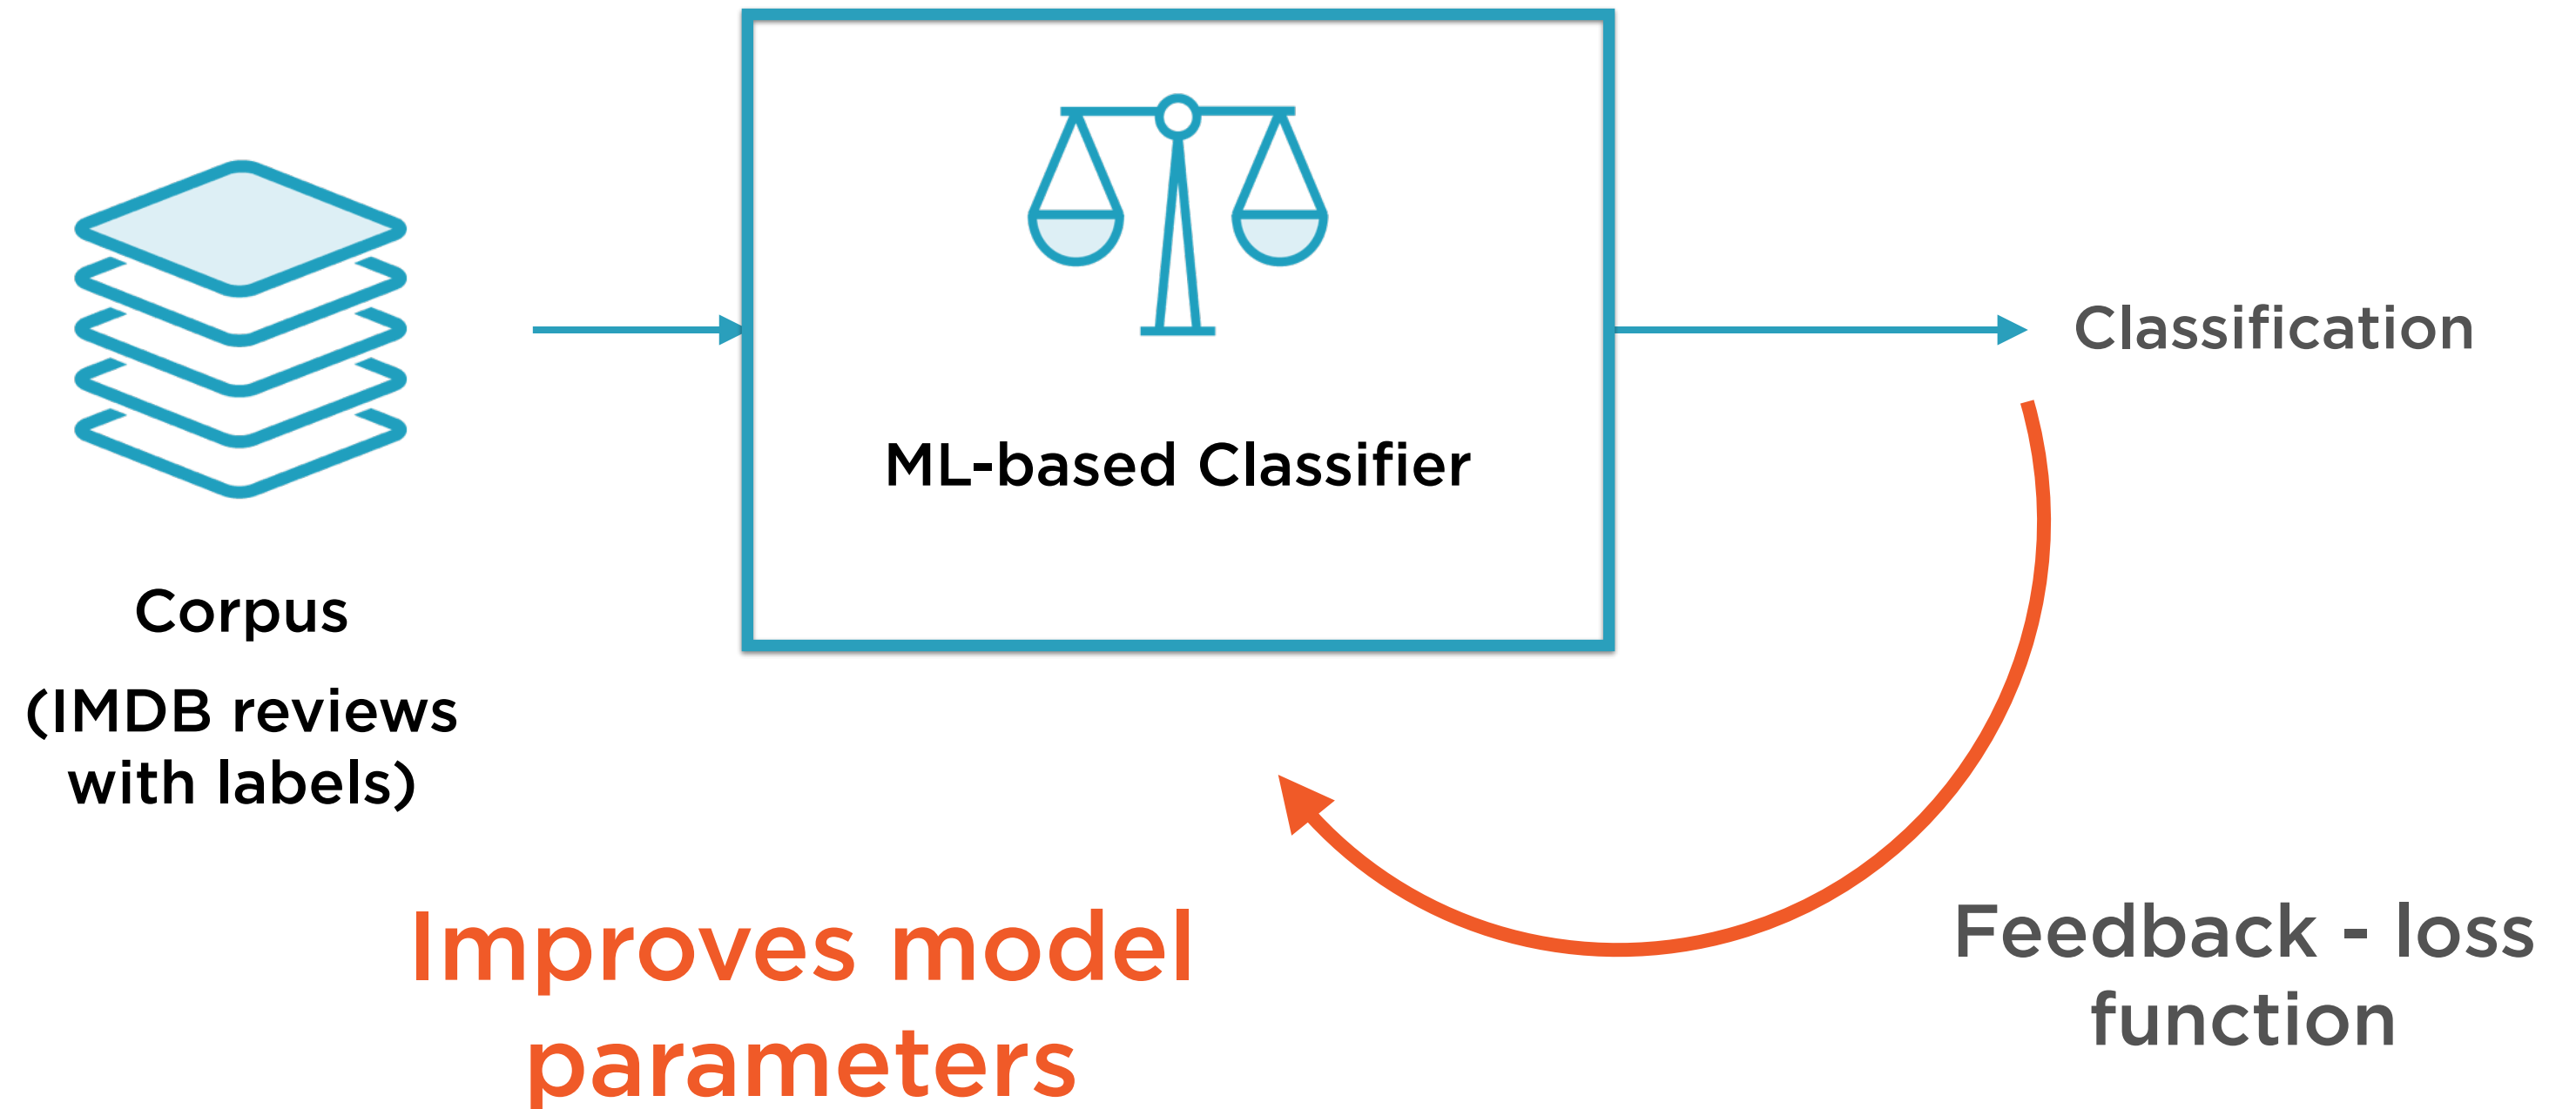
## Training

Feed in a large corpus of data  
classified correctly

## Prediction

Use it to classify new instances  
which it has not seen before

# Training the ML-based Classifier



**“Traditional”** ML-based systems  
rely on experts to decide what  
features to pay attention to

**“Representation”** ML-based systems figure out by themselves what features to pay attention to

Neural networks are examples of such systems

# What is a Neural Network?

## Deep Learning

Algorithms that learn  
what features matter


## Neural Networks

The most common class  
of deep learning  
algorithms


## Neurons

Simple building blocks  
that actually “learn”


# Neural Networks



# Neural Networks

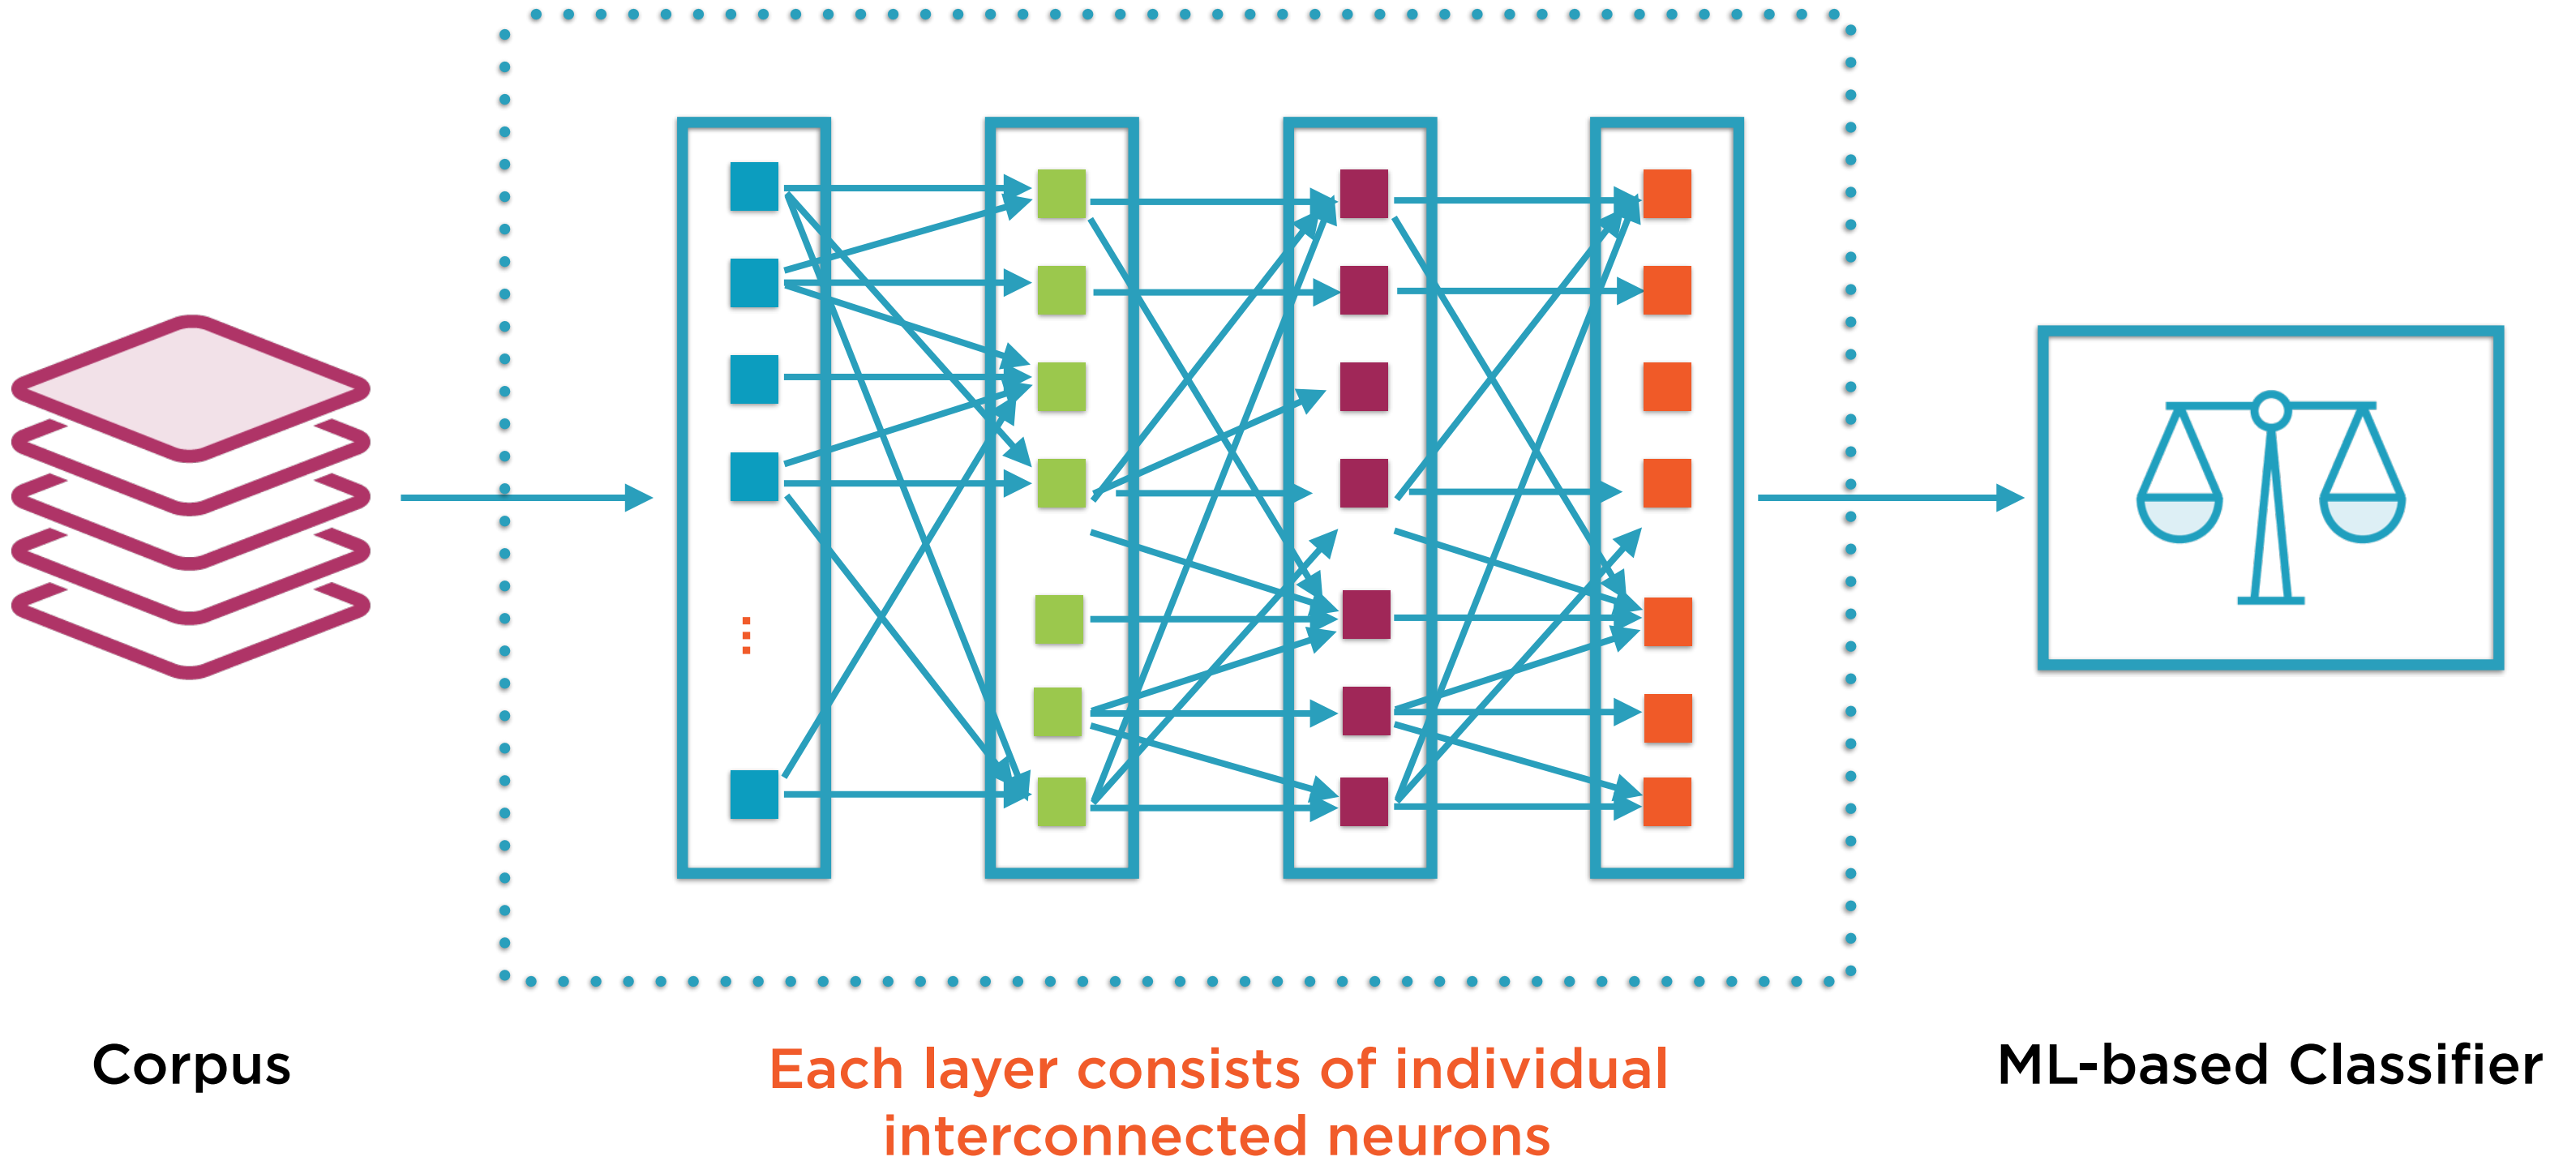


# Neural Networks






# Neural Networks




# Operation of a Single Neuron




**For an active neuron a change in inputs should trigger a corresponding change in the outputs**

# Operation of a Single Neuron




**The outputs of neurons feed into the neurons from the next layer**

# Operation of a Single Neuron



**Each connection is associated with a weight**

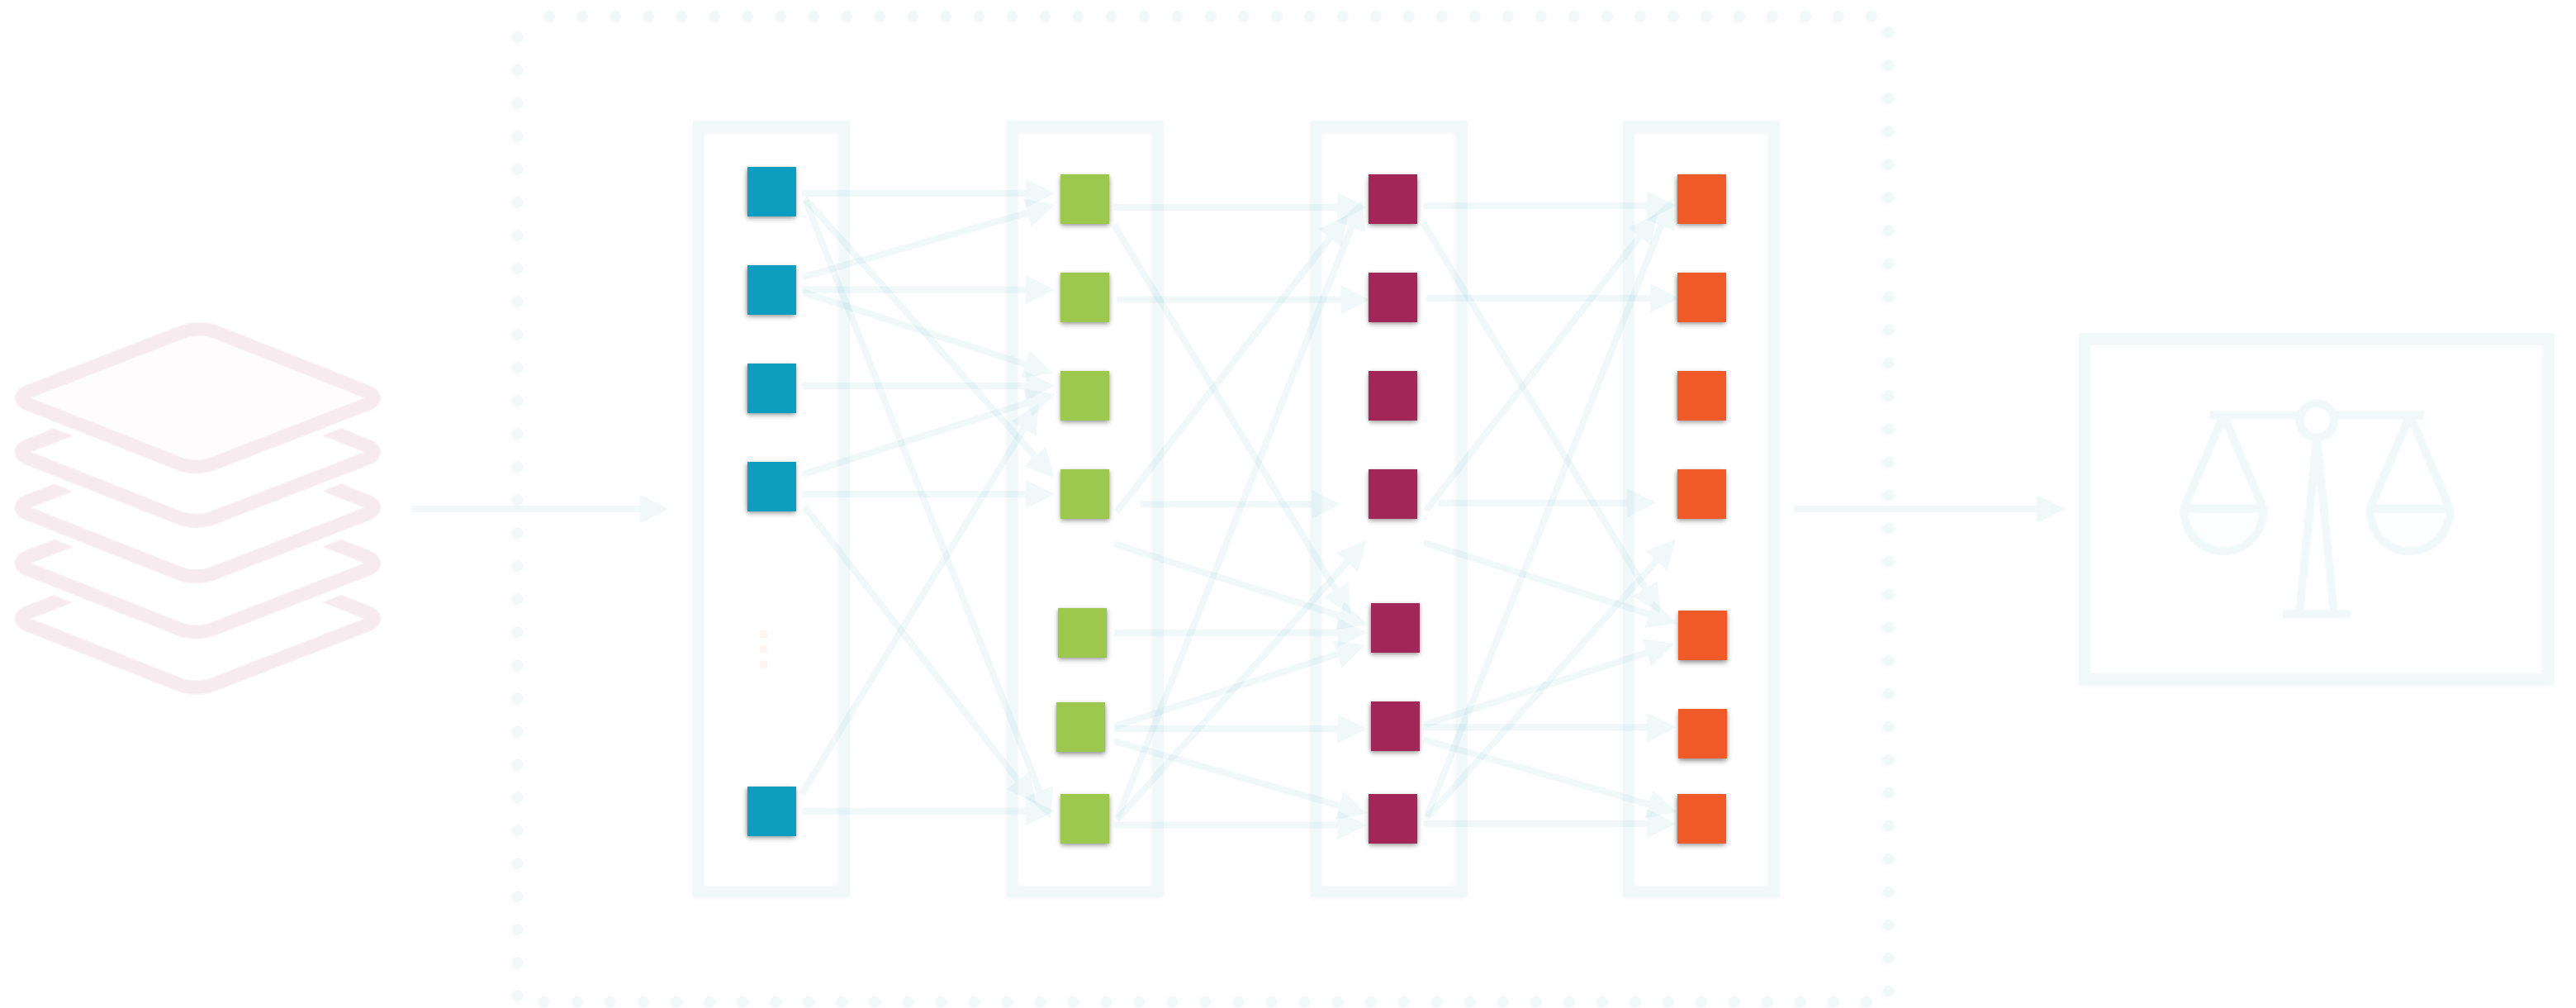
# Operation of a Single Neuron



If the second neuron is sensitive to the output of the first neuron, the **connection between them gets stronger**

**$W$  increases**

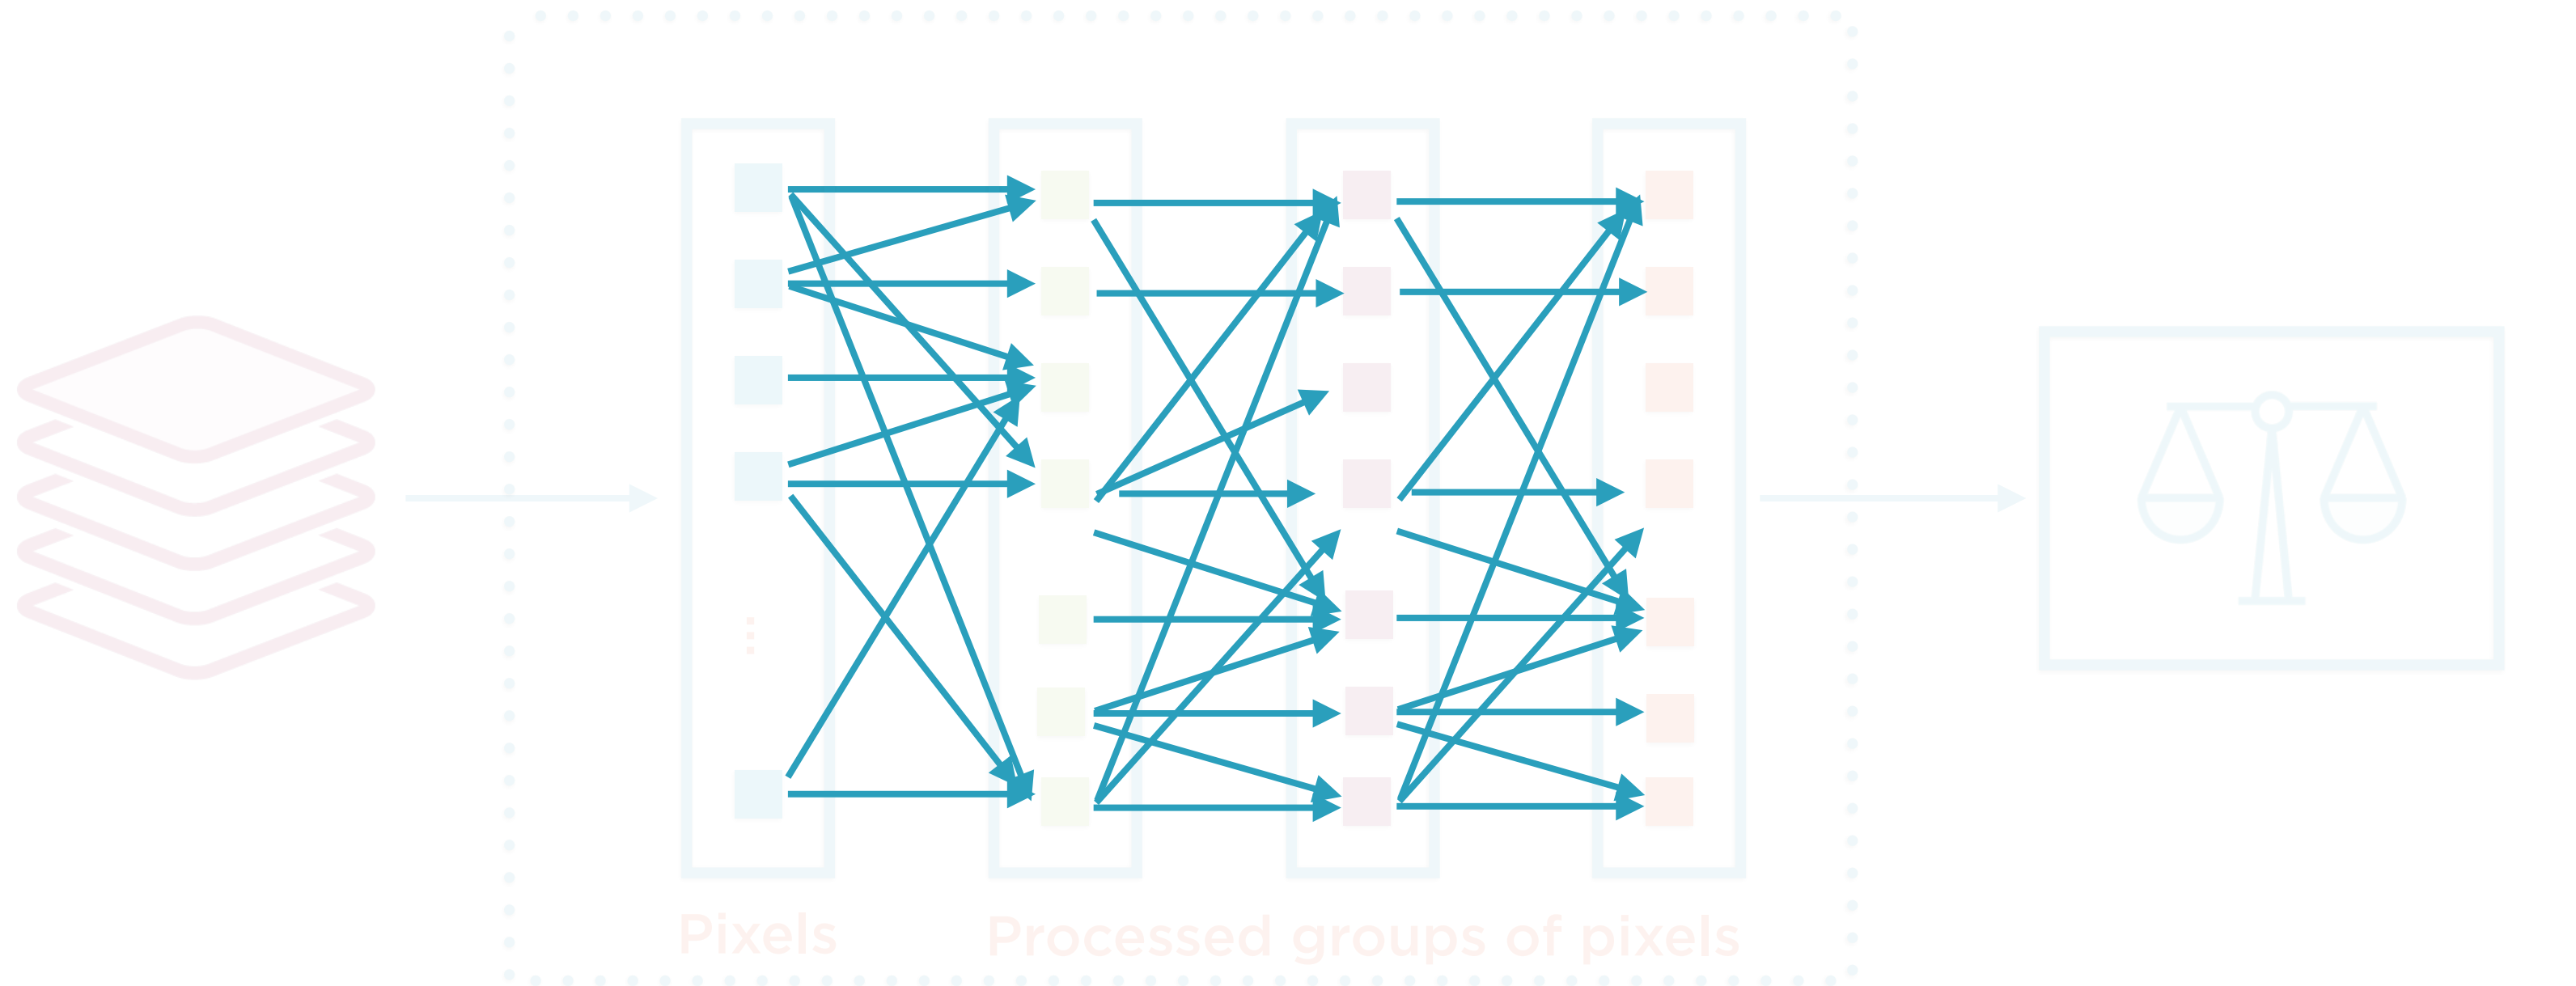
# The Computational Graph



Corpus

**The nodes in the computation graph are** ML-based Classifier  
**neurons (simple building blocks)**

# The Computational Graph




Corpus


**The edges in the computation graph  
are data called tensors**

ML-based Classifier

# A Neural Network



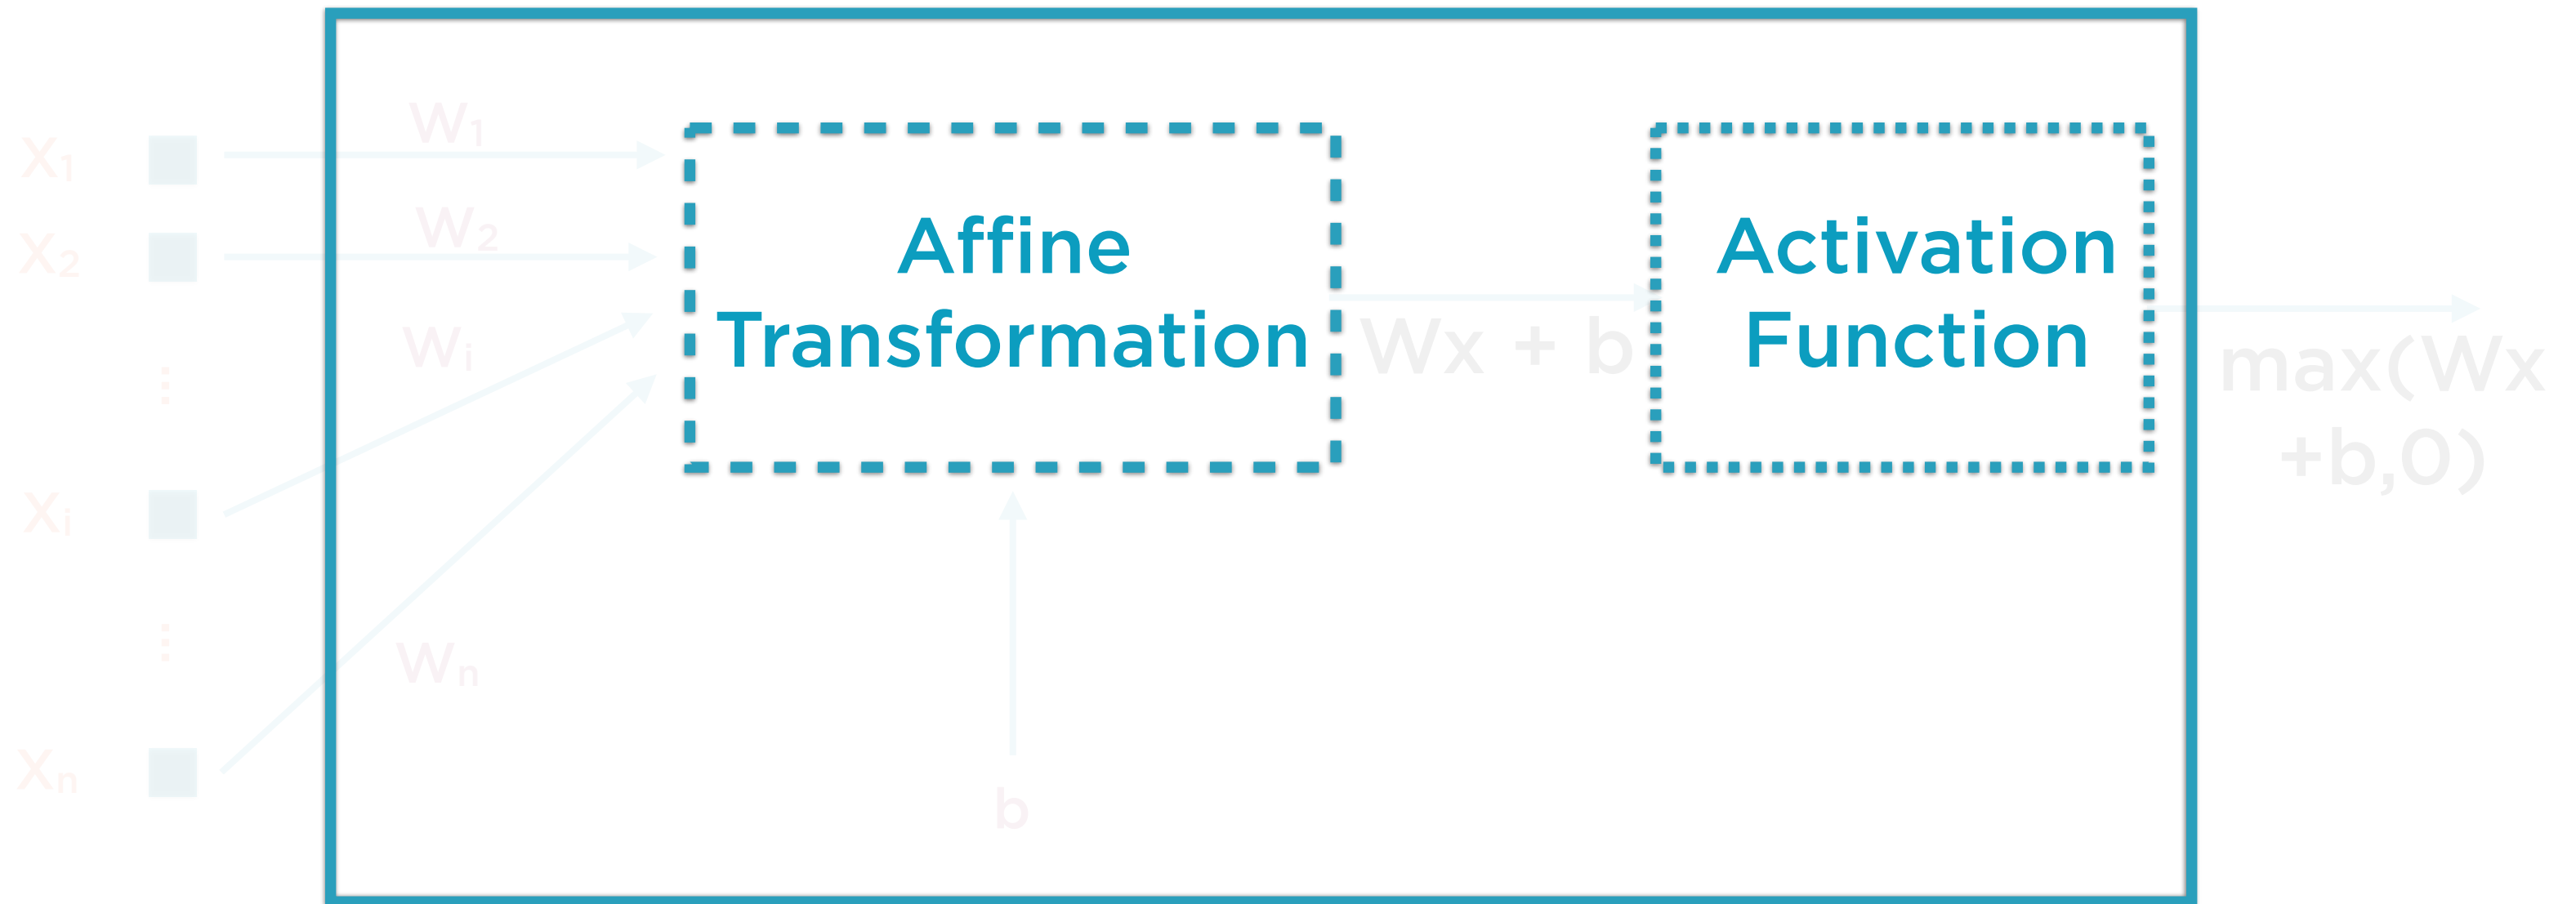
Corpus



Once a neural network is **trained**, all edges have weights which help it make predictions

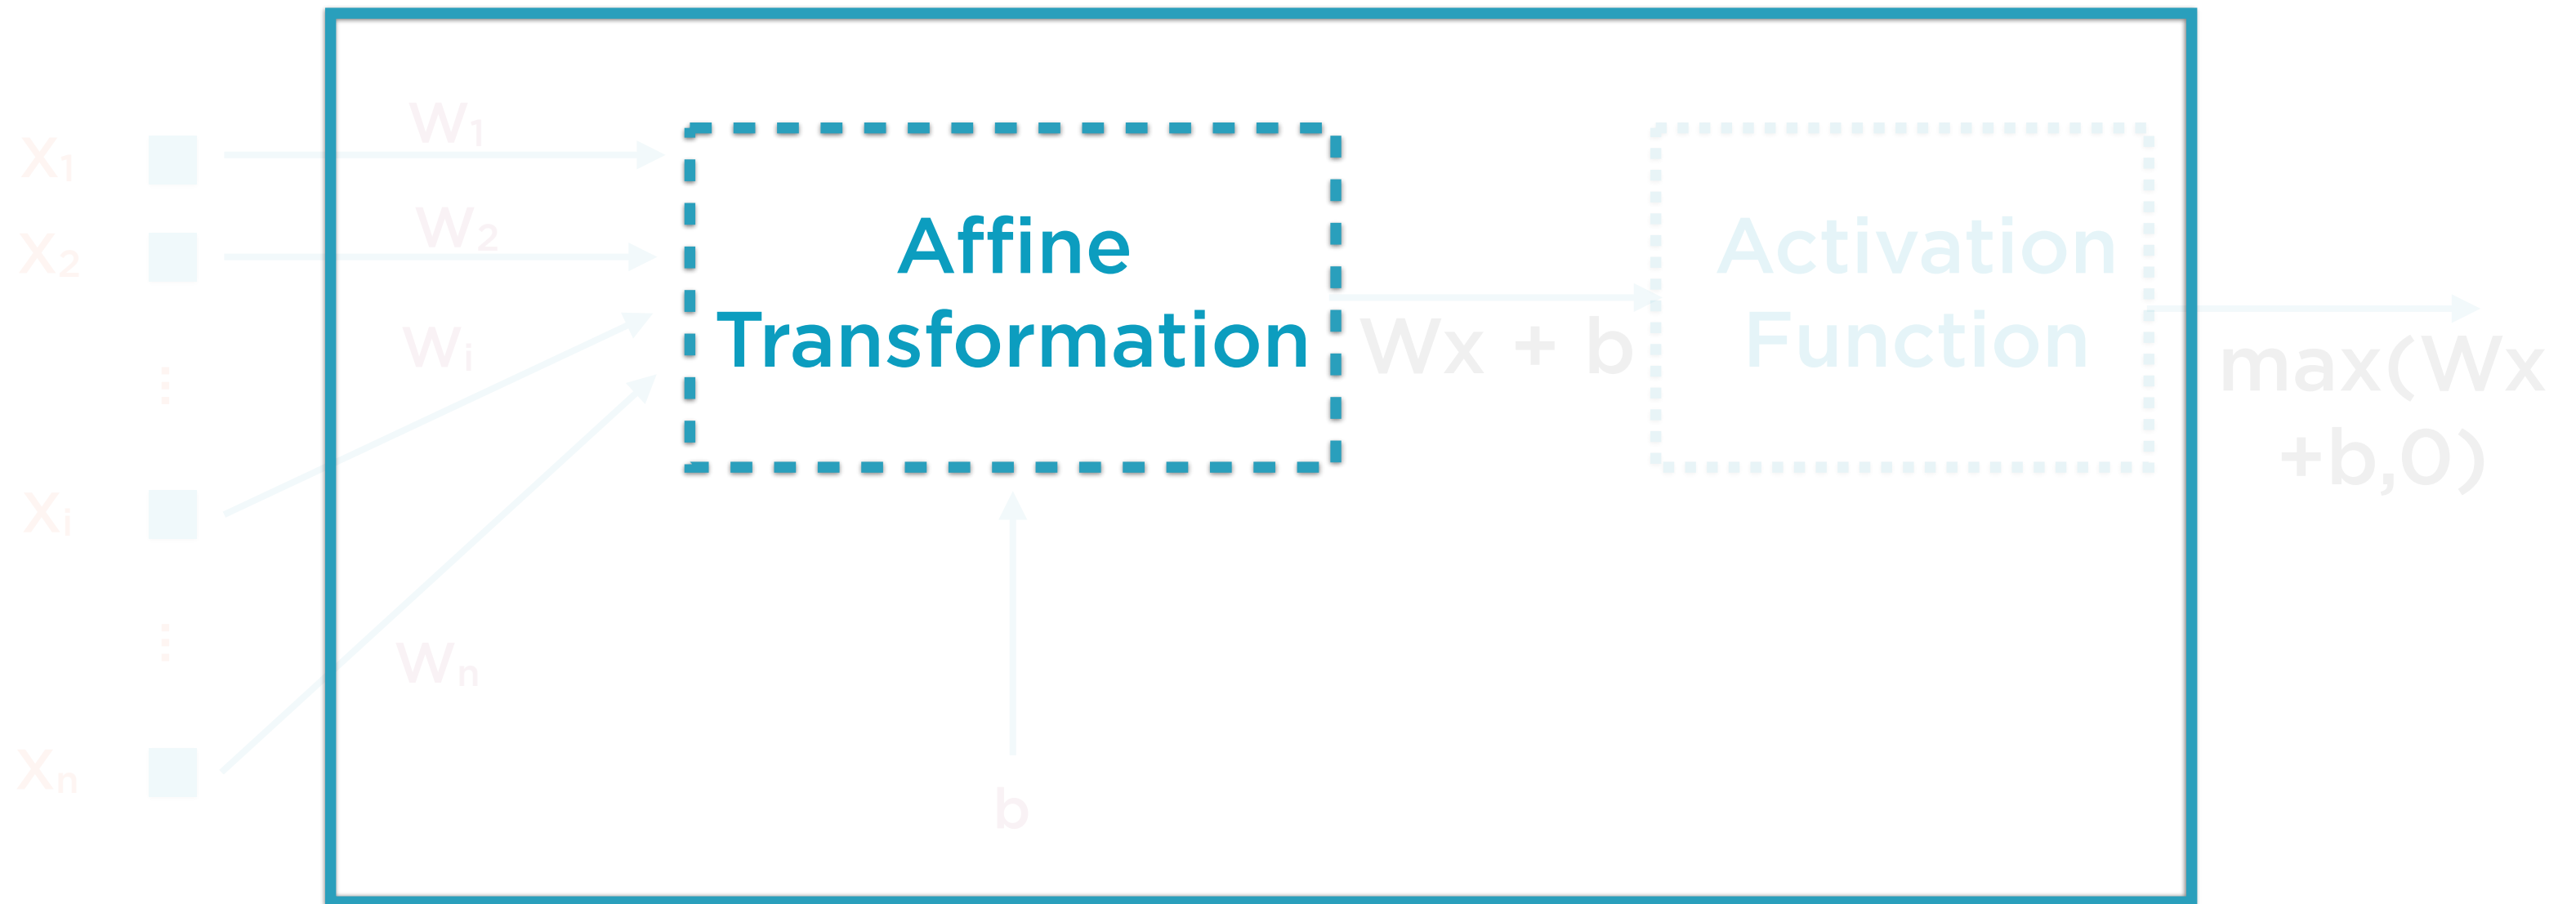


# Operation of a Single Neuron



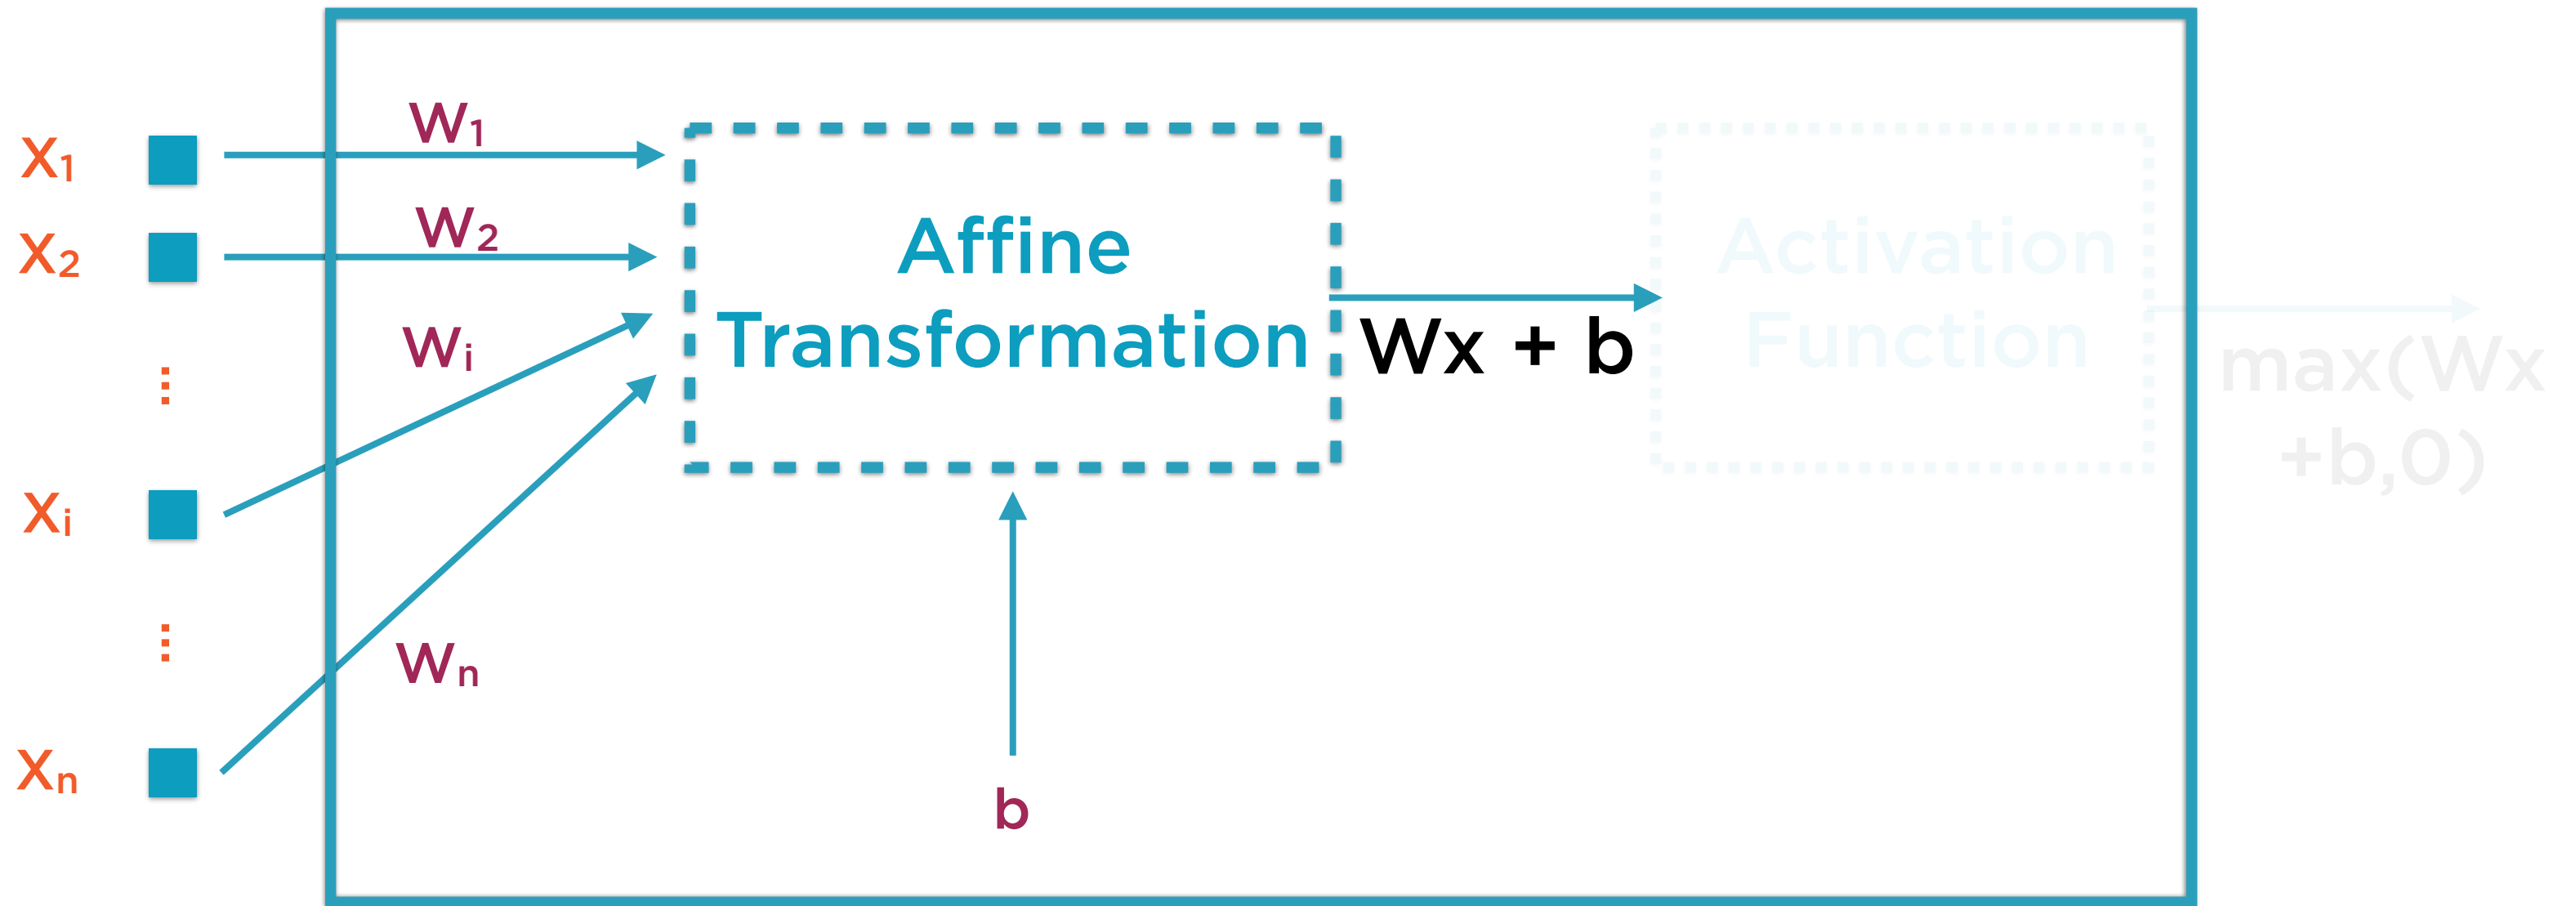
**Each neuron only applies two simple functions to its inputs**

# Operation of a Single Neuron



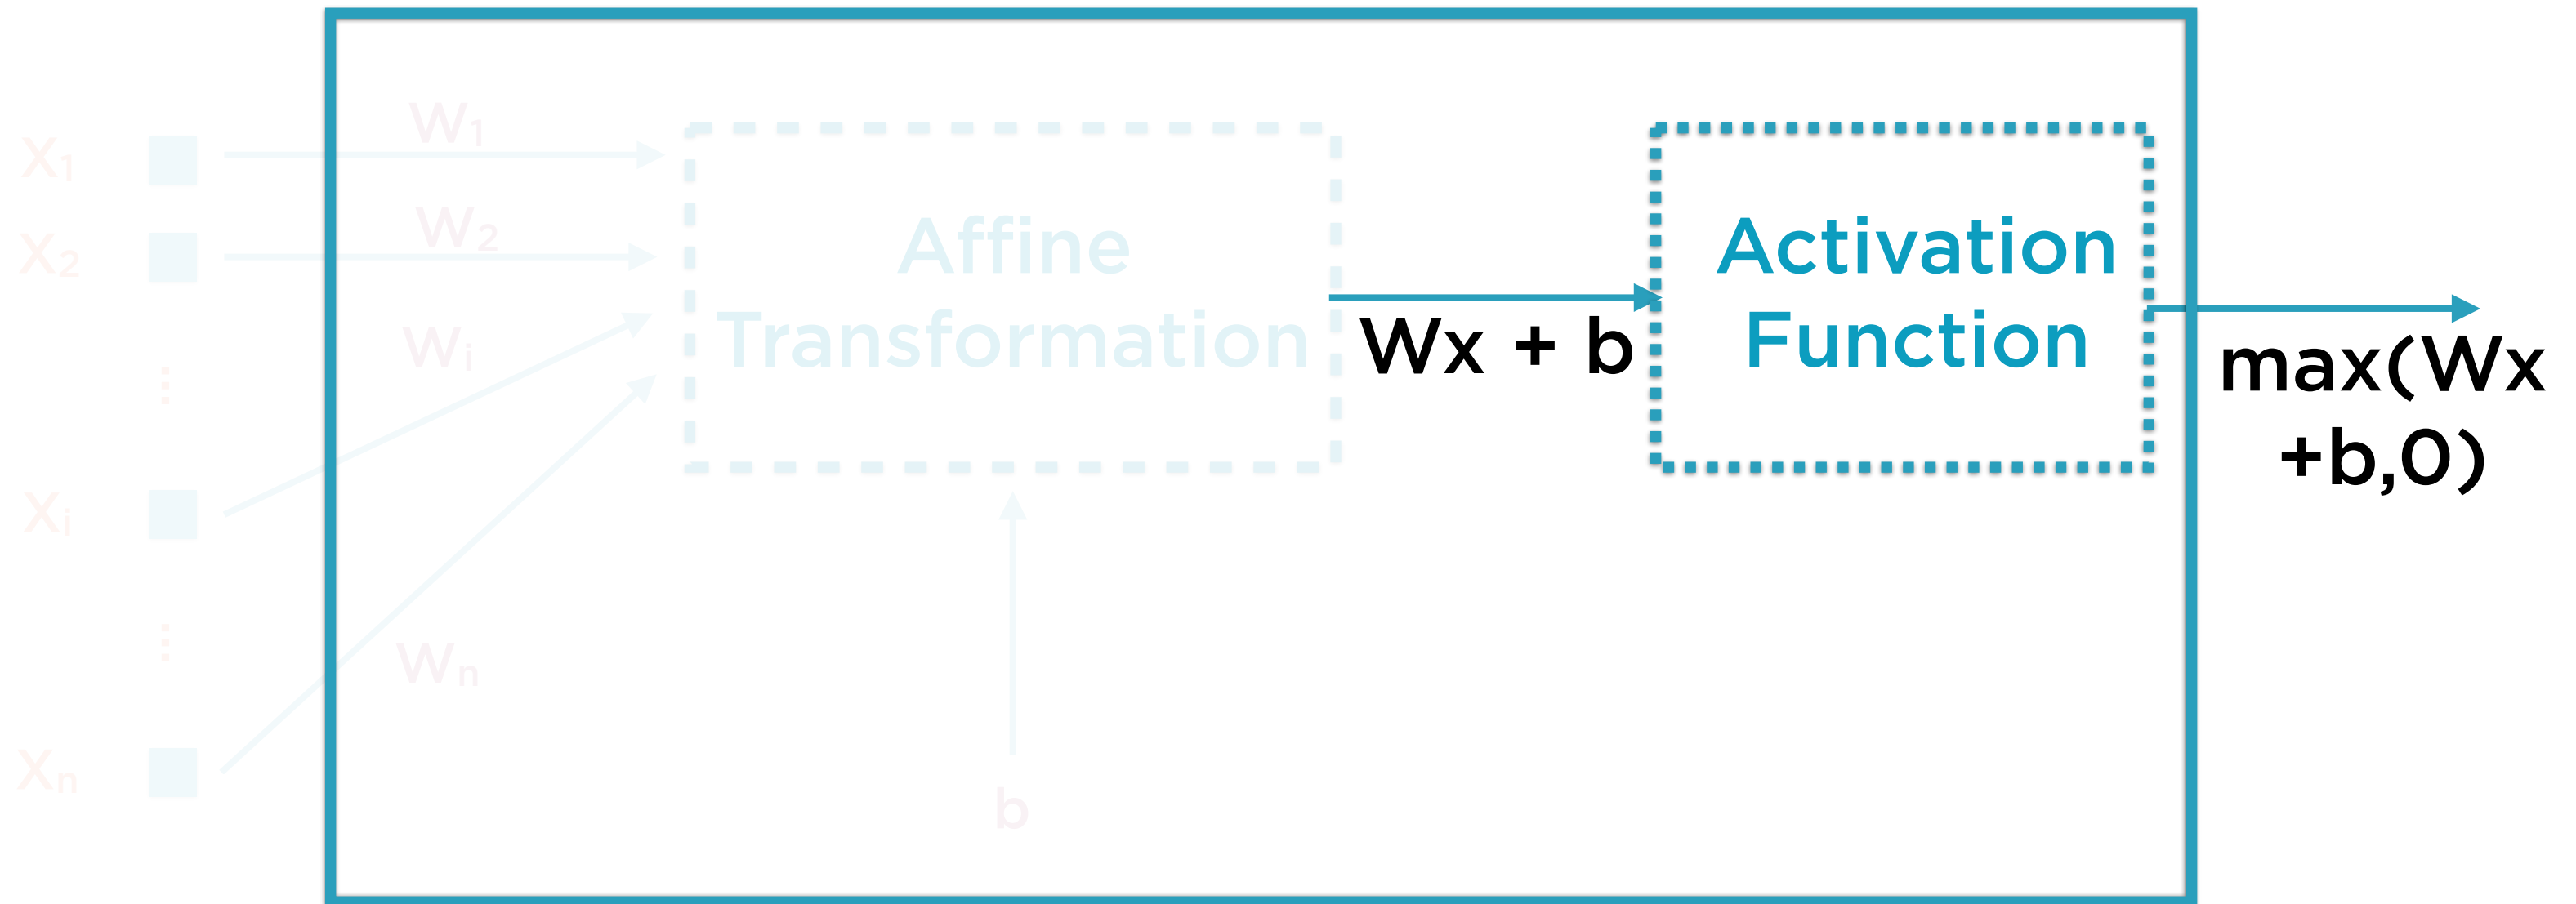
The affine transformation alone can **only** learn **linear** relationships between the inputs and the output

# Operation of a Single Neuron



The affine transformation is just a weighted sum with a bias added:  $W_1x_1 + W_2x_2 + \dots + W_nx_n + b$


# Operation of a Single Neuron




The **combination** of the affine transformation and the activation function can **learn any arbitrary relationship**

# Common Activation Functions


ReLU




logit



tanh

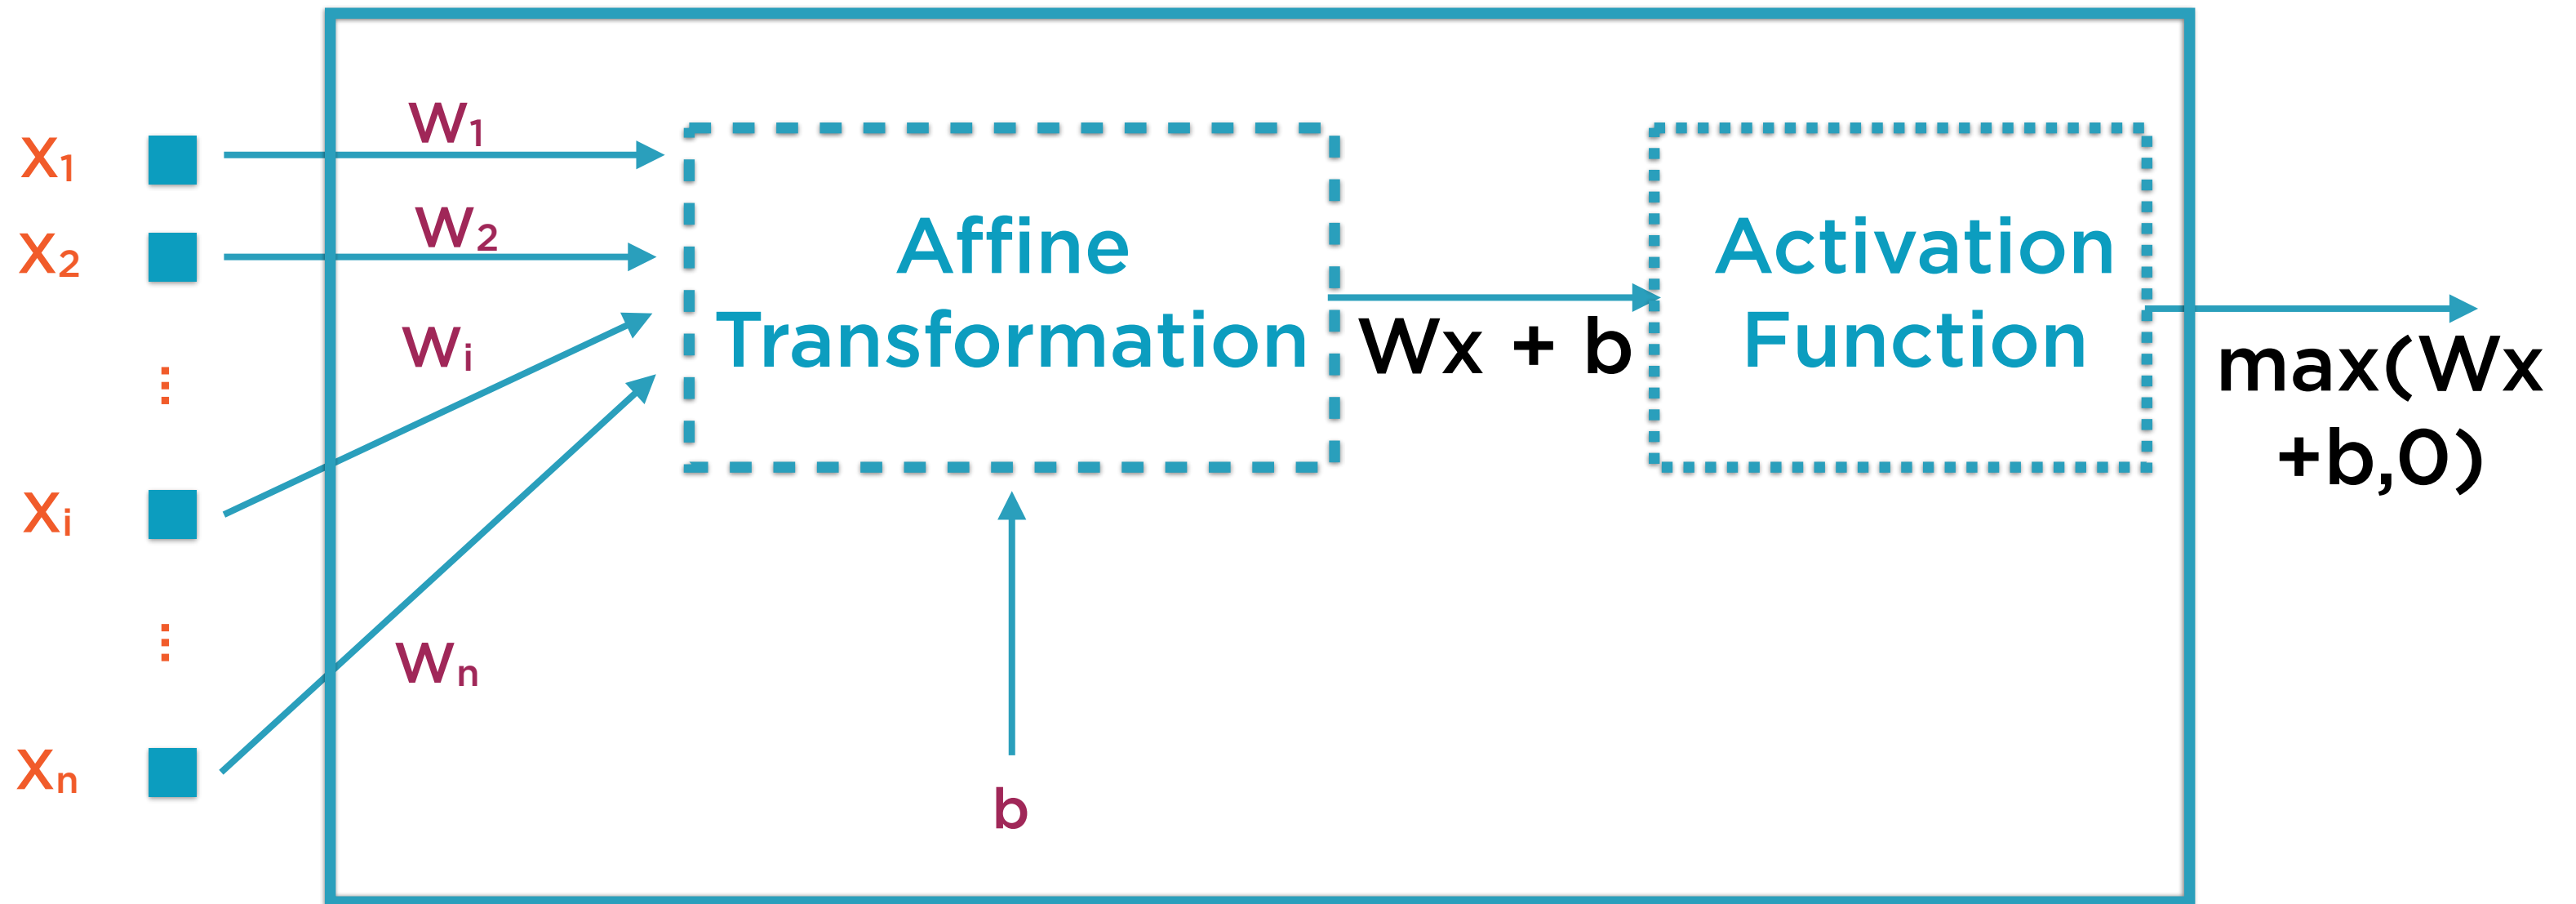


step



Notice how activations functions have a gradient, this gradient allows them to be sensitive to input changes

# Neuron as a Learning Unit



Many of these simple neurons arranged in layers can do magical stuff

The **weights** and **biases** of individual neurons are determined during the **training** process

# Introducing PyTorch

---



# PyTorch

A deep learning framework for fast, flexible experimentation.

*<https://pytorch.org/>*

# PyTorch Use-cases



## Two primary uses


### As alternative to TensorFlow

- Deep learning framework

### As alternative to NumPy

- Tensor computations but with “strong GPU acceleration”

# Deep Learning Framework




**“Tape-based Autograd”**

**Neural network can be redefined  
dynamically**

**Different from TensorFlow, CNTK**


# GPU-ready Tensor Library



**Tensors for either CPU or GPU**

**Powerful, fast NumPy-like functionality**

- slicing
- indexing
- reductions
- linear algebra



# Tight Python Integration


**Deeply tied to Python**

**Approach similar to numpy/scikit-learn**

**Create neural networks in Python**

**Use existing Python libraries...**

**...and debuggers**



# Imperative Execution

**Write code, run immediately**

**No separate build and run phases**


**Debugging easier**

# Ease of Extensibility

**Easily write new neural network layers**

**Different alternatives**

- Torch API
- Scipy
- C/C++ extension API





# PyTorch Background

## **Relatively new**

- Initial release October 2016
- Stable release April 2018

## **Based on Torch**

- Open-source ML library (since 2002!)

## **Facebook connection**

- Developed by Facebook AI researchers



# TensorFlow vs. PyTorch

## TensorFlow

Computation graph is static...

...must be defined before being run

`tf.Session` for separation from Python

Debugging via `tfdg`

Visualization using TensorBoard

Deployment using TF Serving

`tf.device` and `tf.DeviceSpec` to use GPUs  
(relatively hard)

## PyTorch

Computation graph is dynamic...

...can be defined and run as you go

Tightly integrated with Python

Debugging with PyCharm, `pdb`

Visualization using `matplotlib`, `seaborn`

Need to set up REST API e.g. Flask

`torch.nn.DataParallel` to use GPUs  
(relatively easy)

Demo

**Installing PyTorch on your local machine**

# Tensors in PyTorch

---

Tensors in PyTorch: Conceptually  
identical to tensors in TensorFlow

# Tensor

The central unit of data in TensorFlow. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

<https://www.tensorflow.org/>

# Tensor

The central unit of data in TensorFlow. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

<https://www.tensorflow.org/>

# Tensor

The central unit of data in TensorFlow. A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

<https://www.tensorflow.org/>

# Data Is Represented as Tensors



**Scalars** are **0-D** tensors

**3, 6.7, “a”**



# Data Is Represented as Tensors



**Vectors** are **1-D** tensors


**[1, 3, 5, 7, 9]**

# Data Is Represented as Tensors


**Matrices** are **2-D** tensors

**[[1, 3, 5],  
[7, 9, 11]]**

# Data Is Represented as Tensors



**N-Dimensional matrices** are **N-D** tensors

**[[[1, 2], [3, 4], [5, 6],  
[7, 8], [9, 10], [11, 12]]]**

# Characterization of Tensors




## Rank

The number of dimensions in a tensor



## Shape

The number of elements in each dimension



## Data Type

The data type of each element in the tensor




# Rank

Tensor	Rank
4	0
[1, 2, 3]	1
[[1, 2], [3, 4]]	2
[[[1], [2]], [[3], [4]]]	3



# Shape

Tensor	Shape
4	[]
[1, 2, 3]	[3]
[[1, 2, 3], [4, 5, 6]]	[3, 2]
[[[1], [2]], [[3], [4]]]	[2, 2, 1]



Data Type

**int**  
**float**  
**string**  
**boolean**

Rank, shape and data types are  
3 important characteristics  
which define a Tensor



Demo

**Tensors in PyTorch**

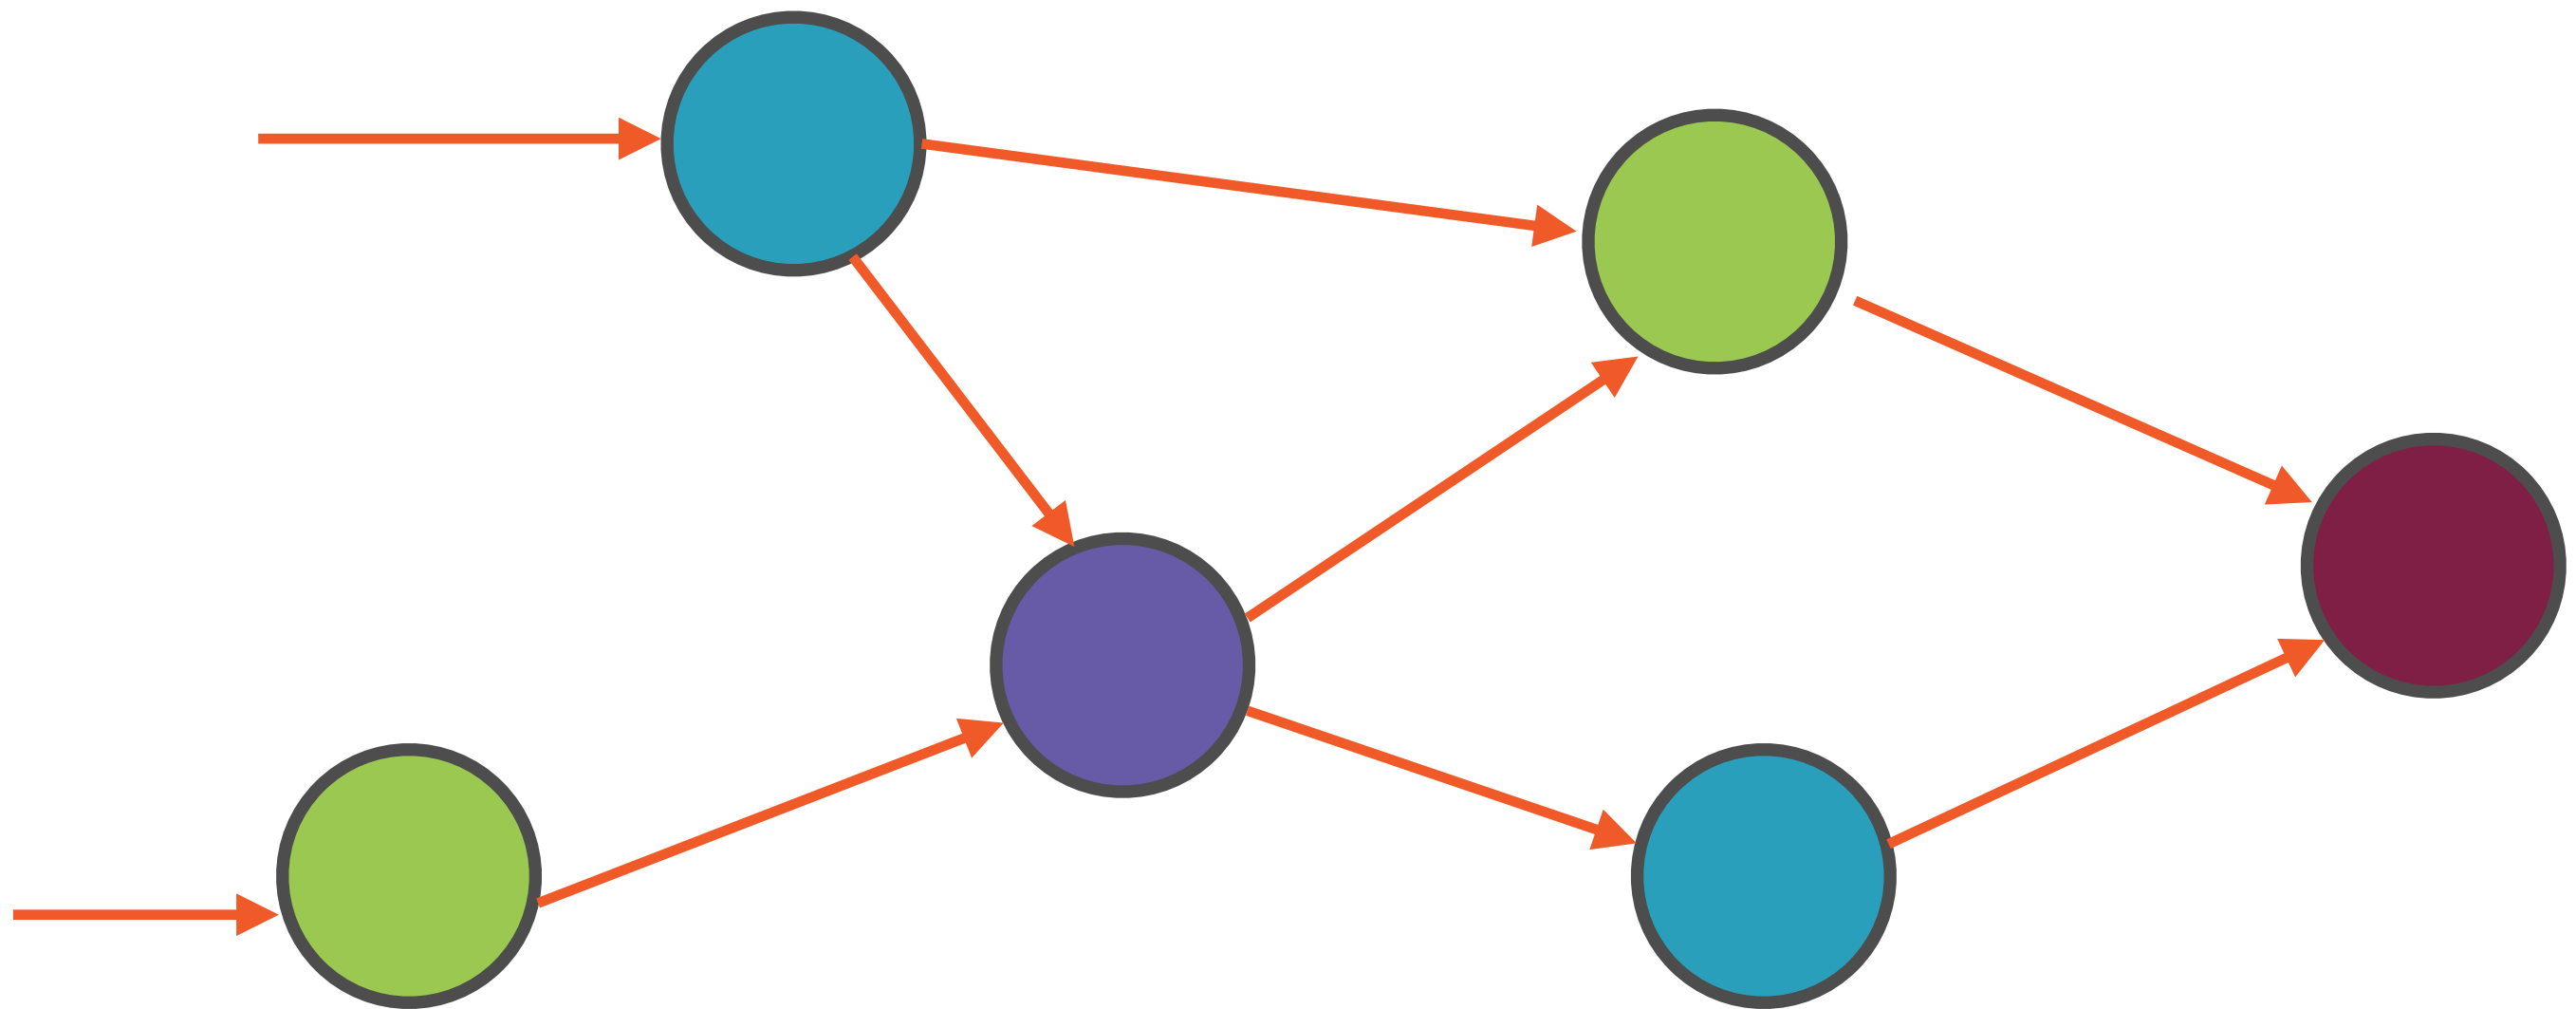
**Operations with Tensors**

# Computation Graphs in PyTorch

---

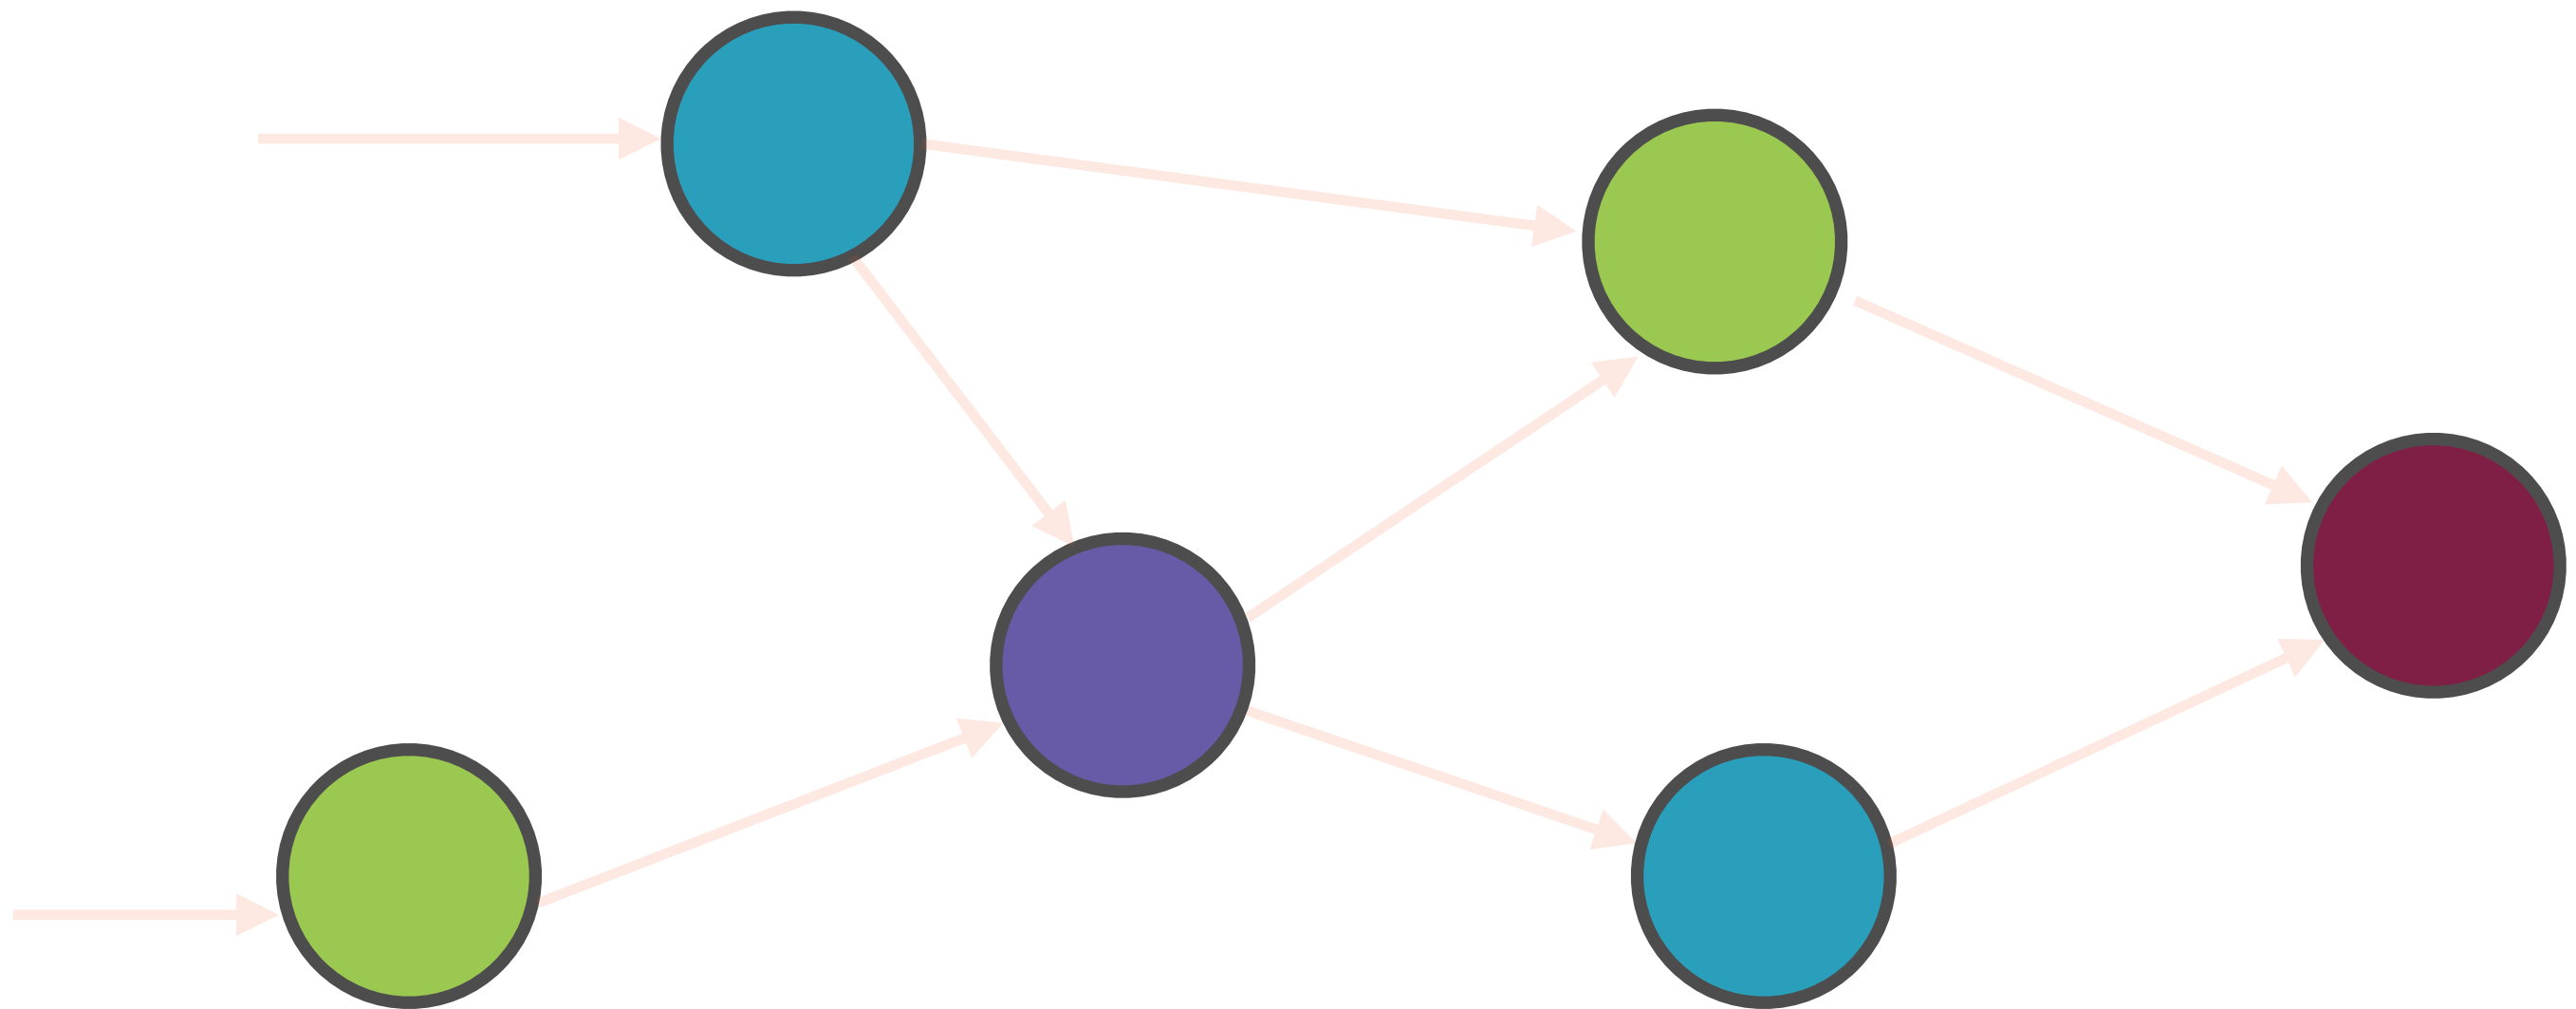
Similar to TF graphs but with one important difference: PyTorch computation graphs are **dynamic**

Everything Is a Graph



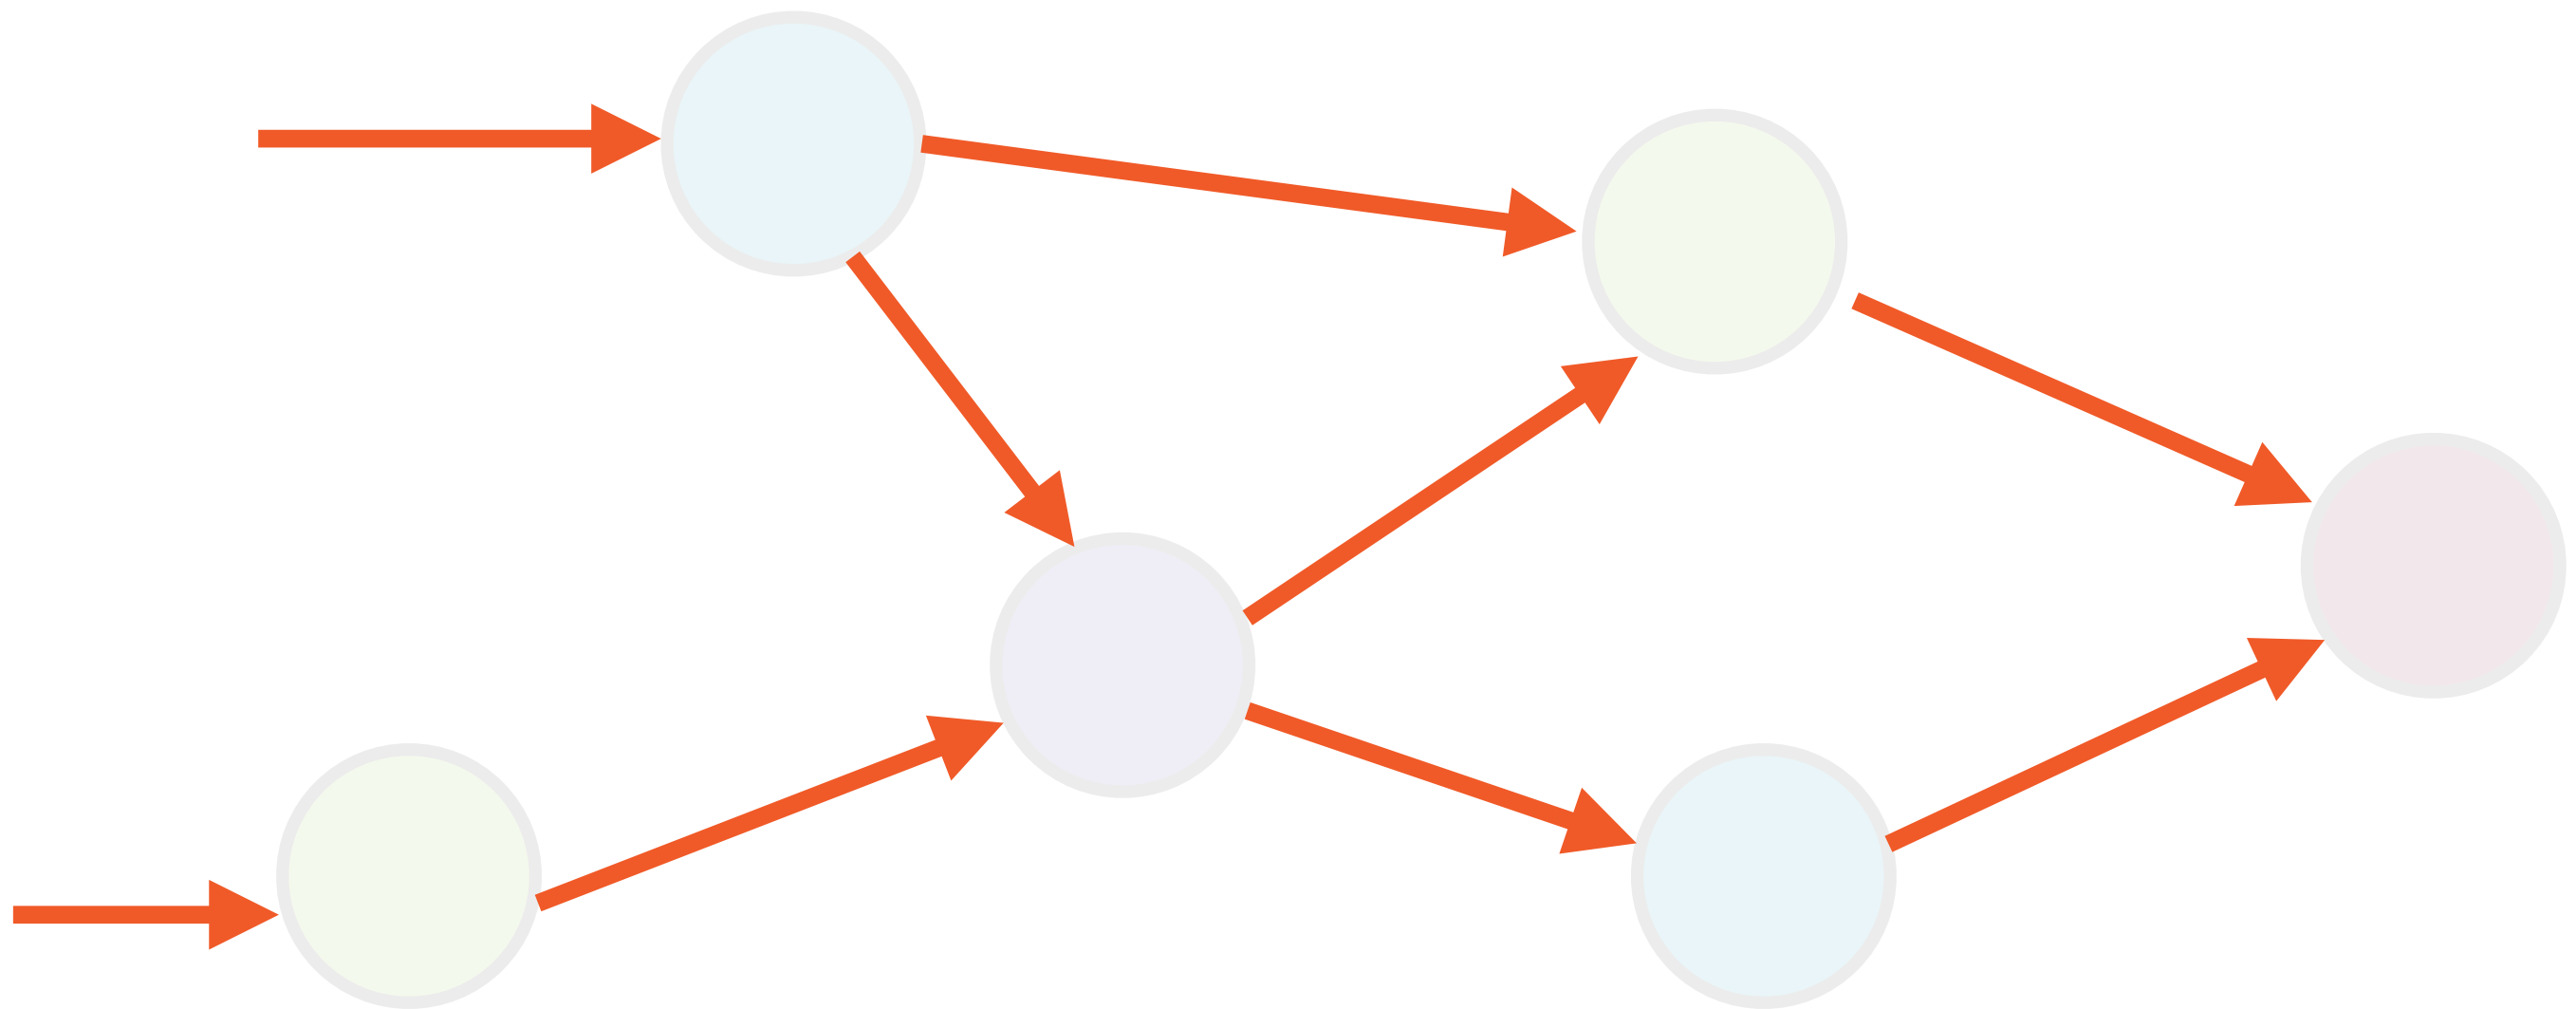
**A network**

# Everything Is a Graph



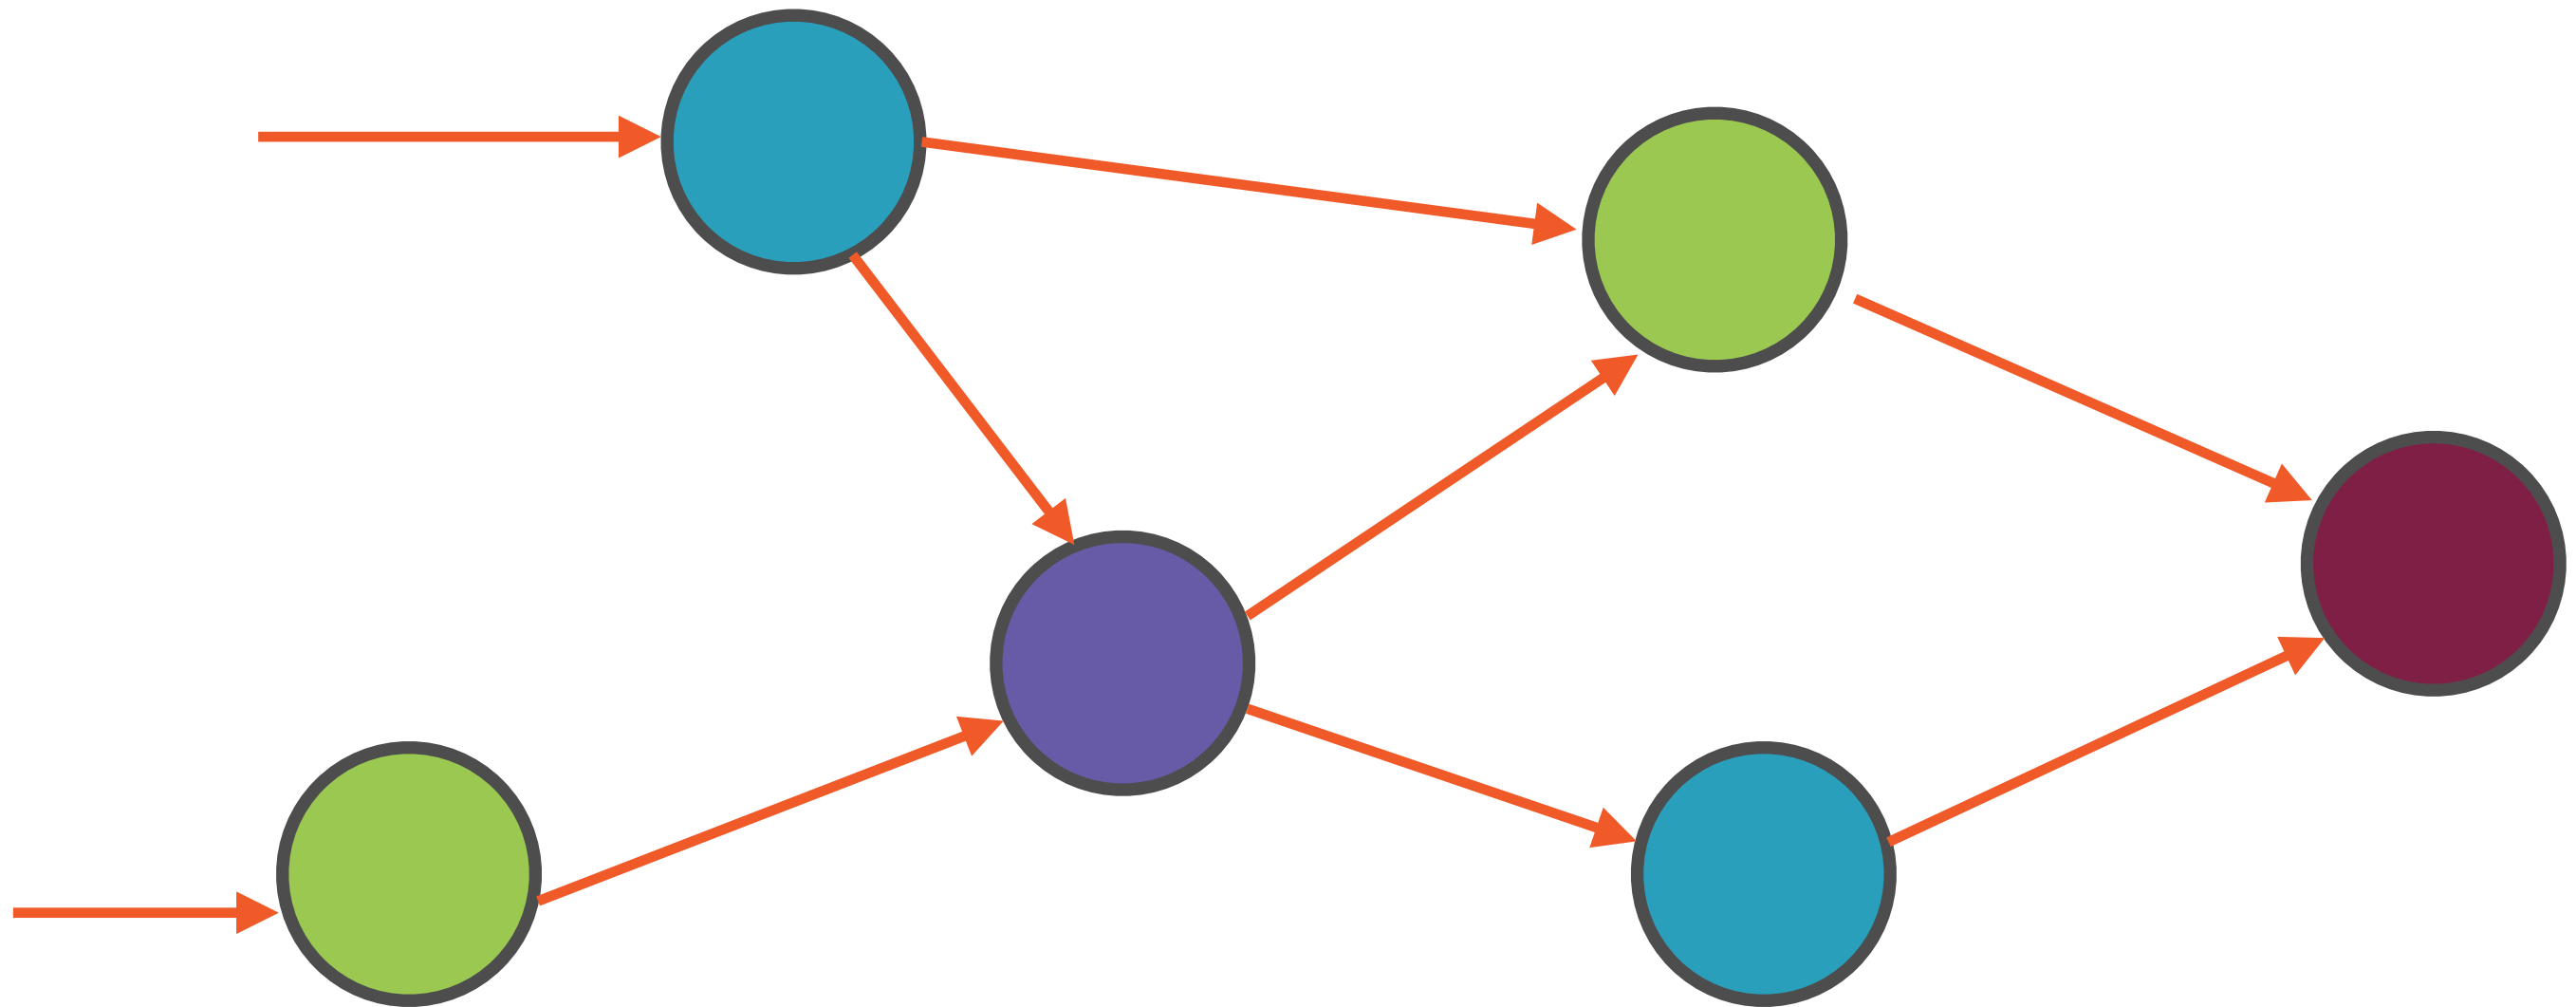
## Computations

# Everything Is a Graph



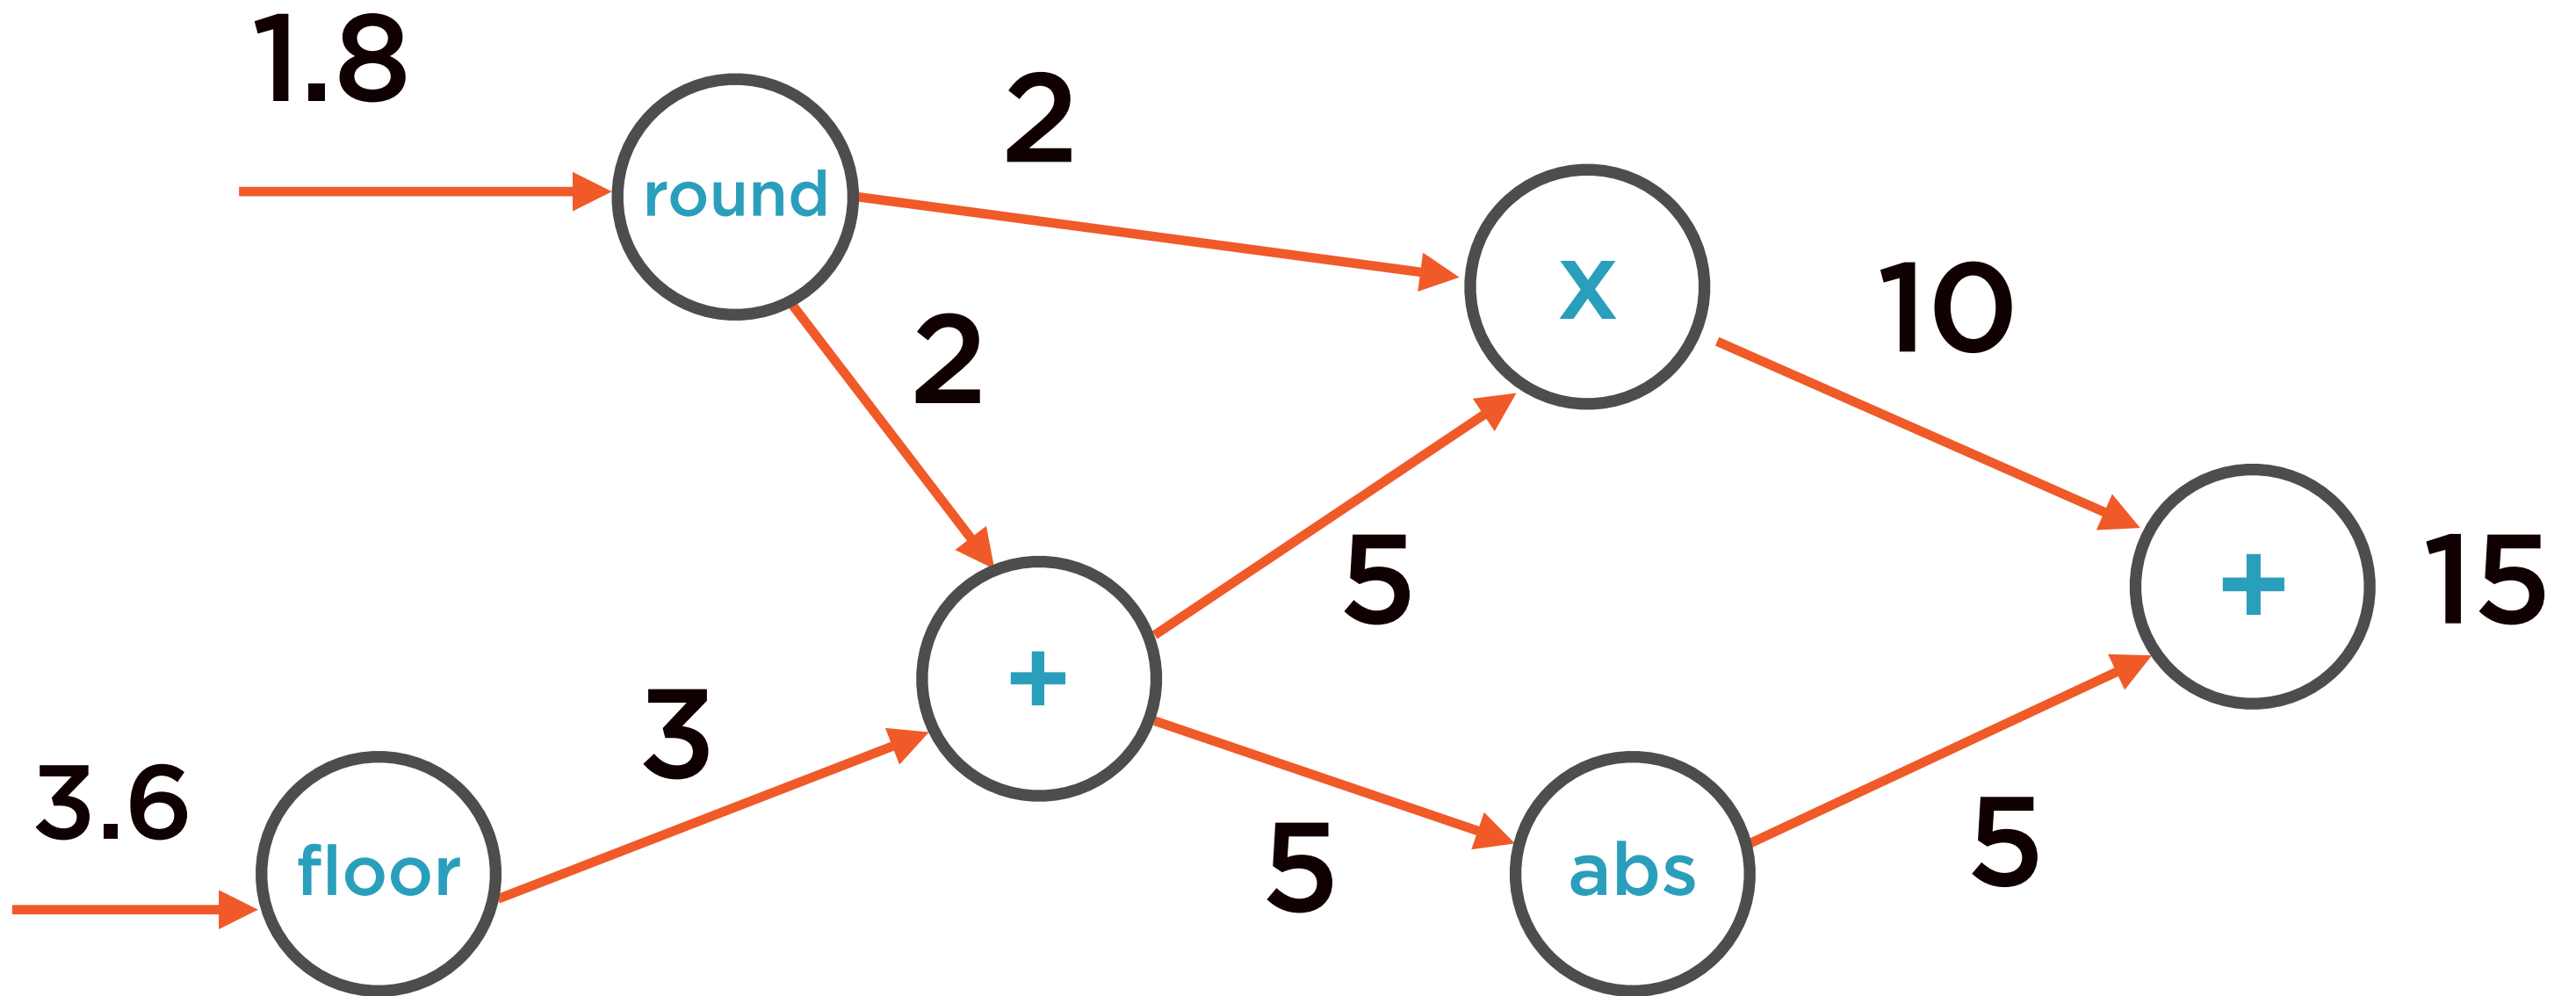
**Tensors**

# Tensors Flow Through the Graph




**...and get transformed along the way**

# Tensors Flow Through the Graph





# TensorFlow: “Define, then Run”



## Building a Graph


Specify the operations and  
the data



## Running a Graph

Execute the graph to get the  
final result

# PyTorch: “Define by Run”



## Building a Graph

Specify the operations and  
the data



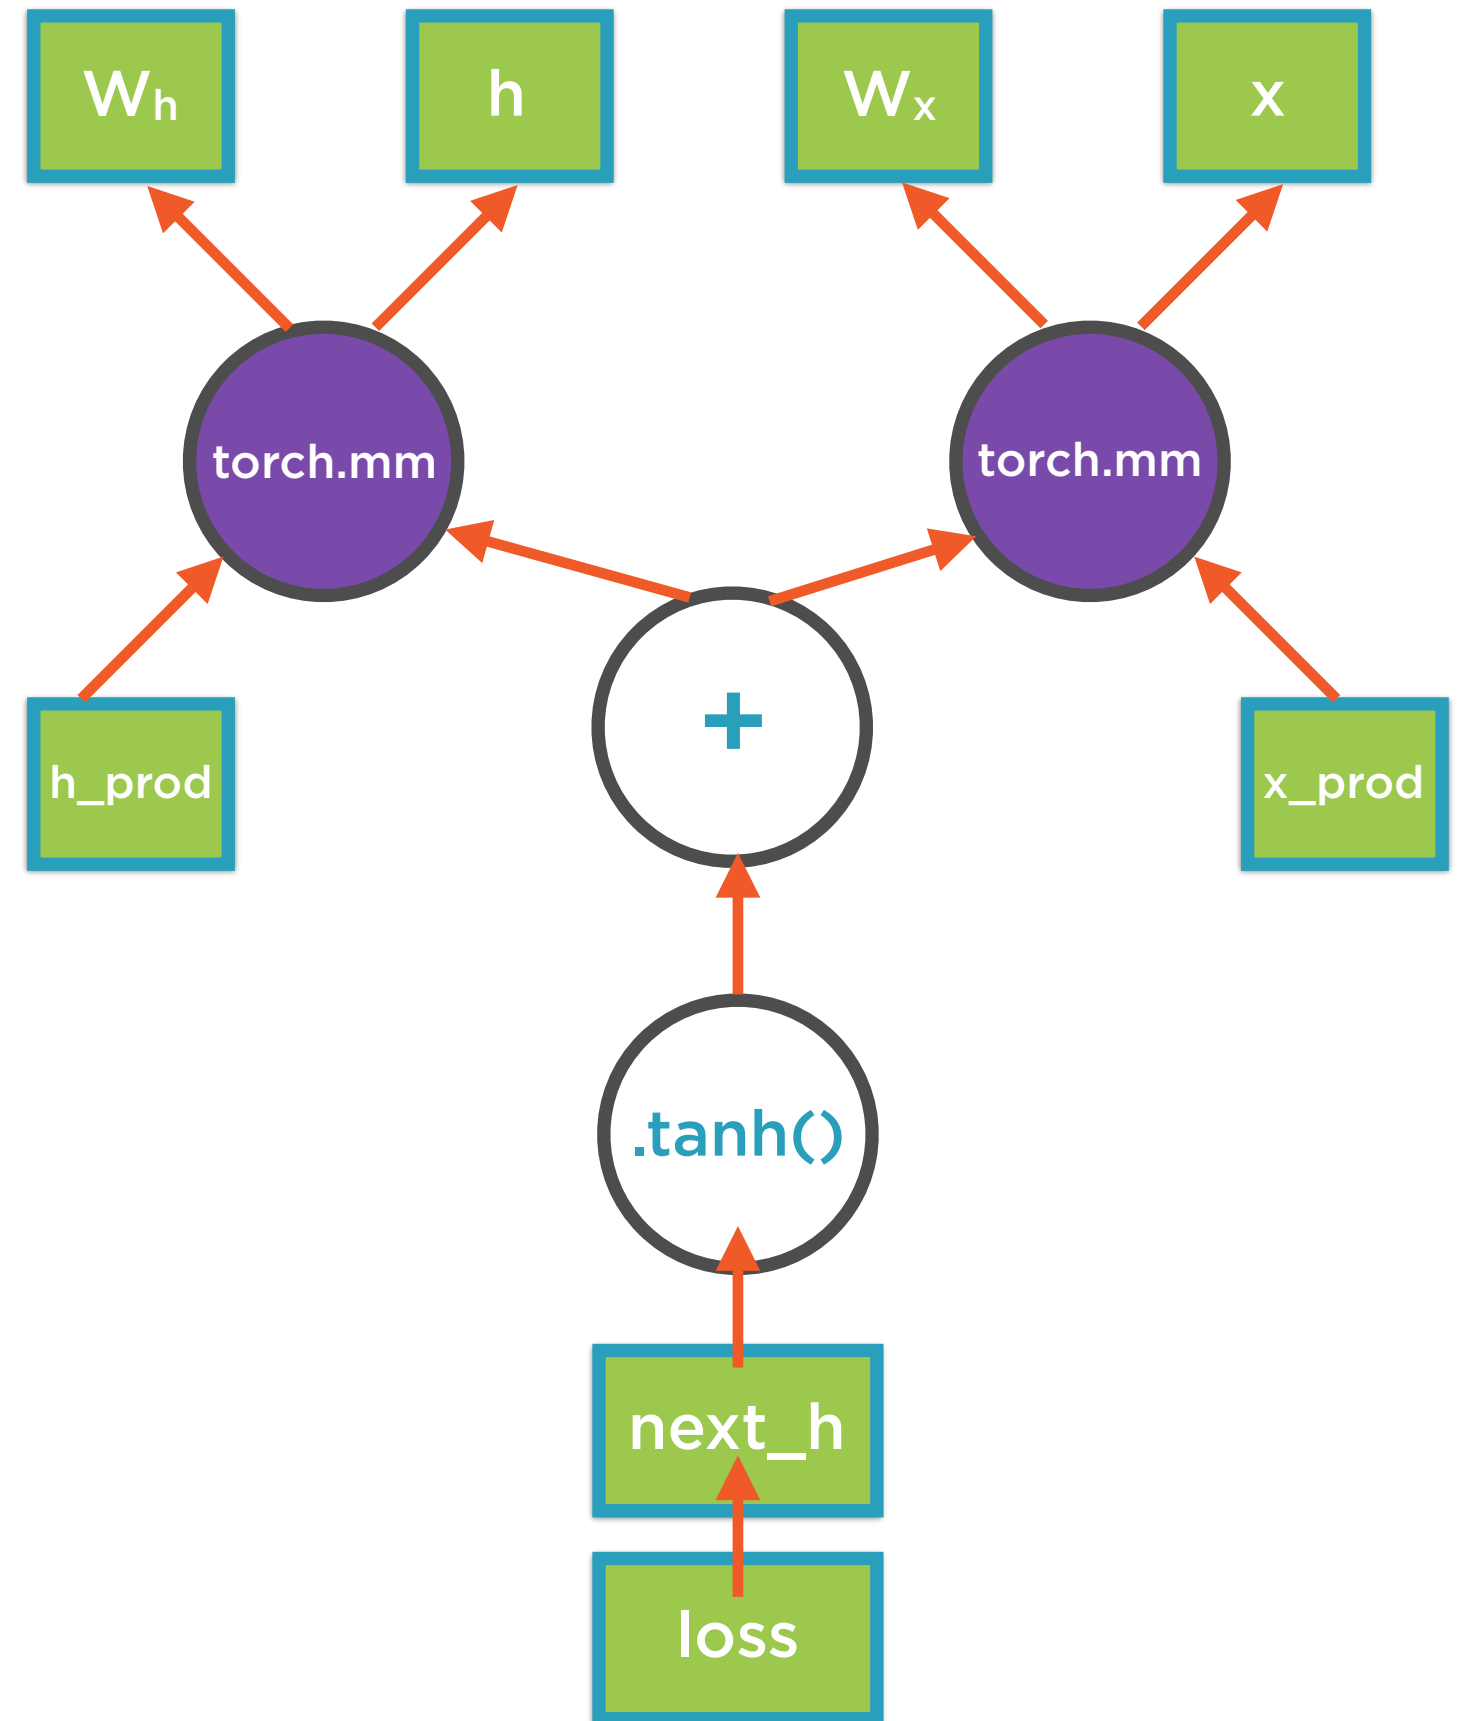
## Running a Graph

Execute the graph to get the  
final result

Build and execute the graph in one  
go - execute as you build

```
from torch.autograd import Variable
x = Variable(torch.randn(1,10))
h = Variable(torch.randn(1,20))
W_h = Variable(torch.randn(20,20))
W_x = Variable(torch.randn(20,10))


h_prod = torch.mm(W_h,h.t())
x_prod = torch.mm(W_x,x.t())
next_h = (h_prod + x_prod).tanh()
loss = next_h.sum()
loss.backward()
```



# Training an NN Using Gradient Descent

---


# Neural Network Model



**Network of interconnected layers**

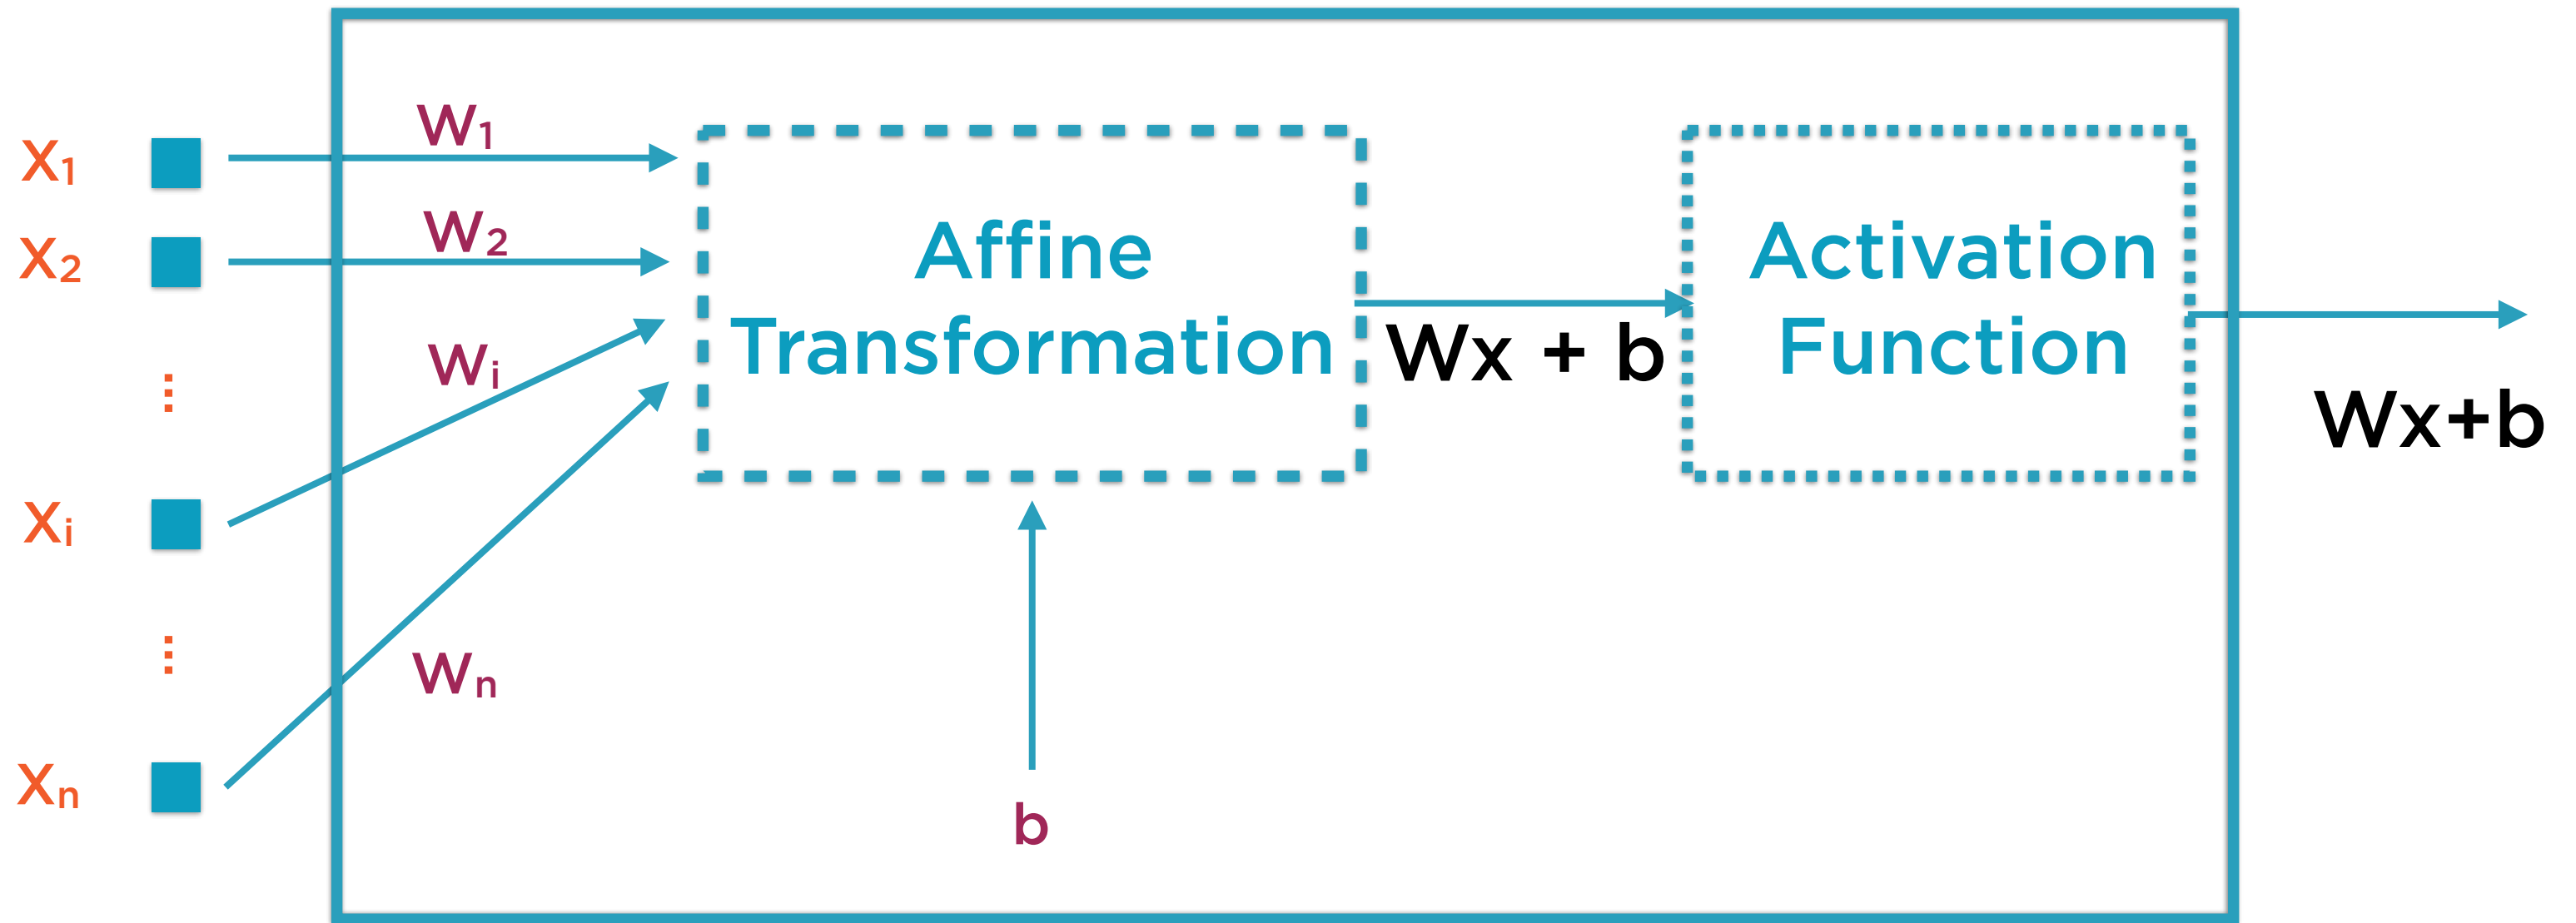
The **weights** and **biases** of individual neurons are determined during the **training** process

# Regression: The Simplest Neural Network

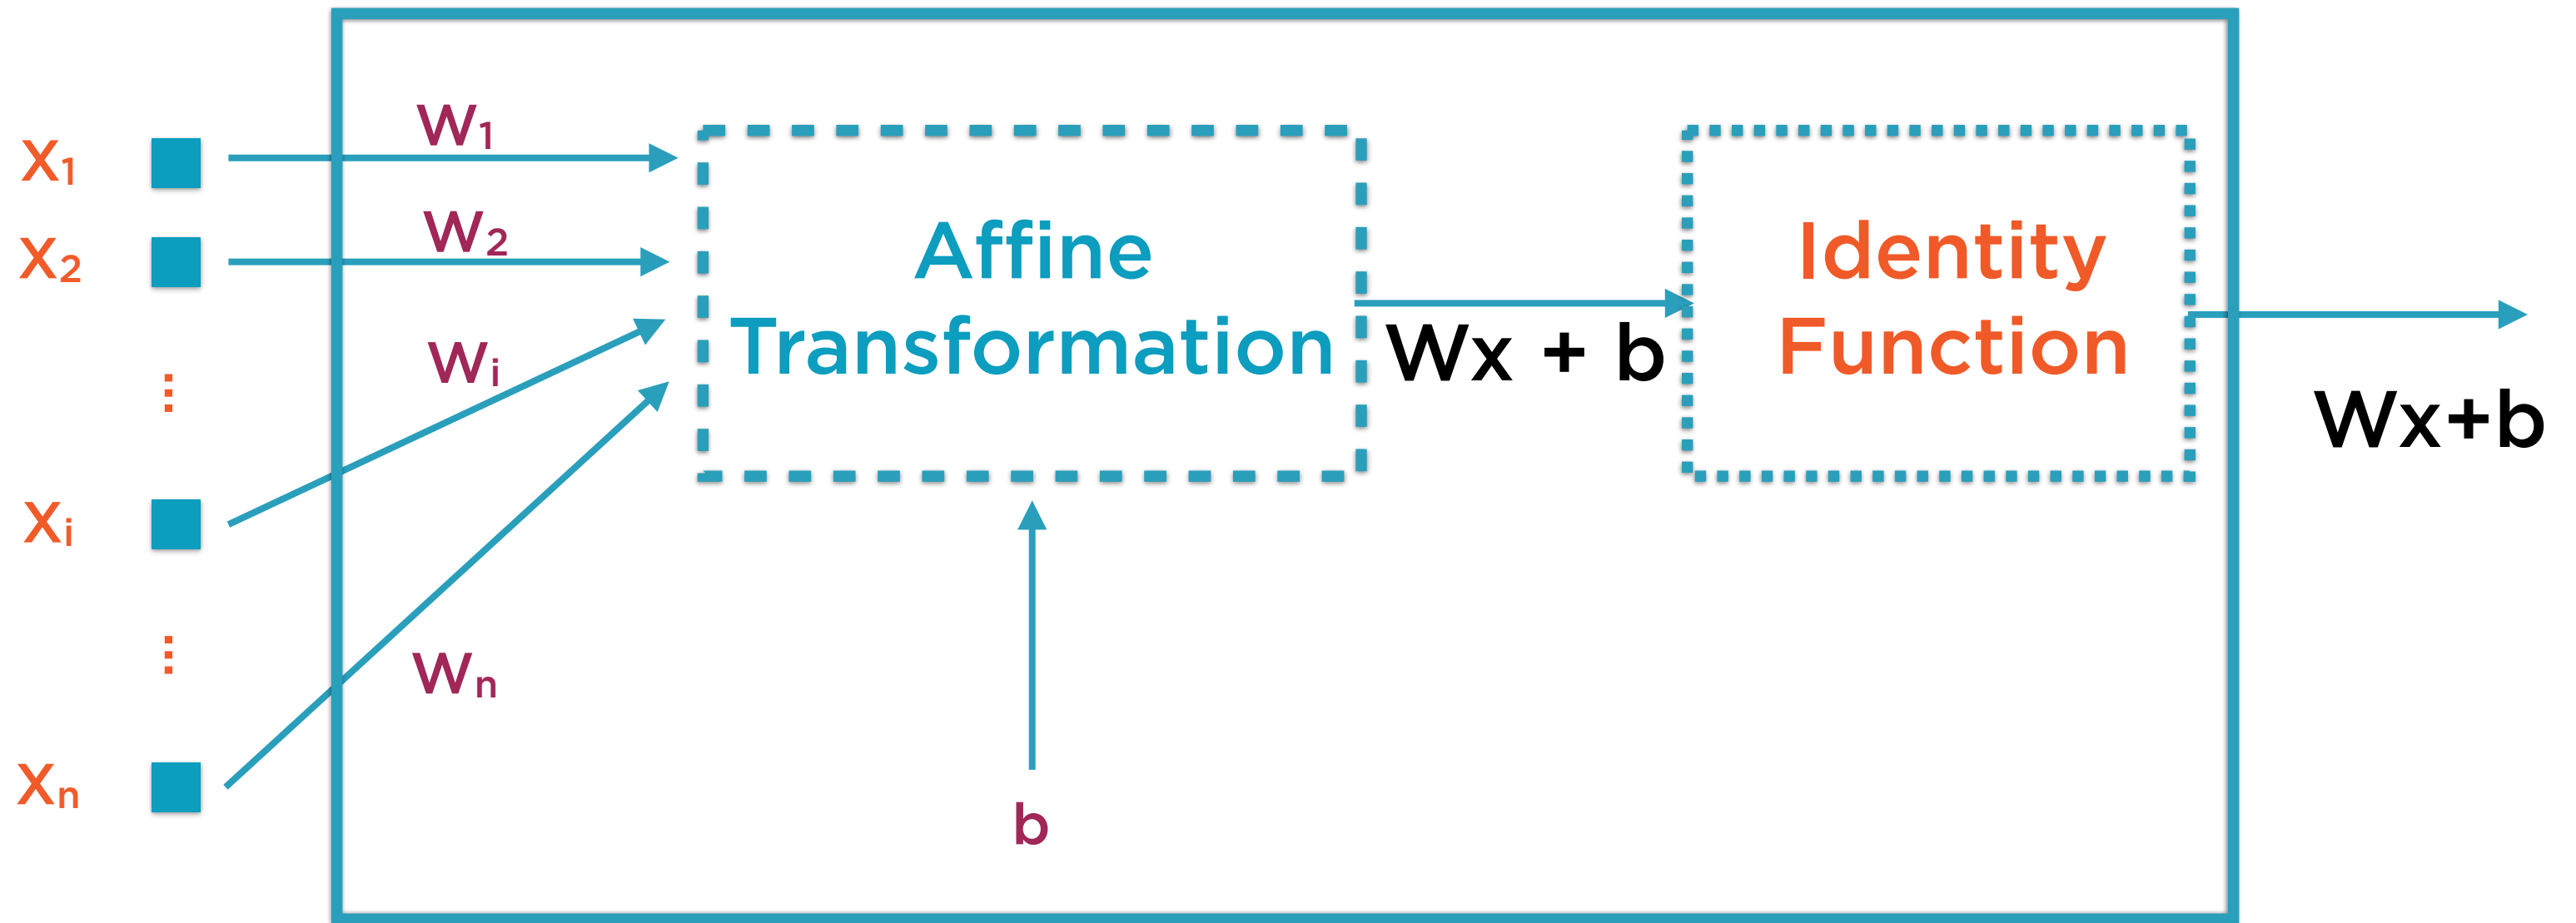




# Regression: The Simplest Neural Network




# Regression: The Simplest Neural Network



# Minimising Least Square Error



Y



Regression Line:  
 $y = A + Bx$

X



The “best fit” line is called the  
regression line

The actual training of a neural network happens via Gradient Descent Optimization

# Linear Regression as an Optimization Problem



## Objective Function

Minimize variance of  
the residuals (MSE)

# Linear Regression as an Optimization Problem



## Objective Function

Minimize variance of  
the residuals (MSE)



## Constraints

Express relationship as  
a straight line

$$y = Wx + b$$

# Linear Regression as an Optimization Problem



## Objective Function


Minimize variance of  
the residuals (MSE)



## Constraints

Express relationship as  
a straight line


$$y = Wx + b$$



## Decision Variables


Values of  $W$  and  $b$

# Minimizing MSE







# Minimizing MSE




# Minimizing MSE




# Minimizing MSE



# Minimizing MSE




# Minimizing MSE




# “Gradient Descent”

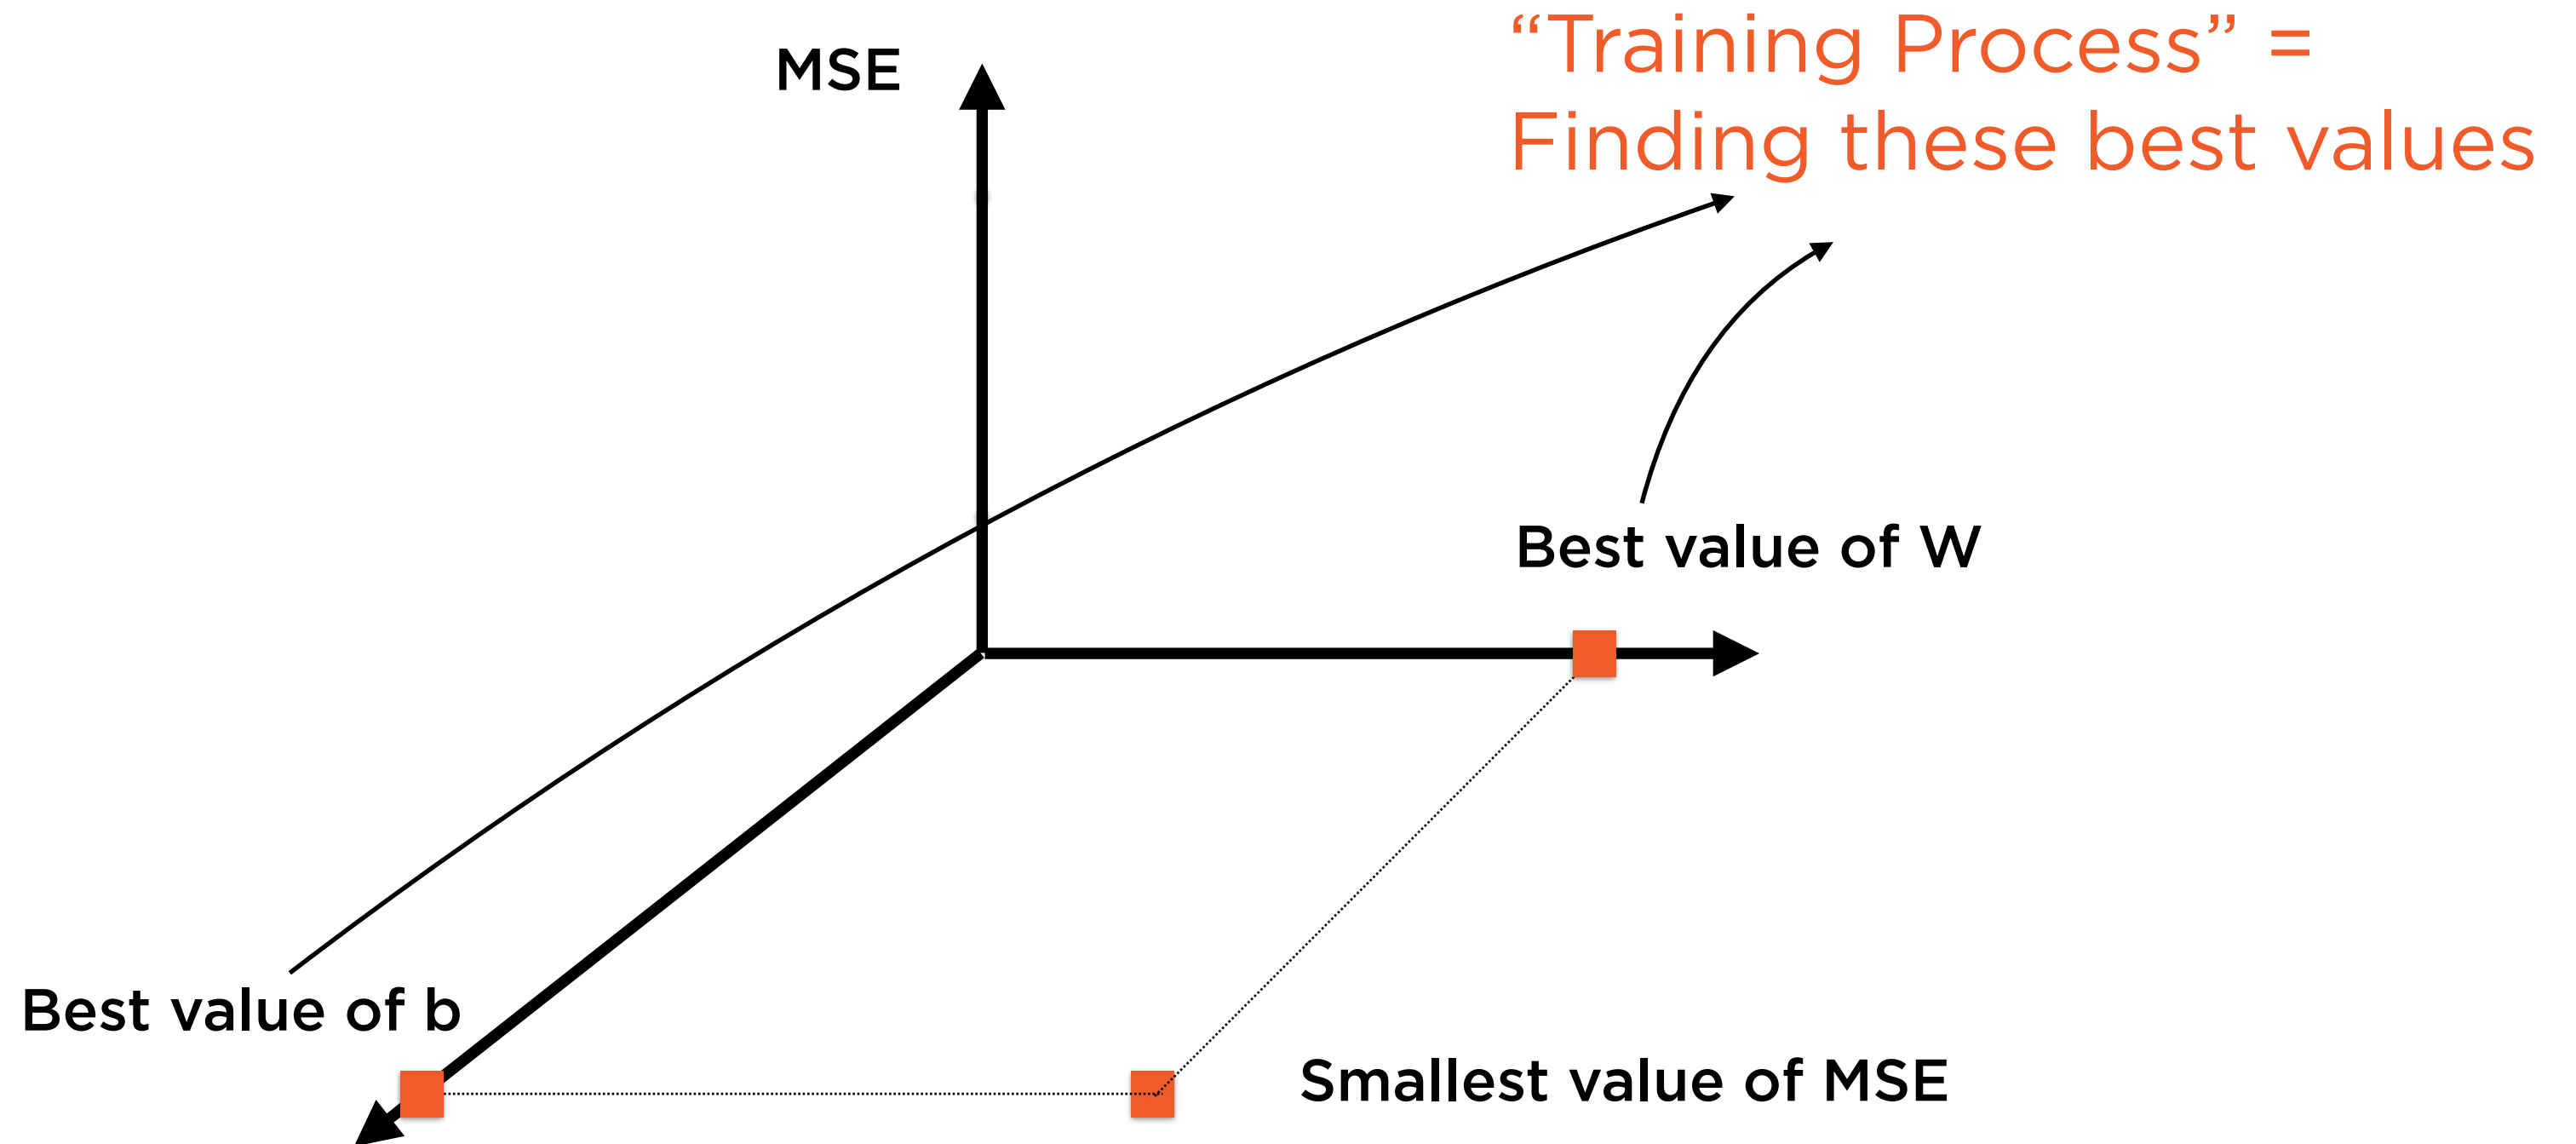
Converging on the “best”  
value using an  
optimization algorithm



# Minimizing MSE




# “Training” the Algorithm






# Start Somewhere



# “Gradient Descent”


Converging on the “best”  
value using an  
optimization algorithm



# Forward and Backward Passes

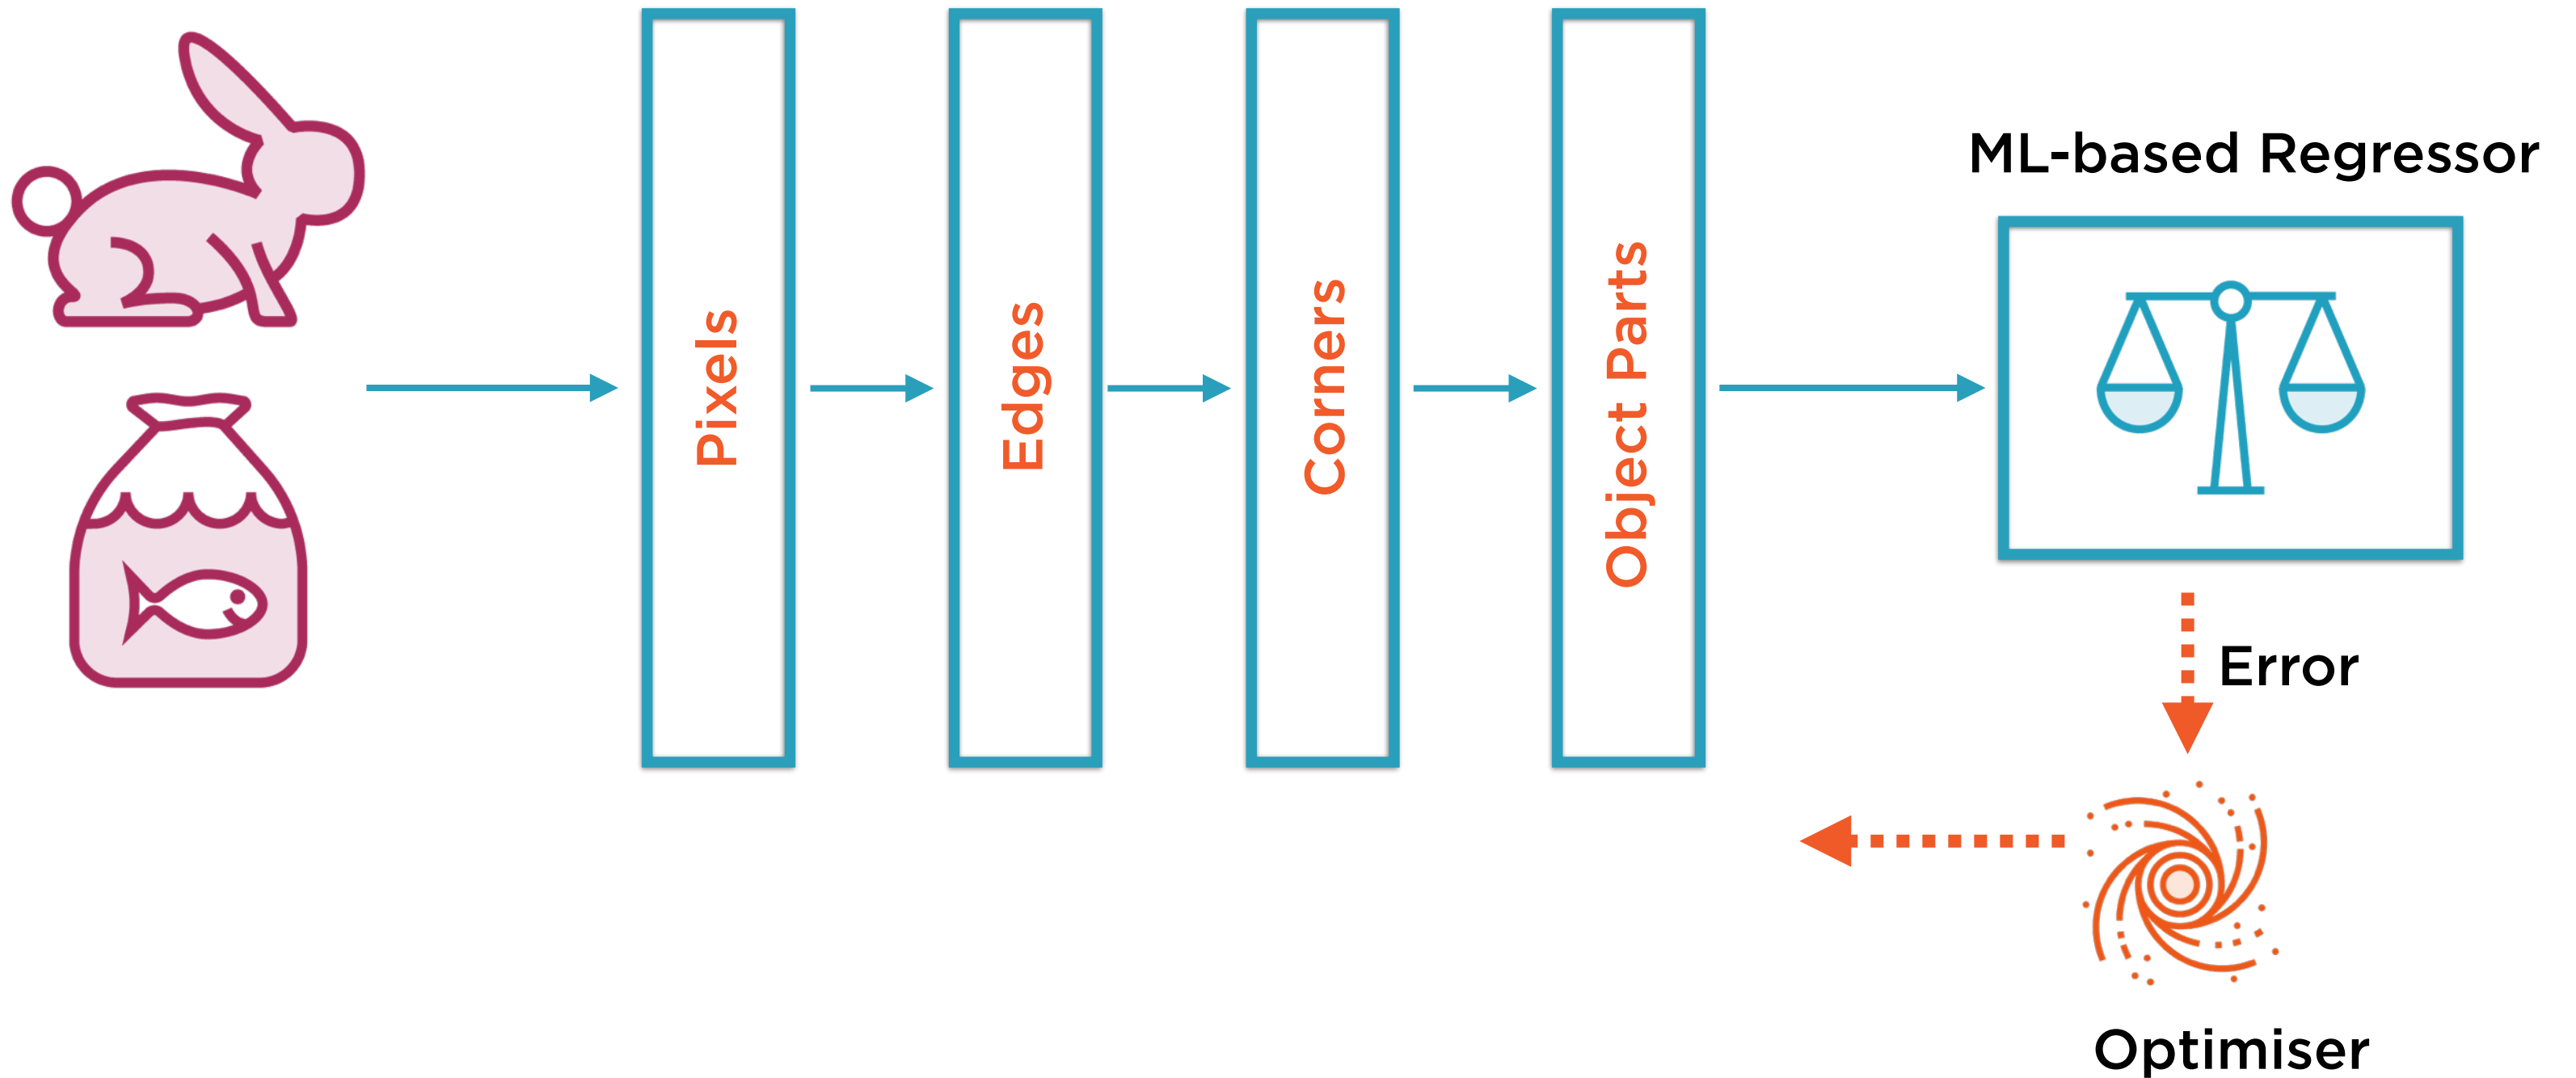
---

# Forward Pass

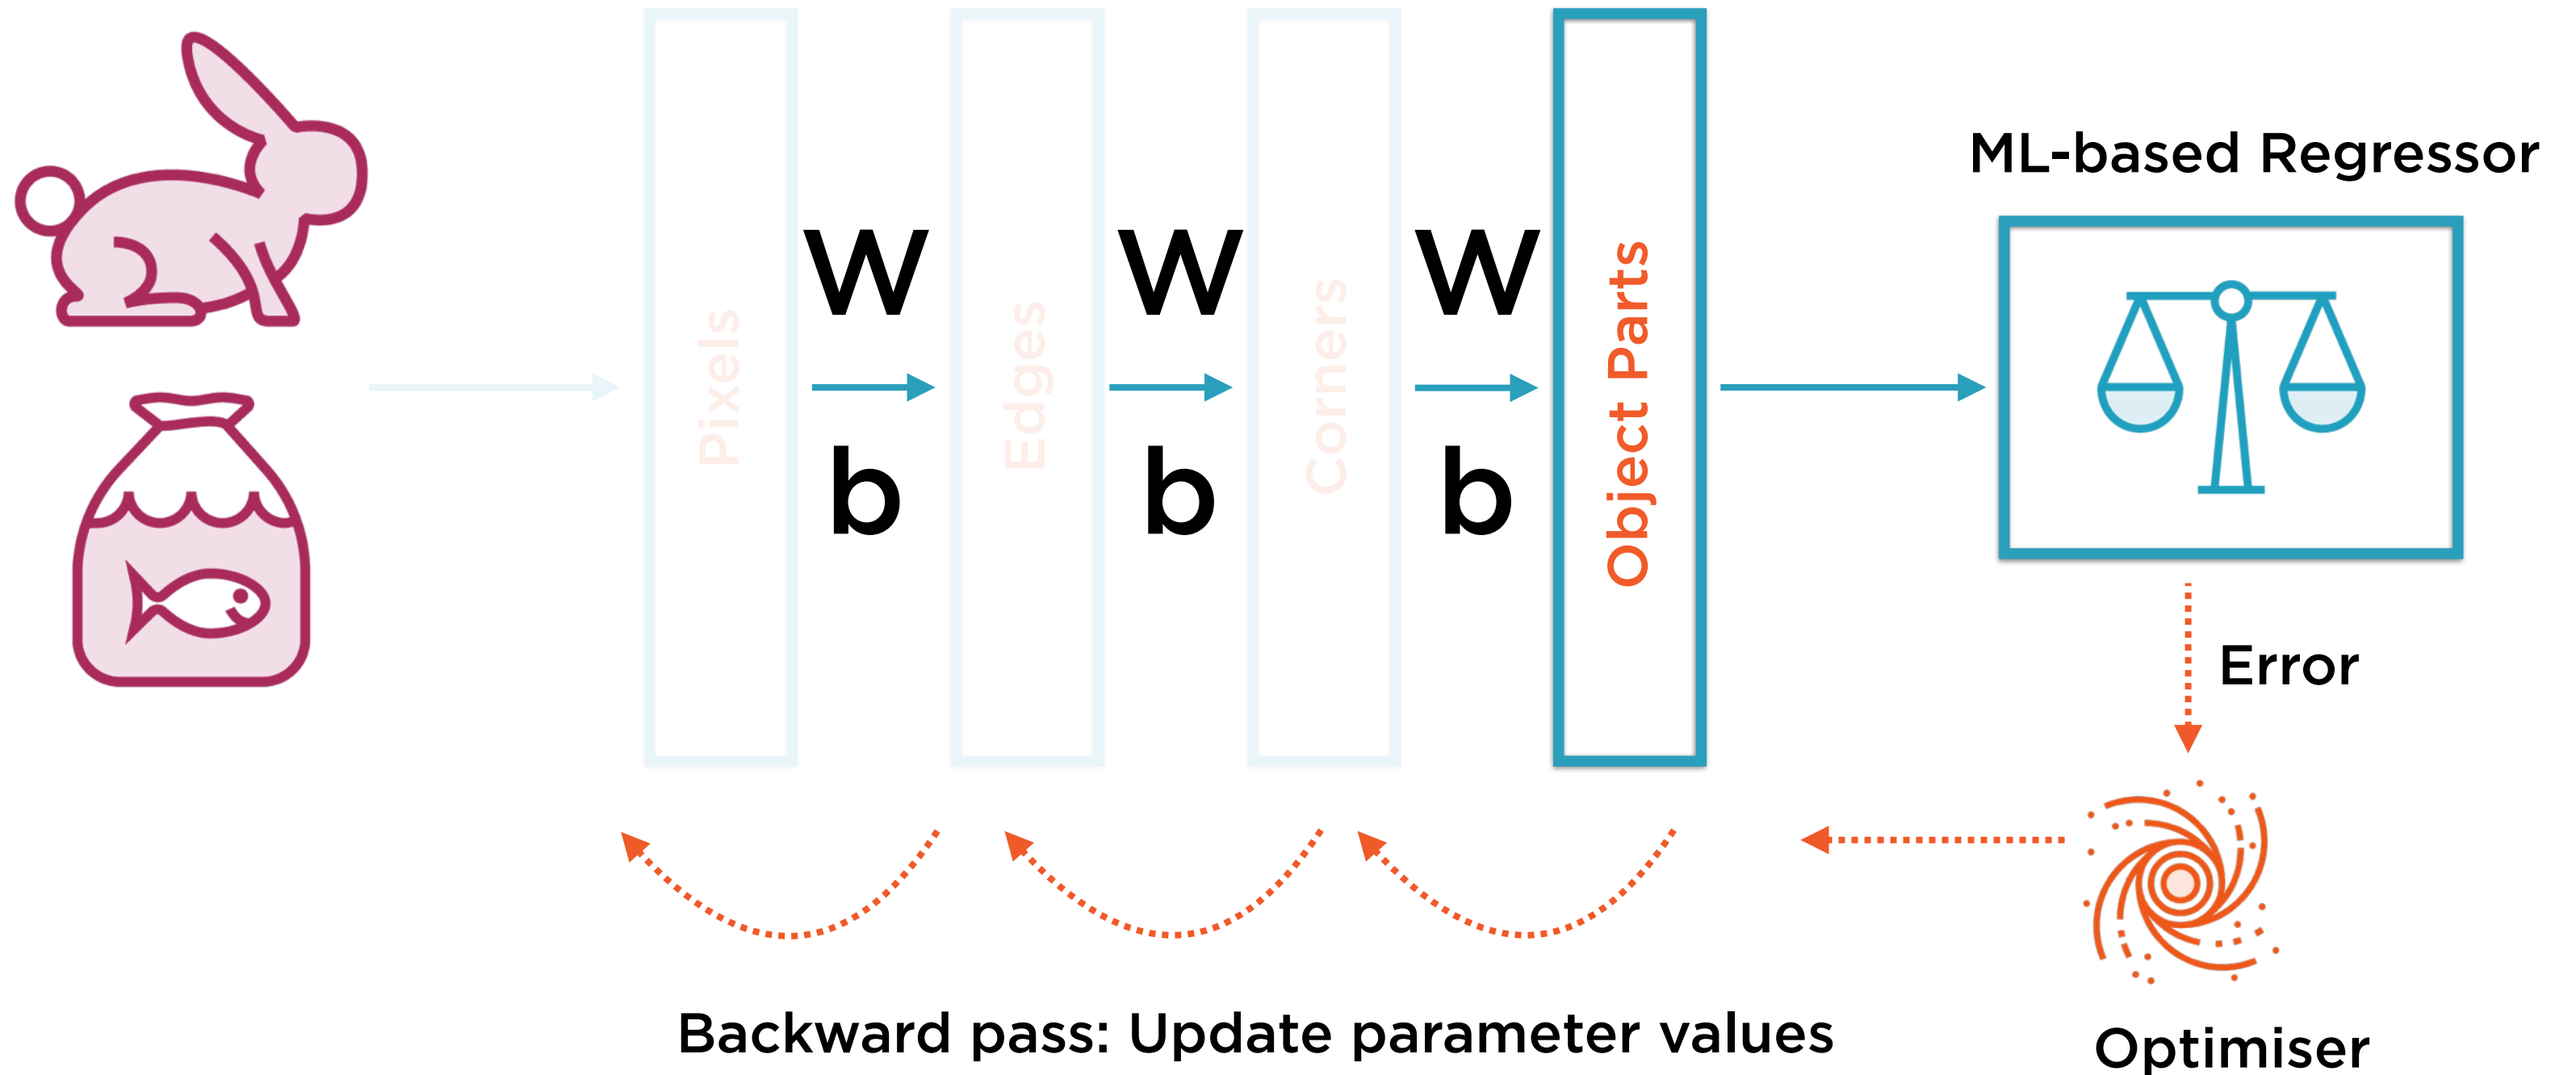


Forward pass: Calculate  $y_{\text{predicted}}$  and error


# Backward Pass




# Backward Pass



# Backward Pass




# Backward Pass






# Backward Pass



# Backward Pass



The backward pass allows the weights and biases of the neurons to converge to their final values

# Overview

**PyTorch is a deep learning framework**

**More tightly integrated with Python than TensorFlow**

**Can use Python libraries, debugger**

**Supports dynamic computation graphs, update the graph for each epoch**

**Uses a forward pass for prediction, backward pass to update weights**