

Building an Image Classification Model



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Convolutional NNs are a deep learning technique easily implemented in PyTorch

ResNet is a famous CNN architecture

Transfer learning is a great way to re-use pre-trained models

PyTorch offers great support for transfer learning

Image classification using transfer learning and ResNet

How Do We See?


“Sometimes the mind can see what is invisible to the eye”

Viewing an Image




**All neurons in the eye don't see the
entire image**

Viewing an Image




**Each neuron has its own local
receptive field**

Viewing an Image




**It reacts only to visual stimuli
located in its receptive field**

Viewing an Image



Some neurons react to more complex patterns
that are **combinations** of lower level patterns

Neural Networks



Sounds like a classic neural network
problem

Two Kinds of Layers in CNNs

Convolution

Local receptive field

Pooling


Subsampling of inputs

Convolution

Convolution


In this context, a sliding window function applied to a matrix

Representing Images as Matrices



= 784 pixels

Representing Images as Matrices



= 36 pixels

Representing Images

6

3

6

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix

3

1	0	1
0	1	0
1	0	1

Kernel

Convolution

6

6

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix

3

3

x1	x0	x1
x0	x1	x0
x1	x0	x1

Kernel

Convolution

6

6

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix

→

x1	x0	x1
x0	x1	x0
x1	x0	x1

4

4

1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution
Result

Convolution

0 _{x1}	x0	0 _{x1}	0	0	0
x0	0.8 _{x1}	x0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Convolution
Result

Convolution

0 _{x1}	x0	0 _{x1}	0	0	0
x0	0.8 _{x1}	x0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution
Result

Convolution

0	0 _{x1}	x0	0 _{x1}	0	0
0.2	x0	0 _{x1}	x0	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution
Result

Convolution

0	0 _{x1}	x0	0 _{x1}	0	0
0.2	x0	0 _{x1}	x0	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution
Result

Convolution

0	0	0 _{x1}	x0	0 _{x1}	0
0.2	0.8	x0	0.3 _{x1}	x0	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution
Result

Convolution

0	0	0 _{x1}	x0	0 _{x1}	0
0.2	0.8	x0	0.3 _{x1}	x0	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution
Result

Convolution

0	0	0	0 _{x1}	x0	0 _{x1}
0.2	0.8	0	x0	0.6 _{x1}	x0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution
Result

Convolution

0	0	0	0 _{x1}	x0	0 _{x1}
0.2	0.8	0	x0	0.6 _{x1}	x0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.6	0
x0	0.9 _{x1}	x0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	x0	0.3 _{x1}	0.6	0
0.2	x0	0 _{x1}	x0	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8 _{x1}	x0	0.3 _{x1}	0.6	0
0.2	x0	0 _{x1}	x0	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0 _{x1}	x0	0.6 _{x1}	0
0.2	0.9	x0	0.3 _{x1}	x0	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0 _{x1}	x0	0.6 _{x1}	0
0.2	0.9	x0	0.3 _{x1}	x0	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 _{x1}	x0	0 _{x1}
0.2	0.9	0	x0	0.8 _{x1}	x0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
x0	0.8 _{x1}	x0	0.8	0.9	0
0 _{x1}	x0	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 _{x1}	x0	0 _{x1}	0.3	0.8	0
x0	0.8 _{x1}	x0	0.8	0.9	0
0 _{x1}	x0	0 _{x1}	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9_{x1}	x0	0.3_{x1}	0.8	0
0.3	x0	0.7_{x1}	x0	0.9	0
0	0_{x1}	x0	0.2_{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9 _{x1}	x0	0.3 _{x1}	0.8	0
0.3	x0	0.7 _{x1}	x0	0.9	0
0	0 _{x1}	x0	0.2 _{x1}	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	x0	0.8 _{x1}	x0	0
0	0	0 _{x1}	x0	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 _{x1}	x0	0.8 _{x1}	0
0.3	0.8	x0	0.8 _{x1}	x0	0
0	0	0 _{x1}	x0	0.8 _{x1}	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	x0	0.9 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 _{x1}	x0	0 _{x1}
0.3	0.8	0.7	x0	0.9 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
x0	0 _{x1}	x0	0.2	0.8	0
0 _{x1}	x0	0 _{x1}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 _{x1}	x0	0.7 _{x1}	0.8	0.9	0
x0	0 _{x1}	x0	0.2	0.8	0
0 _{x1}	x0	0 _{x1}	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	x0	0 _{x1}	x0	0.8	0
0	0 _{x1}	x0	0.2 _{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 _{x1}	x0	0.8 _{x1}	0.9	0
0	x0	0 _{x1}	x0	0.8	0
0	0 _{x1}	x0	0.2 _{x1}	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	x0	0.2 _{x1}	x0	0
0	0	0 _{x1}	x0	0.2 _{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 _{x1}	x0	0.9 _{x1}	0
0	0	x0	0.2 _{x1}	x0	0
0	0	0 _{x1}	x0	0.2 _{x1}	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	x0	0.8 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution
Result

Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 _{x1}	x0	0 _{x1}
0	0	0	x0	0.8 _{x1}	x0
0	0	0	0.2 _{x1}	x0	0 _{x1}

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution
Result

Convolutional Layers

Parameter Explosion



Consider a 100 x 100 pixel image (10,000 pixels)

If first layer = 10,000 neurons

Interconnections ~ $O(10,000 * 10,000)$

100 million parameters to train neural network!

Parameter Explosion



Dense, fully connected neural networks can't cope


Convolutional neural networks to the rescue

Inspirations for CNNs



Two Dimensions

Data comes in expressed
in 2D



Local Receptive Fields

Neurons focus on narrow
portions

CNN Layers



Convolution layers - zoom in on specific bits of input


Successive layers aggregate inputs into higher level features

Pixels >> Lines >> Contours/Edges >> Object


Feature Maps



Image




Pixels




Feature
Map

Feature maps are convolutional layers generated by applying a convolutional kernel to the input

Feature Maps




Pixels




Convolutional
Layer

Neurons




Feature Maps



Number of neurons
in receptive field =
kernel size

Feature Maps

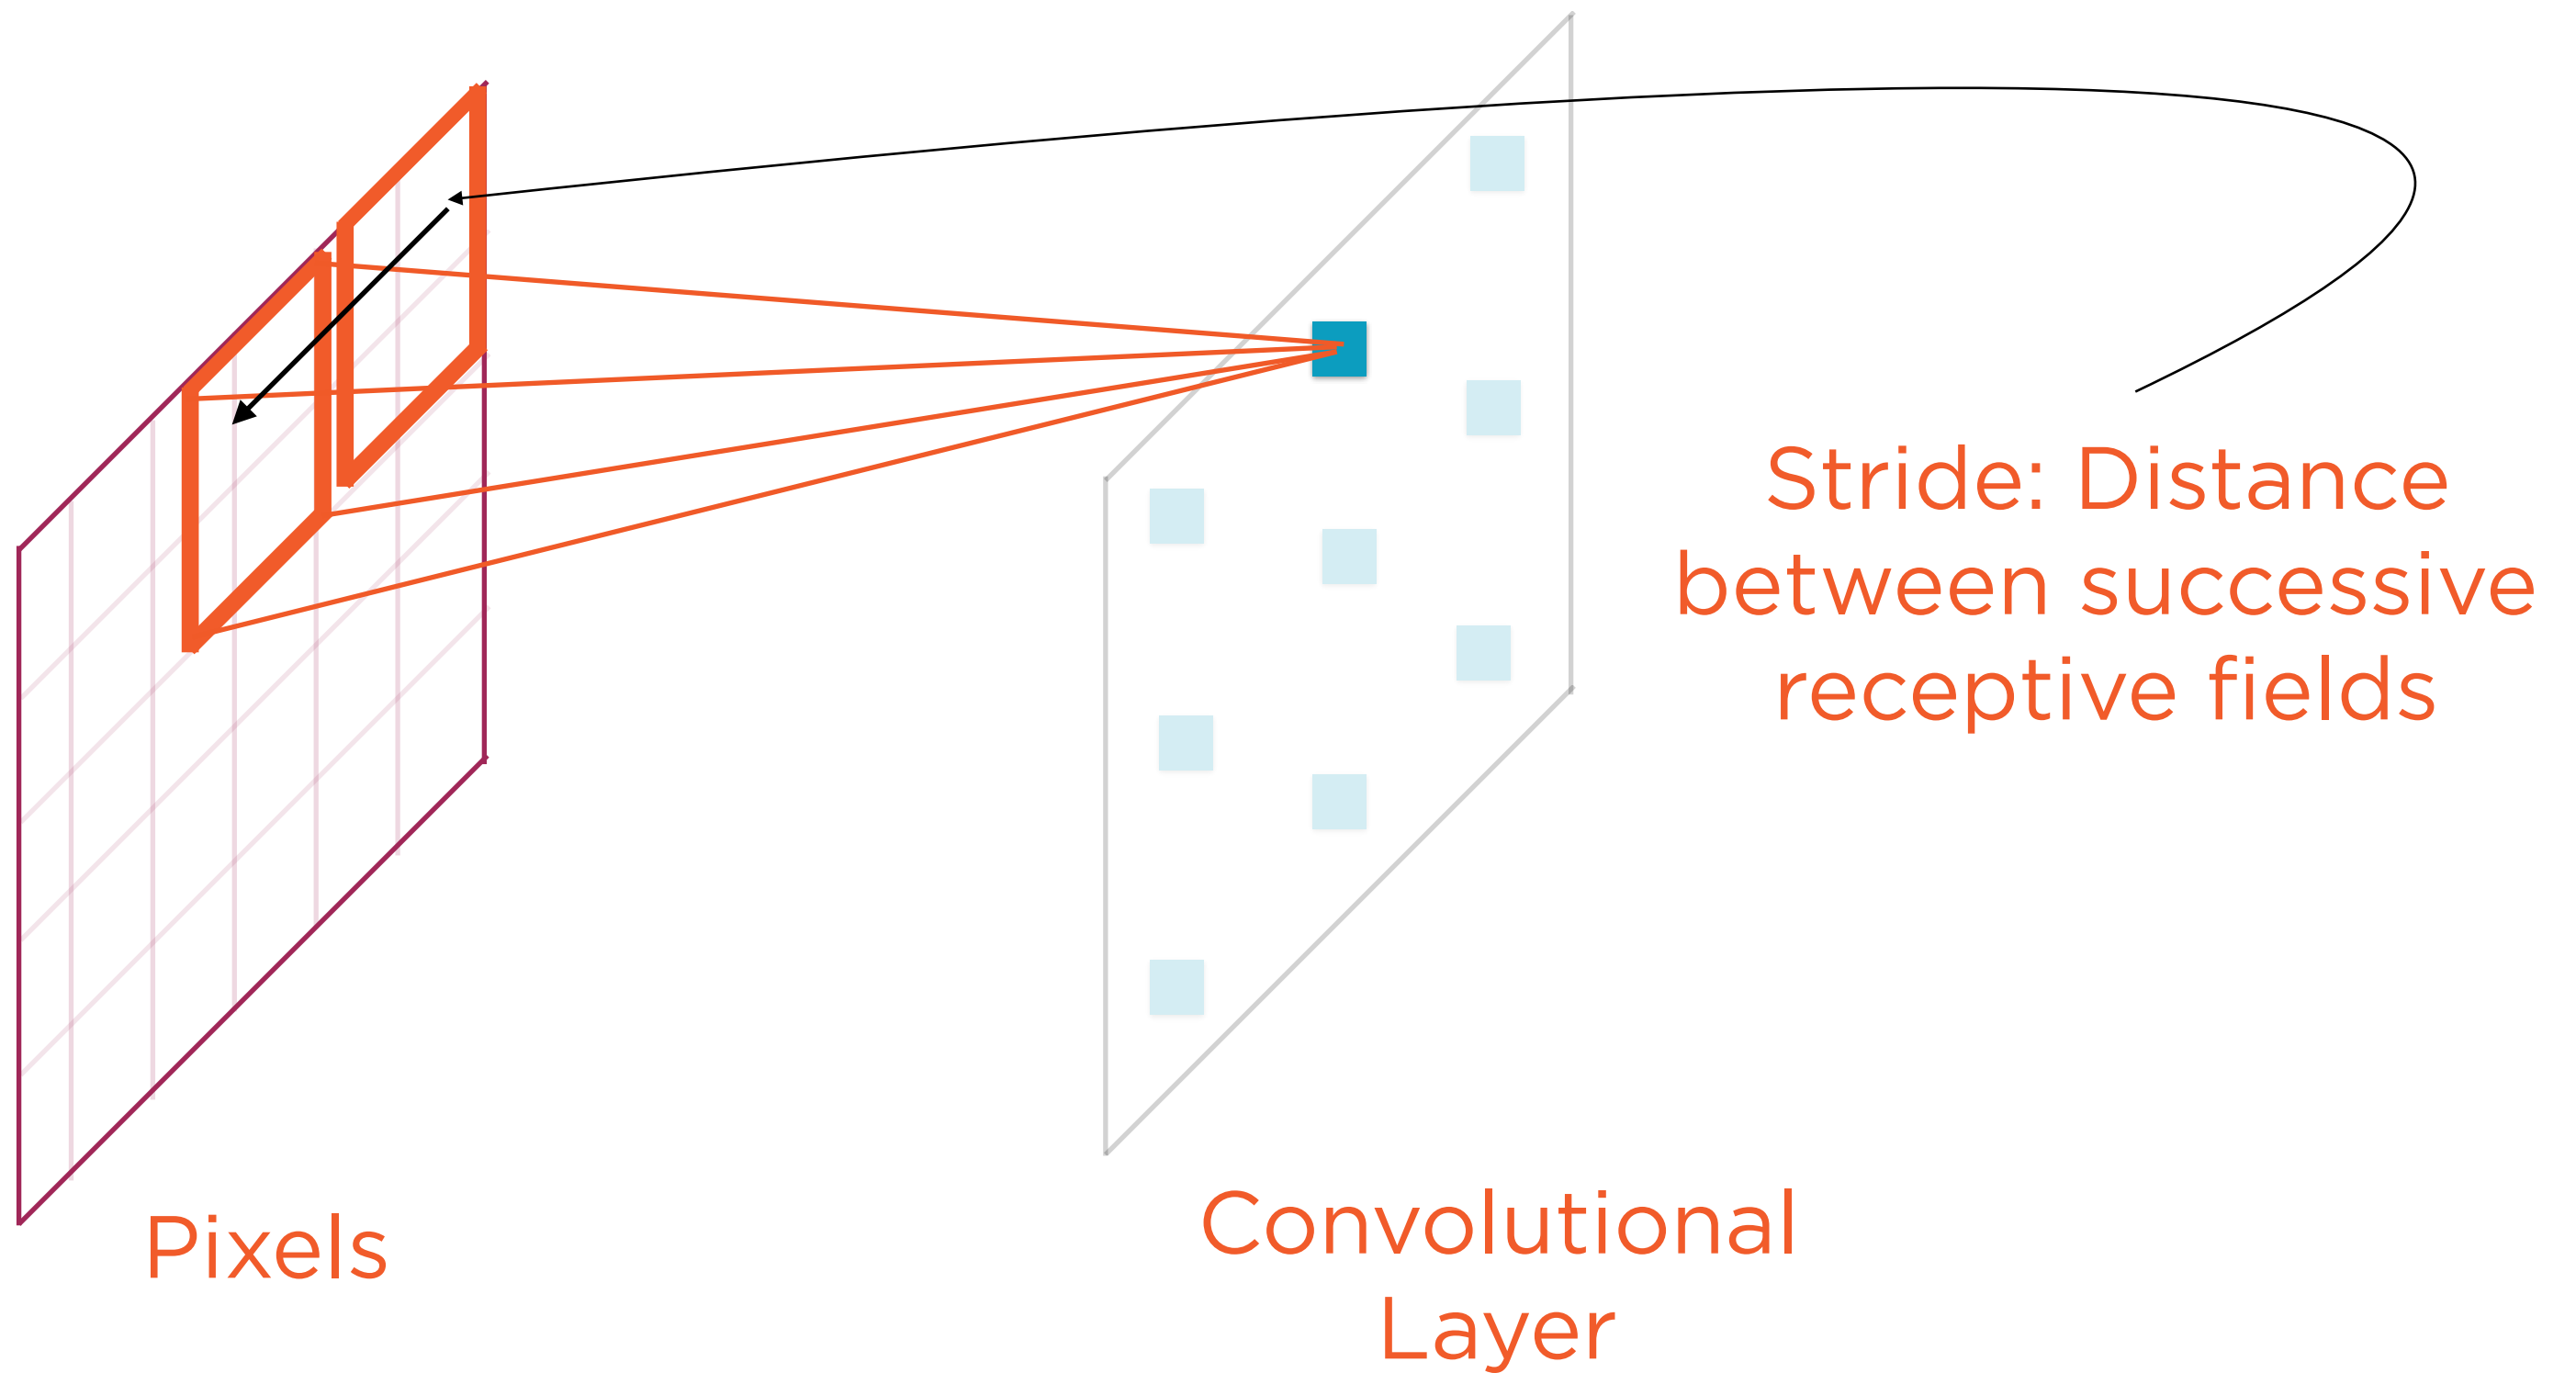
Neuron i




Pixels

Convolutional
Layer


Feature Maps




Feature Maps




Feature Maps



Feature Maps




Feature Maps




Feature Maps

Zero padding
may be needed
at the edges




Feature Maps




Sparse, not
Dense

Feature Maps



Notice also that neurons are not connected to all pixels

CNNs are **sparse** neural networks




Convolutional Layer


Each convolutional layer consists of several feature maps of equal sizes

The different feature maps have different parameters


CNNs



Feature
Map



Convolutional
Layer



CNN

Pooling Layers

Pooling

4

4

0.2	0.8	0.3	0.6
0.2	0.9	0.3	0.8
0.3	0.8	0.8	0.9
0	0	0.2	0.8

Matrix



Max,
2x2 filter,
stride = 2


2

2

0.9	0.8
0.8	0.9

Pooling
Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling
Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling
Result


Pooling



Matrix




Max,
2x2 filter,
stride = 2



Pooling
Result


Pooling




Matrix



Max,
2x2 filter,
stride = 2



Pooling
Result




Pooling Layers

Neurons in a pooling layer have no weights or biases

A pooling neuron simply applies some aggregation function to all inputs

Max, sum, average...




Pooling Layers

Why use them?


- greatly reduce memory usage during training
- mitigate overfitting (via subsampling)
- make NN recognize features independent of location (location invariance)

Batch Normalization

Training via Back Propagation



Vanishing and Exploding Gradients




Back propagation fails if

- gradients are **vanishing**
- gradients are **exploding**


“Vanishing Gradient Problem”

Gradient becomes zero
and stops changing



“Exploding Gradient Problem”

Gradient changes abruptly and “explodes”



Coping with Vanishing/Exploding Gradients

Proper initialisation

**Non-saturating activation
function**

Batch normalization

Gradient clipping

Coping with Vanishing/Exploding Gradients


Proper initialisation

Non-saturating activation
function

Batch normalization

Gradient clipping

Batch Normalization




Just before applying activation function

First, “normalize” inputs

Second, “scale and shift” inputs

Batch Normalization




“Normalize” inputs

- subtract mean
- divide by standard deviation

“Scale and shift” inputs

- scale = multiply by constant
- shift = add constant

Batch Normalization




Supported in PyTorch

Many other benefits

- allows much larger learn rate
- reduces overfitting
- speeds convergence of training


CNN Architectures

Typical CNN Architecture




Alternating groups of convolutional and pooling layers

Typical CNN Architecture




Each group of convolutional layers usually followed by a ReLU layer

Typical CNN Architecture




The outputs of each layer are also images

Typical CNN Architecture




However successive outputs are smaller and smaller (due to pooling layers)

Typical CNN Architecture



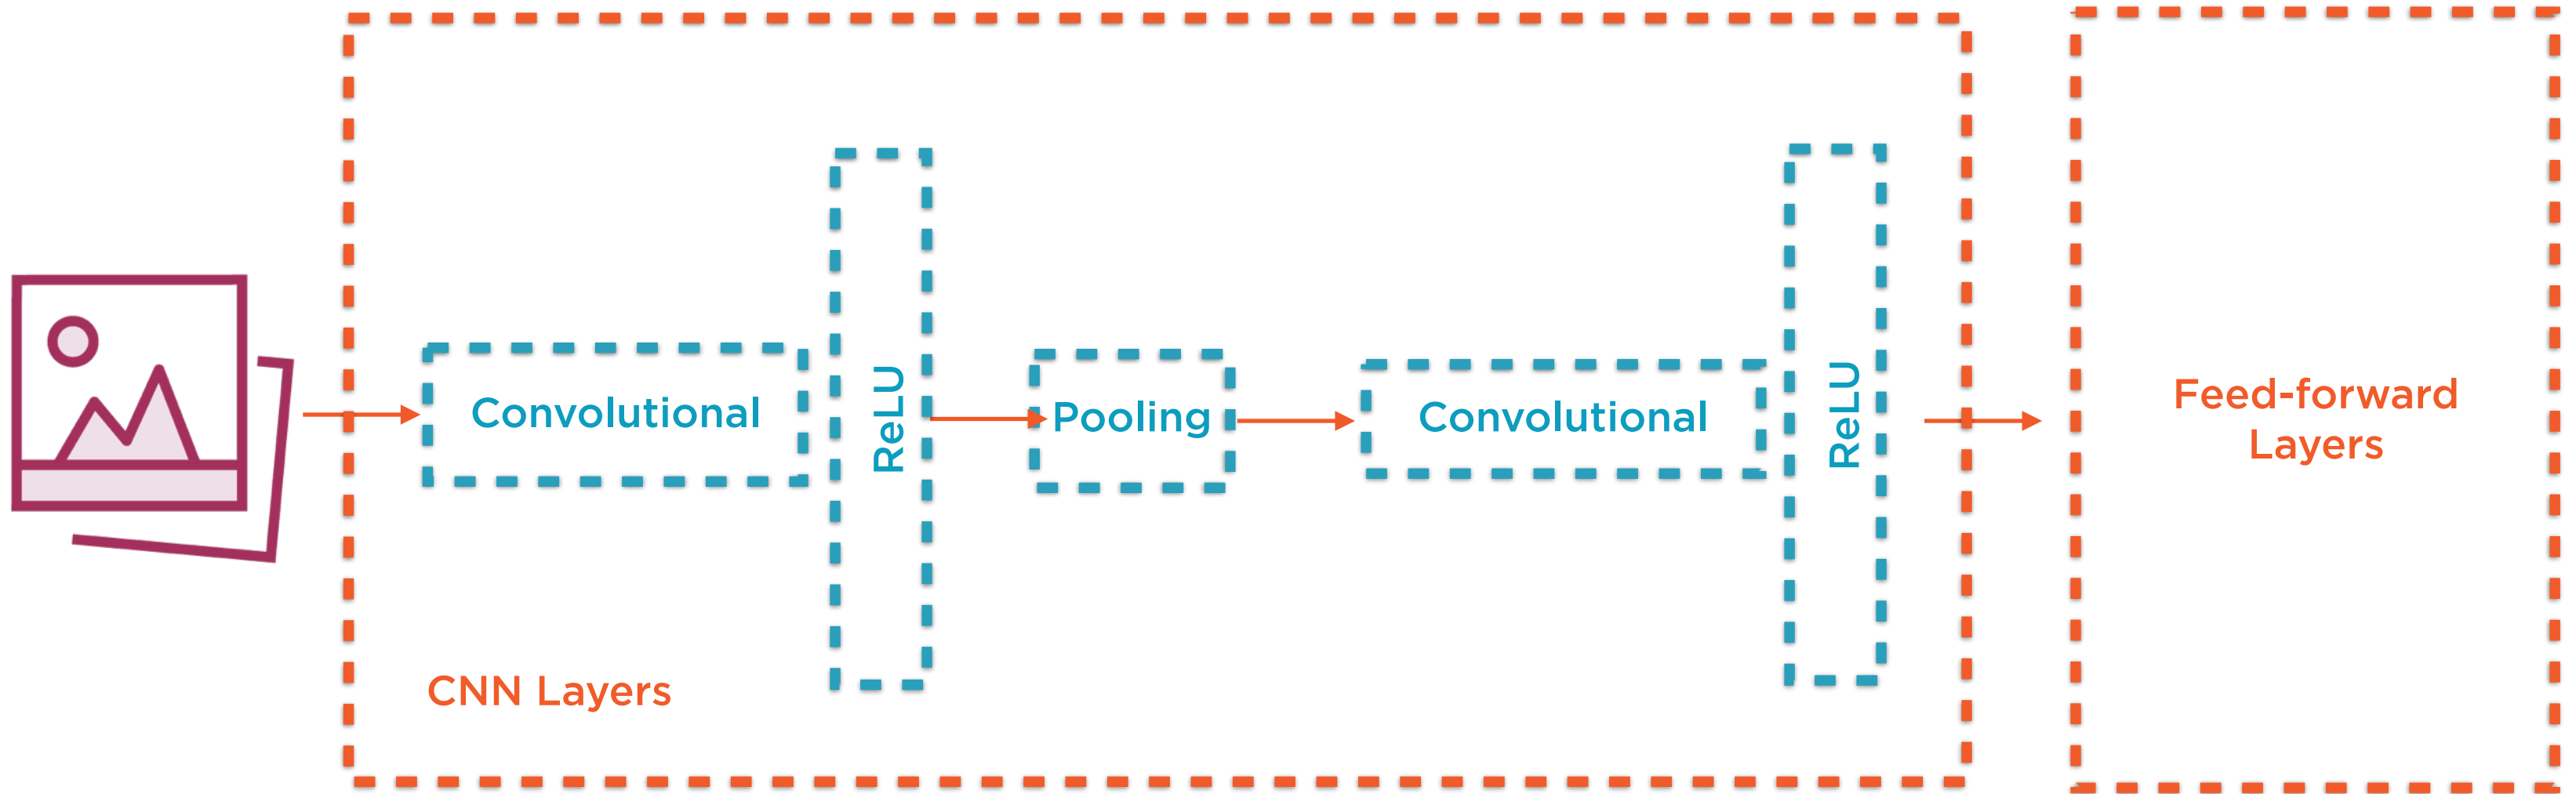
As well as deeper and deeper (due to feature maps in the convolutional layers)

Typical CNN Architecture




This entire set of layers is then fed into a regular, feed-forward NN

Typical CNN Architecture




This entire set of layers is then fed into a regular, feed-forward NN

Typical CNN Architecture



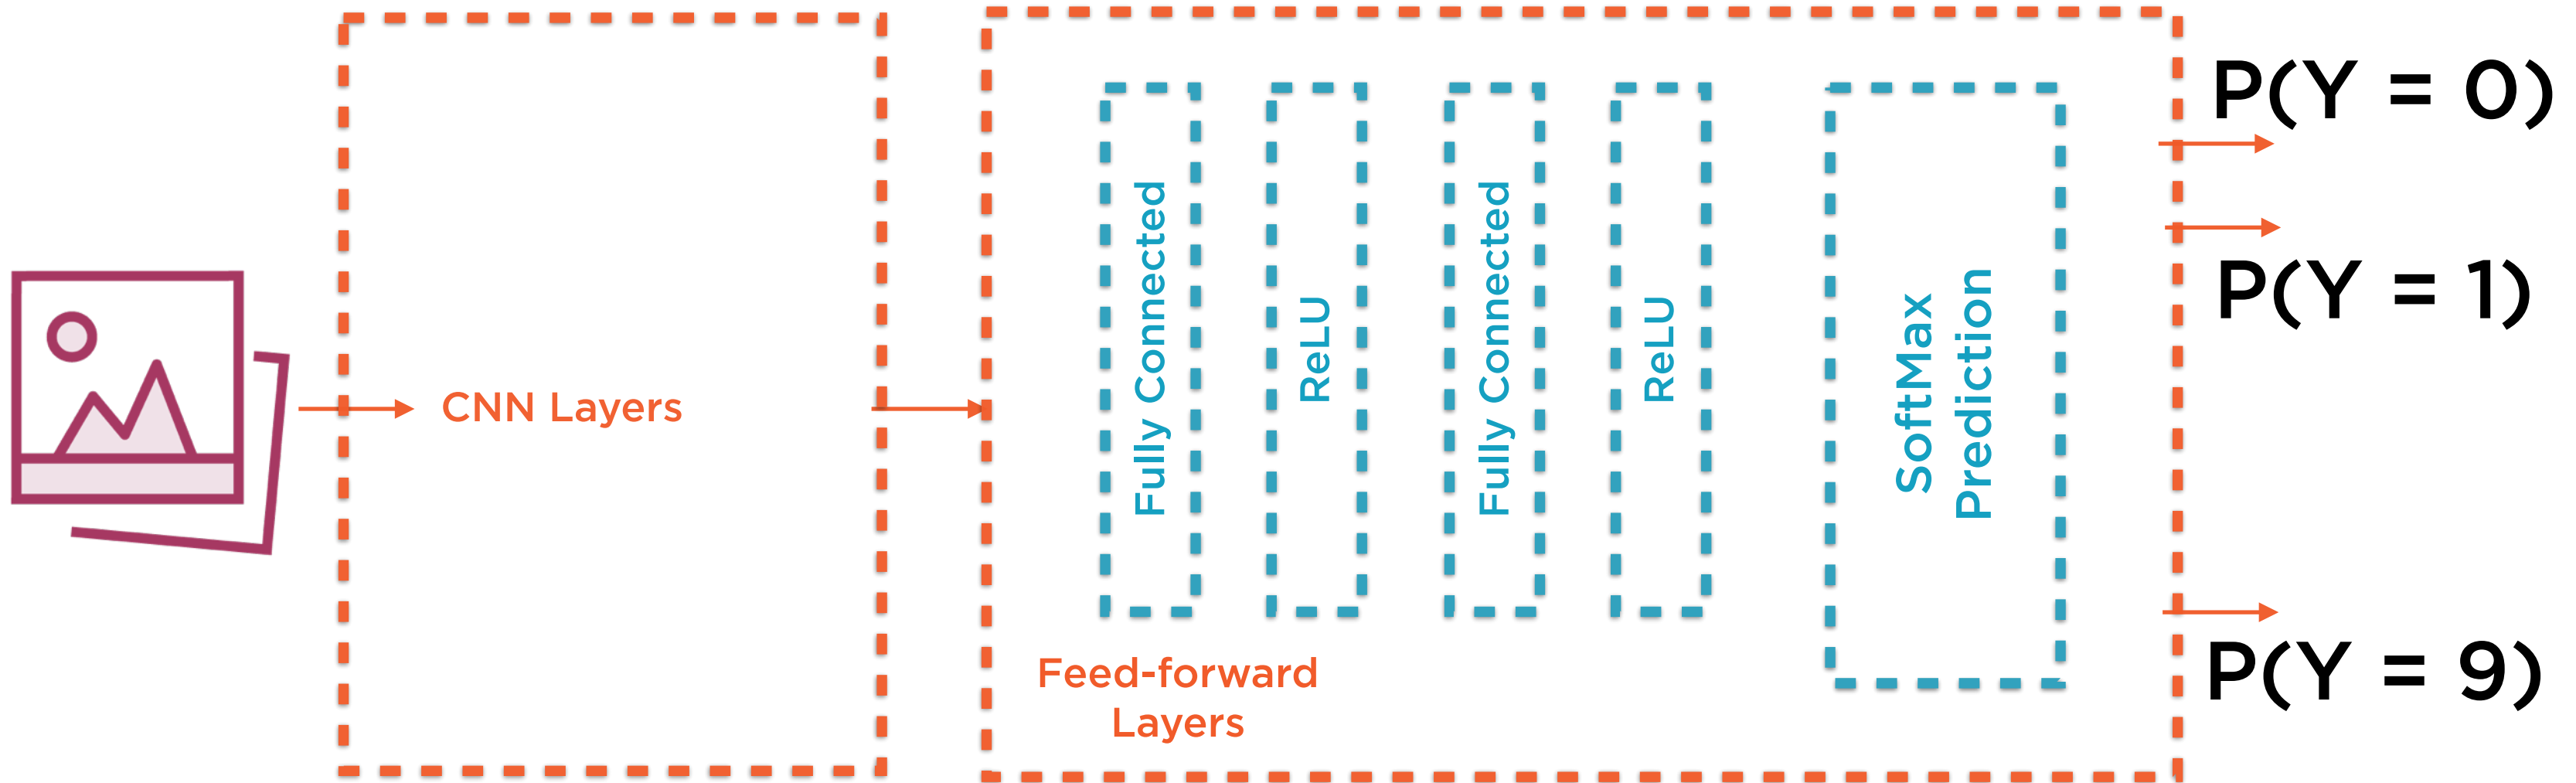
This feed-forward has a few fully connected layers with ReLU activation

Typical CNN Architecture




Finally a SoftMax prediction layer

Typical CNN Architecture



This is the output layer, emitting probabilities

Typical CNN Architectures



Input is an image

Outputs are probabilities

Demo

**Building a CNN to classify images from
the CIFAR-10 dataset**

Transfer Learning


Transfer Learning

The practice of re-using a trained neural network that solves a problem similar to yours, freezing the lower layers and only re-training the higher layers

Avoid designing NN
architecture from scratch

Transfer Learning


The practice of **re-using a trained neural network** that solves a problem similar to yours, freezing the lower layers and only re-training the higher layers



Also saves on time and effort
of re-training from scratch

Transfer Learning

The practice of **re-using a trained neural network** that solves a problem similar to yours, freezing the lower layers and only re-training the higher layers



Only makes sense for common,
widely studied use-cases...

Transfer Learning

The practice of re-using a trained neural network **that solves a problem similar to yours**, freezing the lower layers and only re-training the higher layers



...in which basic problem structure stays same, but details vary

Transfer Learning

The practice of re-using a trained neural network **that solves a problem similar to yours**, freezing the lower layers and only re-training the higher layers




Image recognition, language translation are classic examples

Transfer Learning

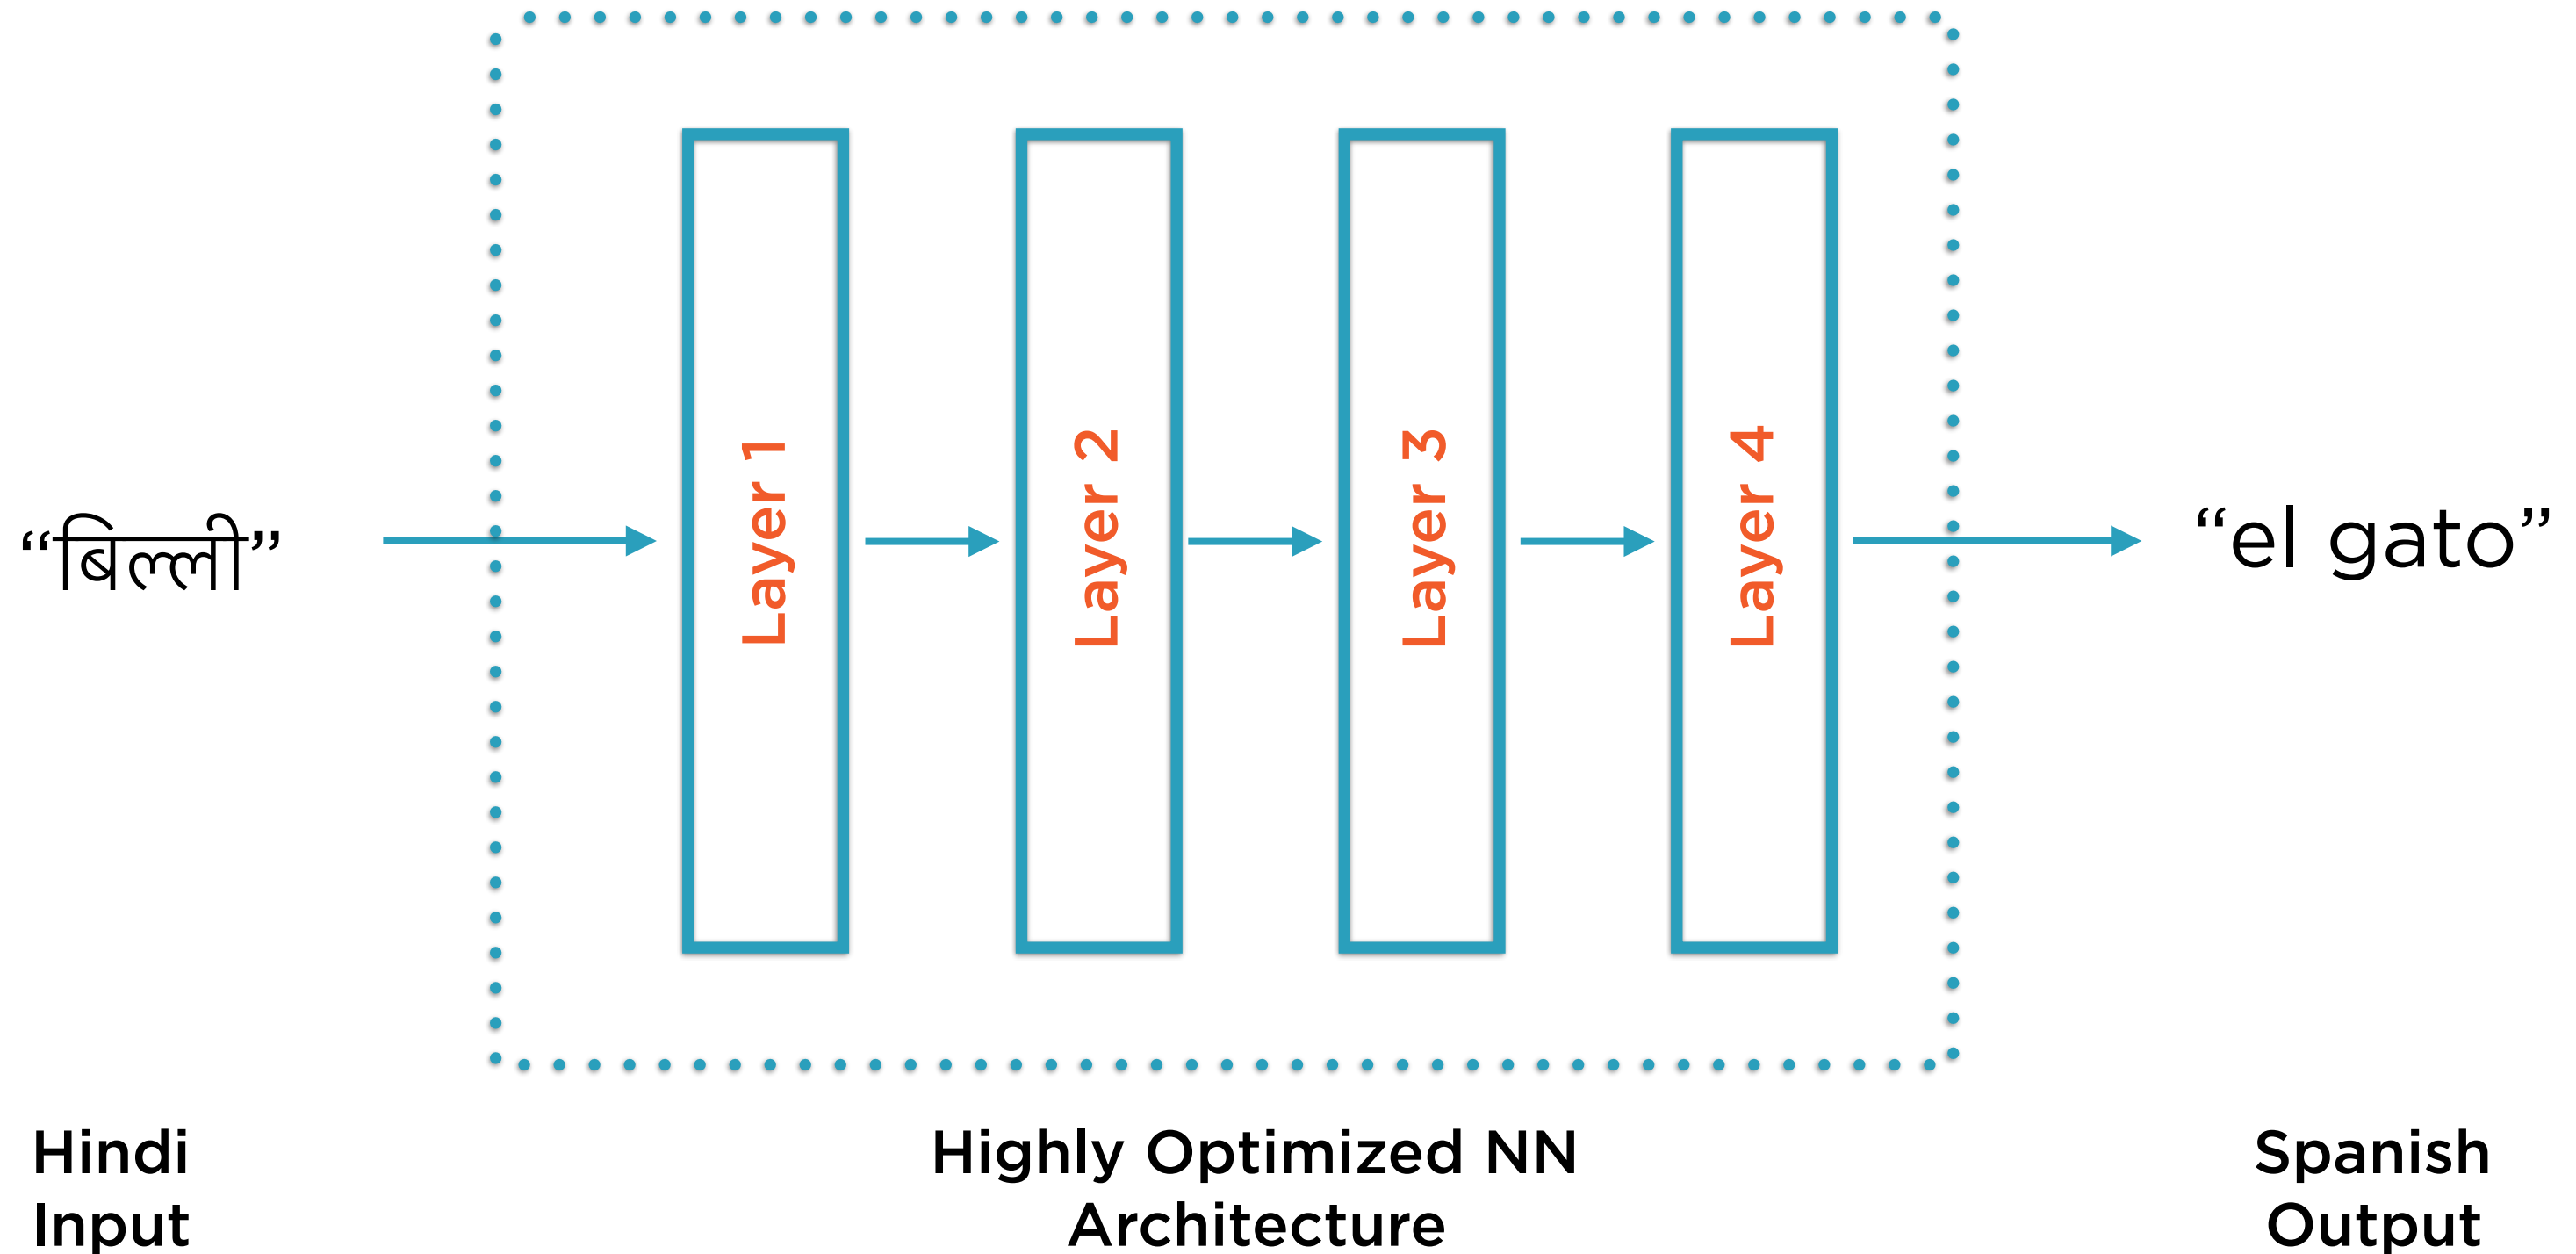
The practice of re-using a trained neural network **that solves a problem similar to yours**, freezing the lower layers and only re-training the higher layers



Original Model: English to French



Transfer Learning: Hindi to Spanish



Transfer Learning: Hindi to Spanish

Re-use
Architecture

“बिल्ली”

Layer 1

Layer 2

Layer 3


Layer 4

“el gato”

Hindi
Input

Highly Optimized NN
Architecture


Spanish
Output



Lower layers mostly
perform feature extraction

Transfer Learning


The practice of re-using a trained neural network that solves a problem similar to yours, **freezing the lower layers** and only re-training the higher layers




Re-use as-is without even
changing parameter weights

Transfer Learning


The practice of re-using a trained neural network that solves a problem similar to yours, **freezing the lower layers** and only re-training the higher layers



Original Model: English to French



Original Model: English to French



Transfer Learning: Hindi to Spanish

Re-use
lower layers
as-is

“बिल्ली”

Layer 1

Layer 2

Layer 3


Layer 4

“el gato”

Hindi
Input


Highly Optimized NN
Architecture

Spanish
Output




Transfer Learning: Hindi to Spanish

i.e. same
parameter
weights



Transfer Learning: Hindi to Spanish


i.e. “freeze”
lower layers



Re-use as-is without even
changing parameter weights

Transfer Learning


The practice of re-using a trained neural network that solves a problem similar to yours, **freezing the lower layers** and only re-training the higher layers



Transfer learning for “fixed
feature extraction”

Transfer Learning


The practice of re-using a trained neural network that solves a problem similar to yours, **freezing the lower layers** and only re-training the higher layers




Can't avoid this - higher layers are more “high-level”

Transfer Learning

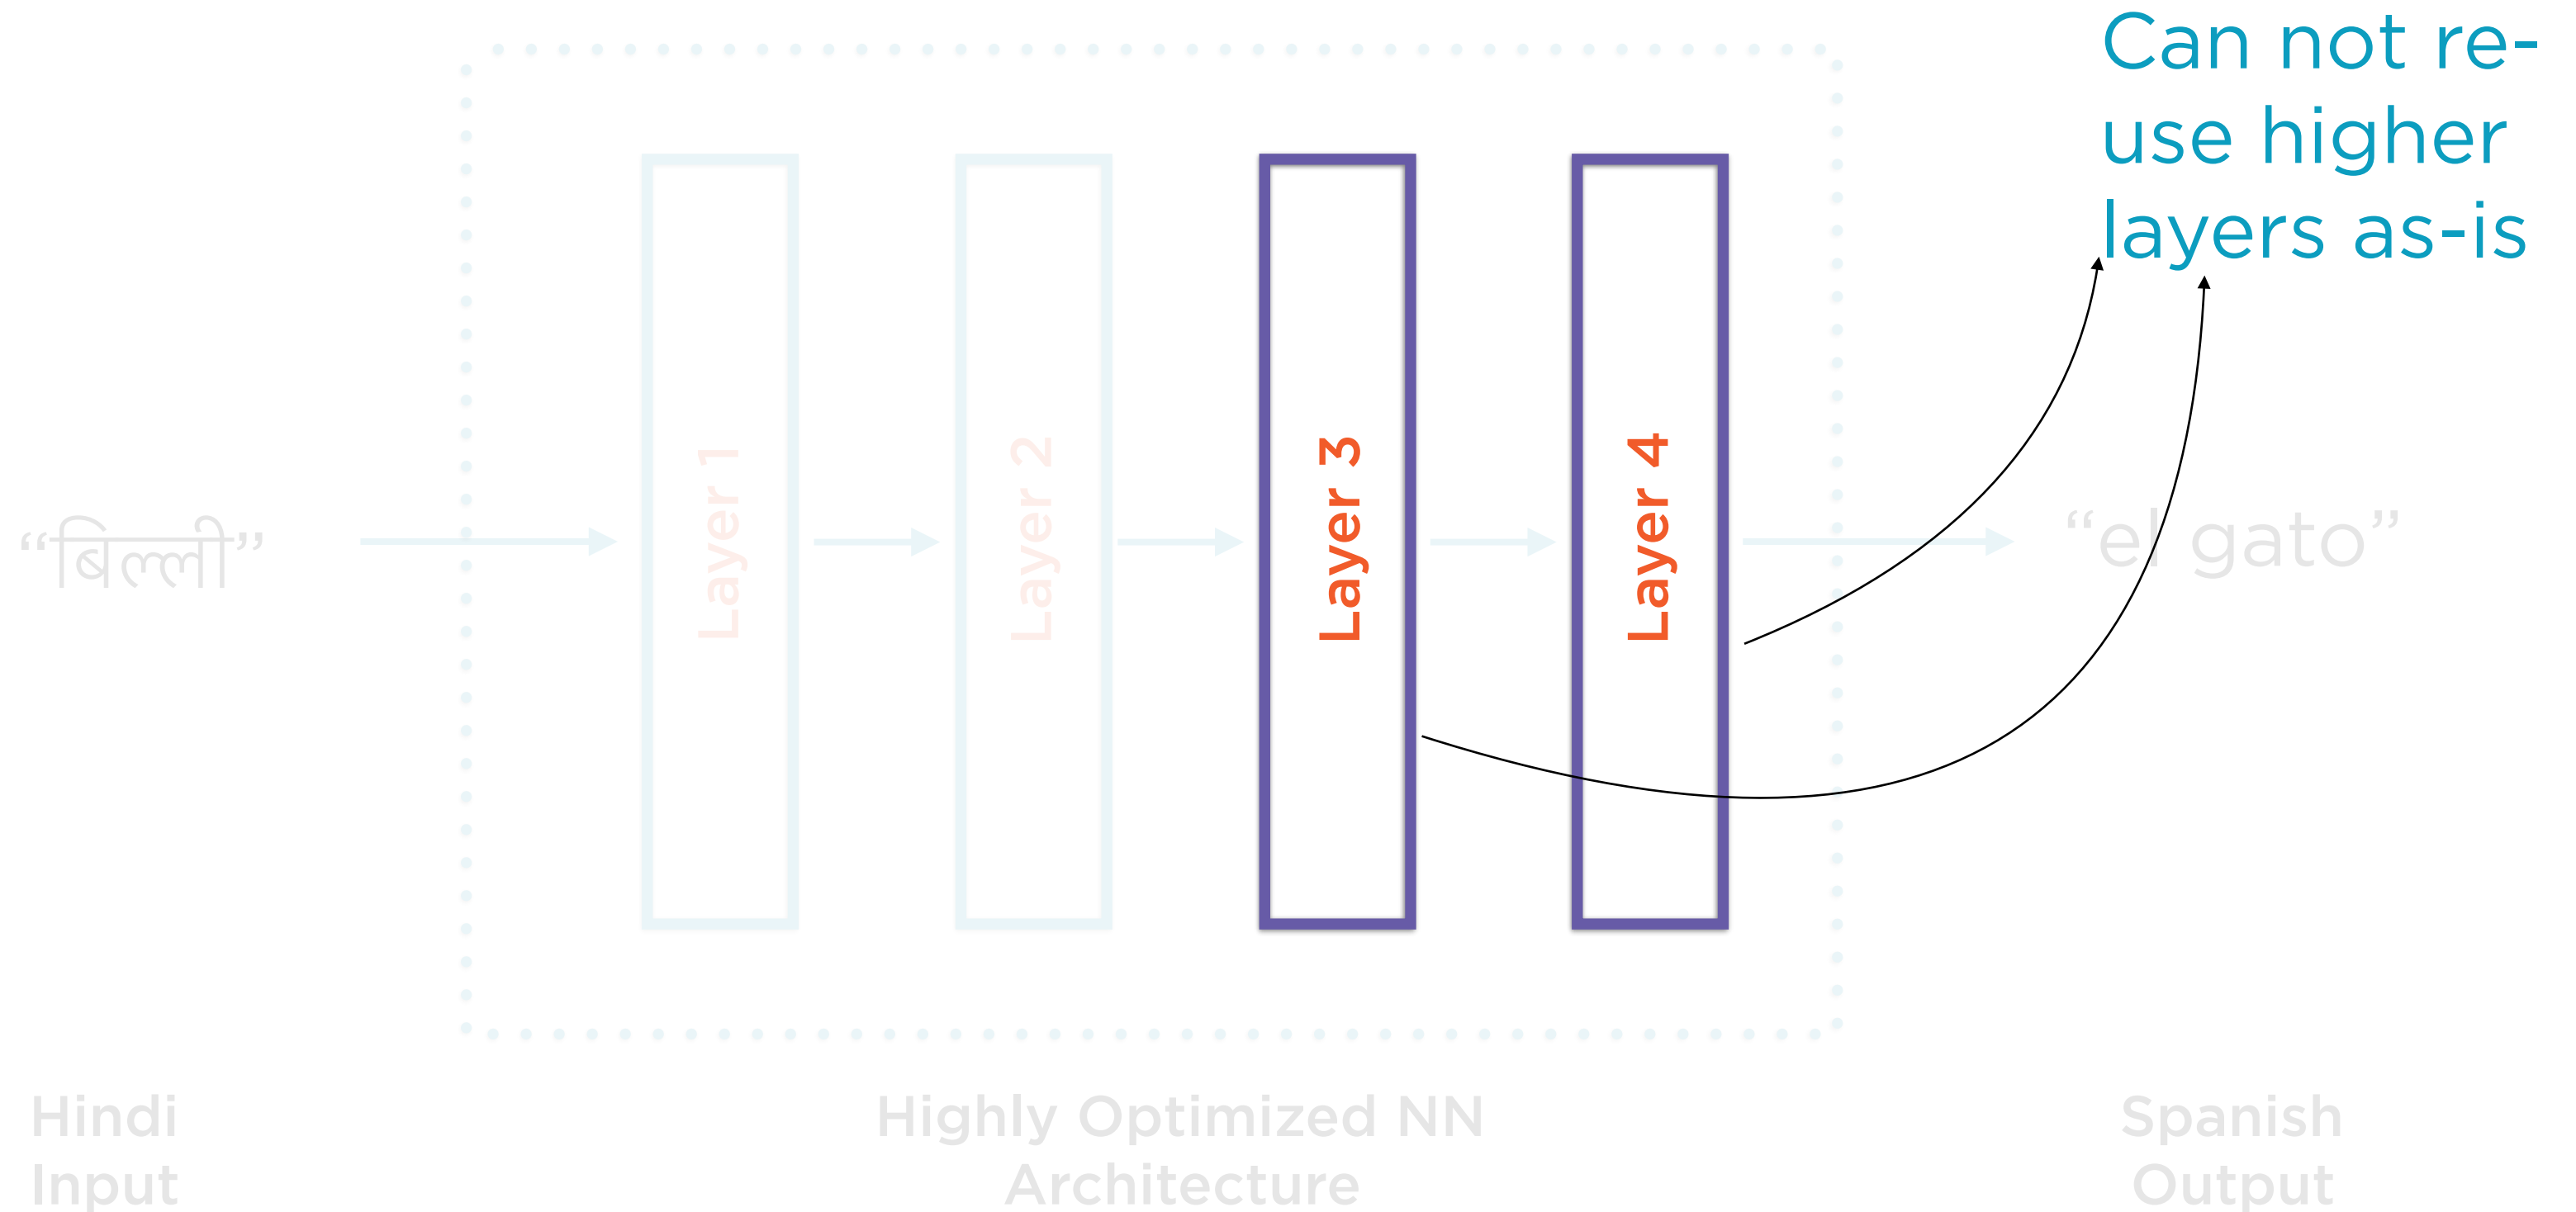
The practice of re-using a trained neural network that solves a problem similar to yours, freezing the lower layers and **only re-training the higher layers**



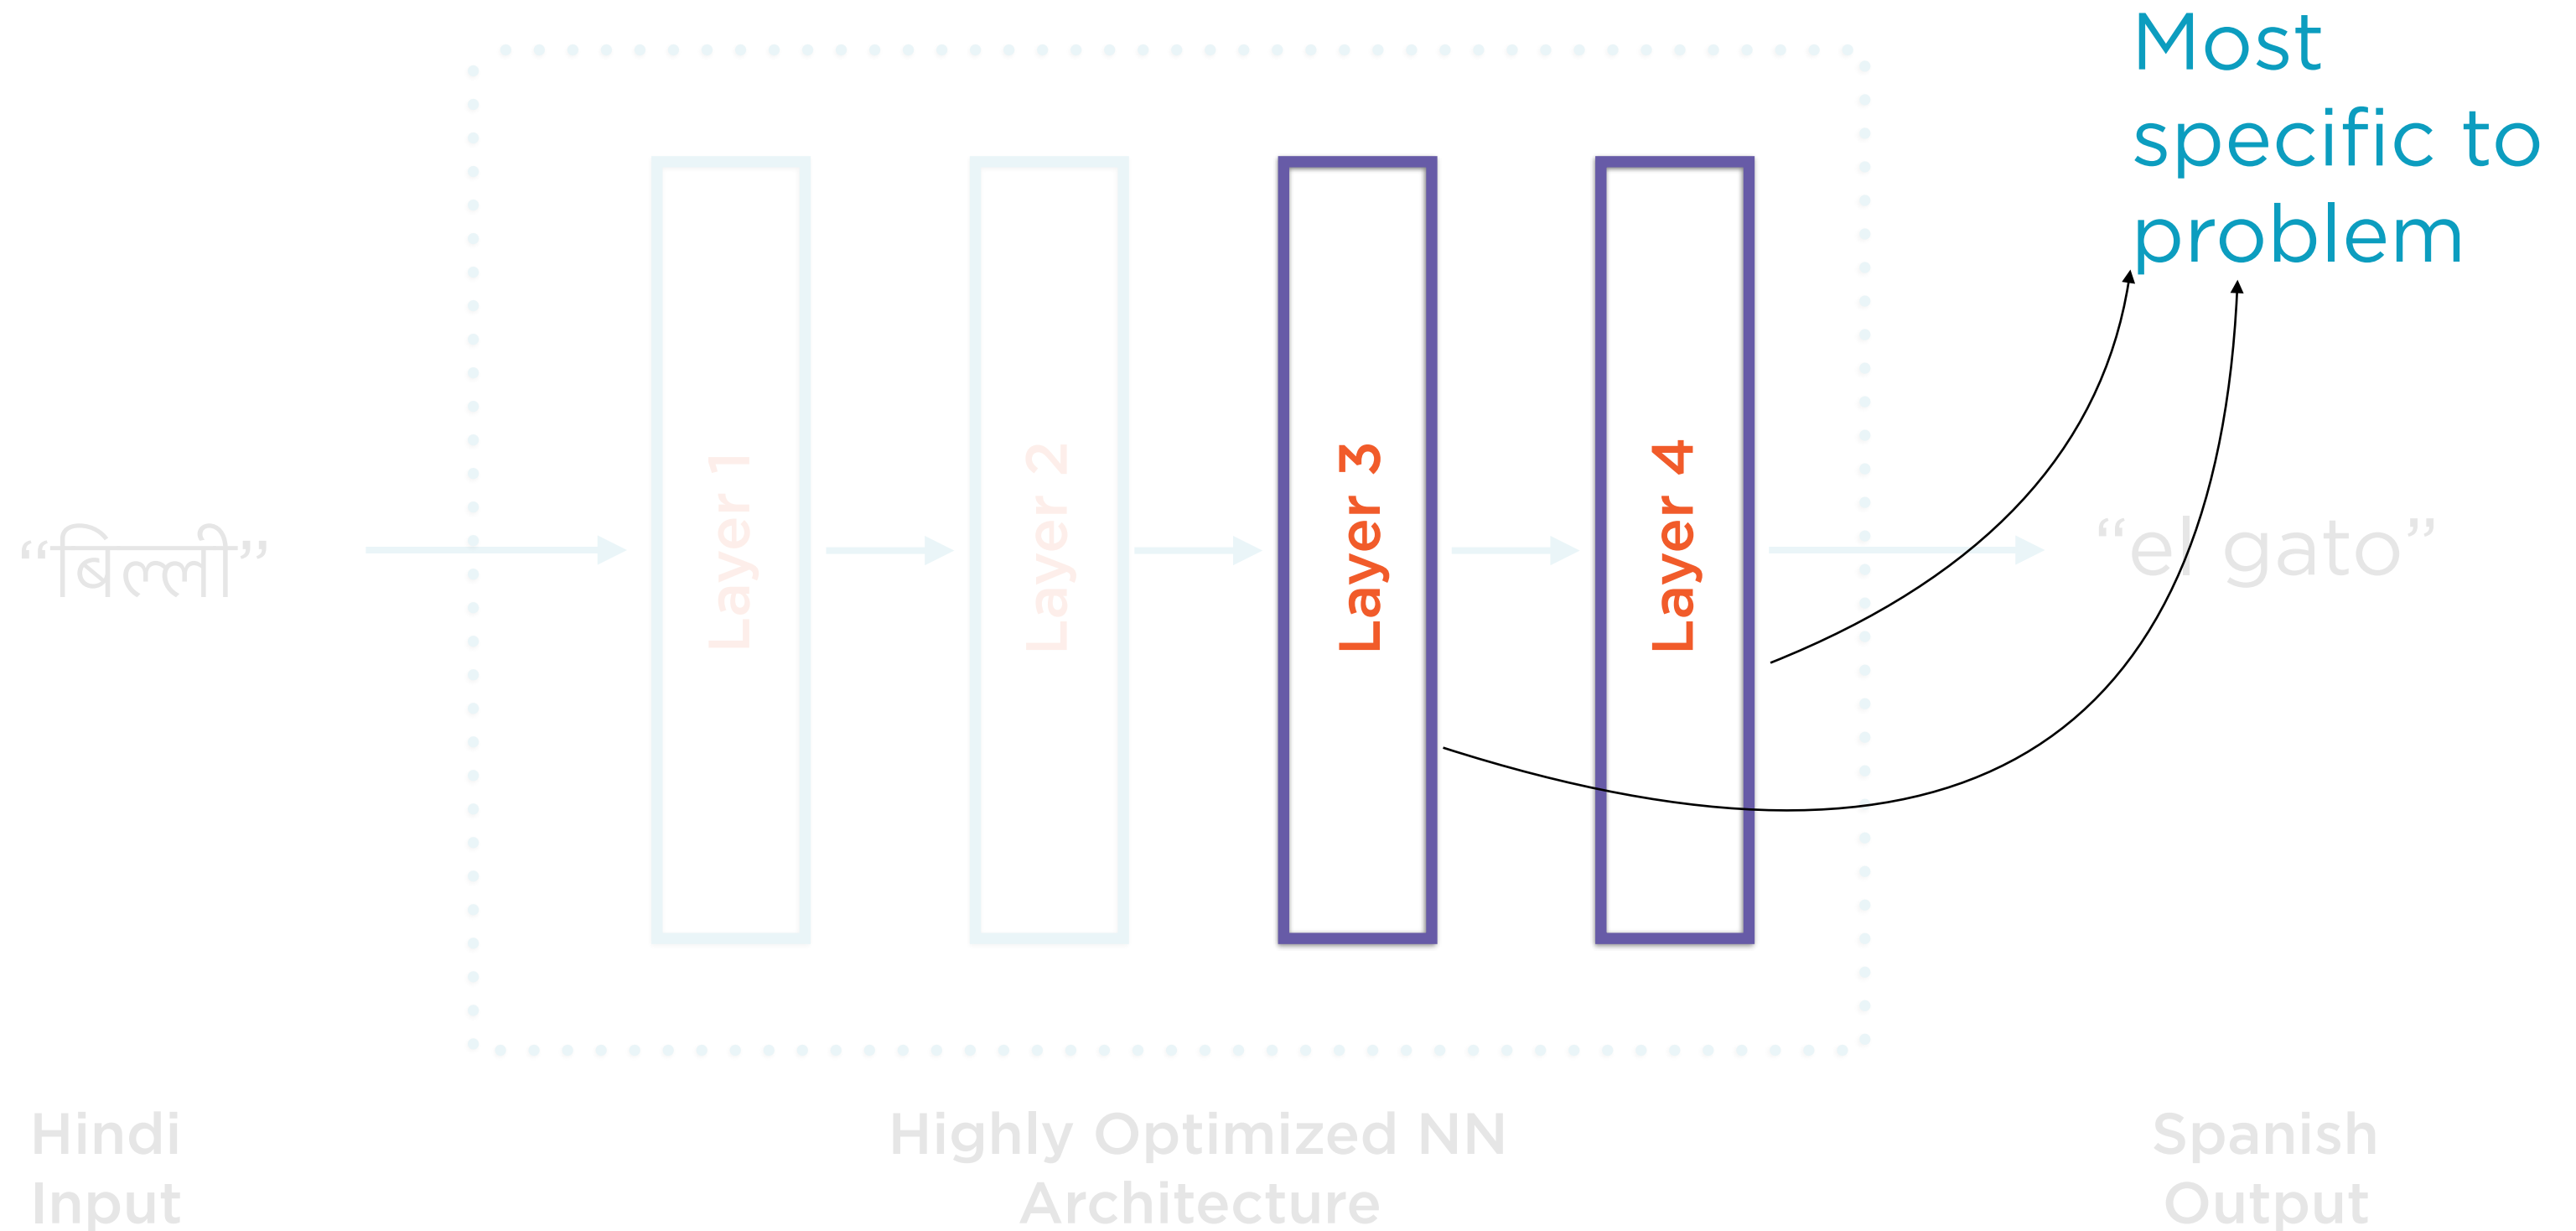
Original Model: English to French



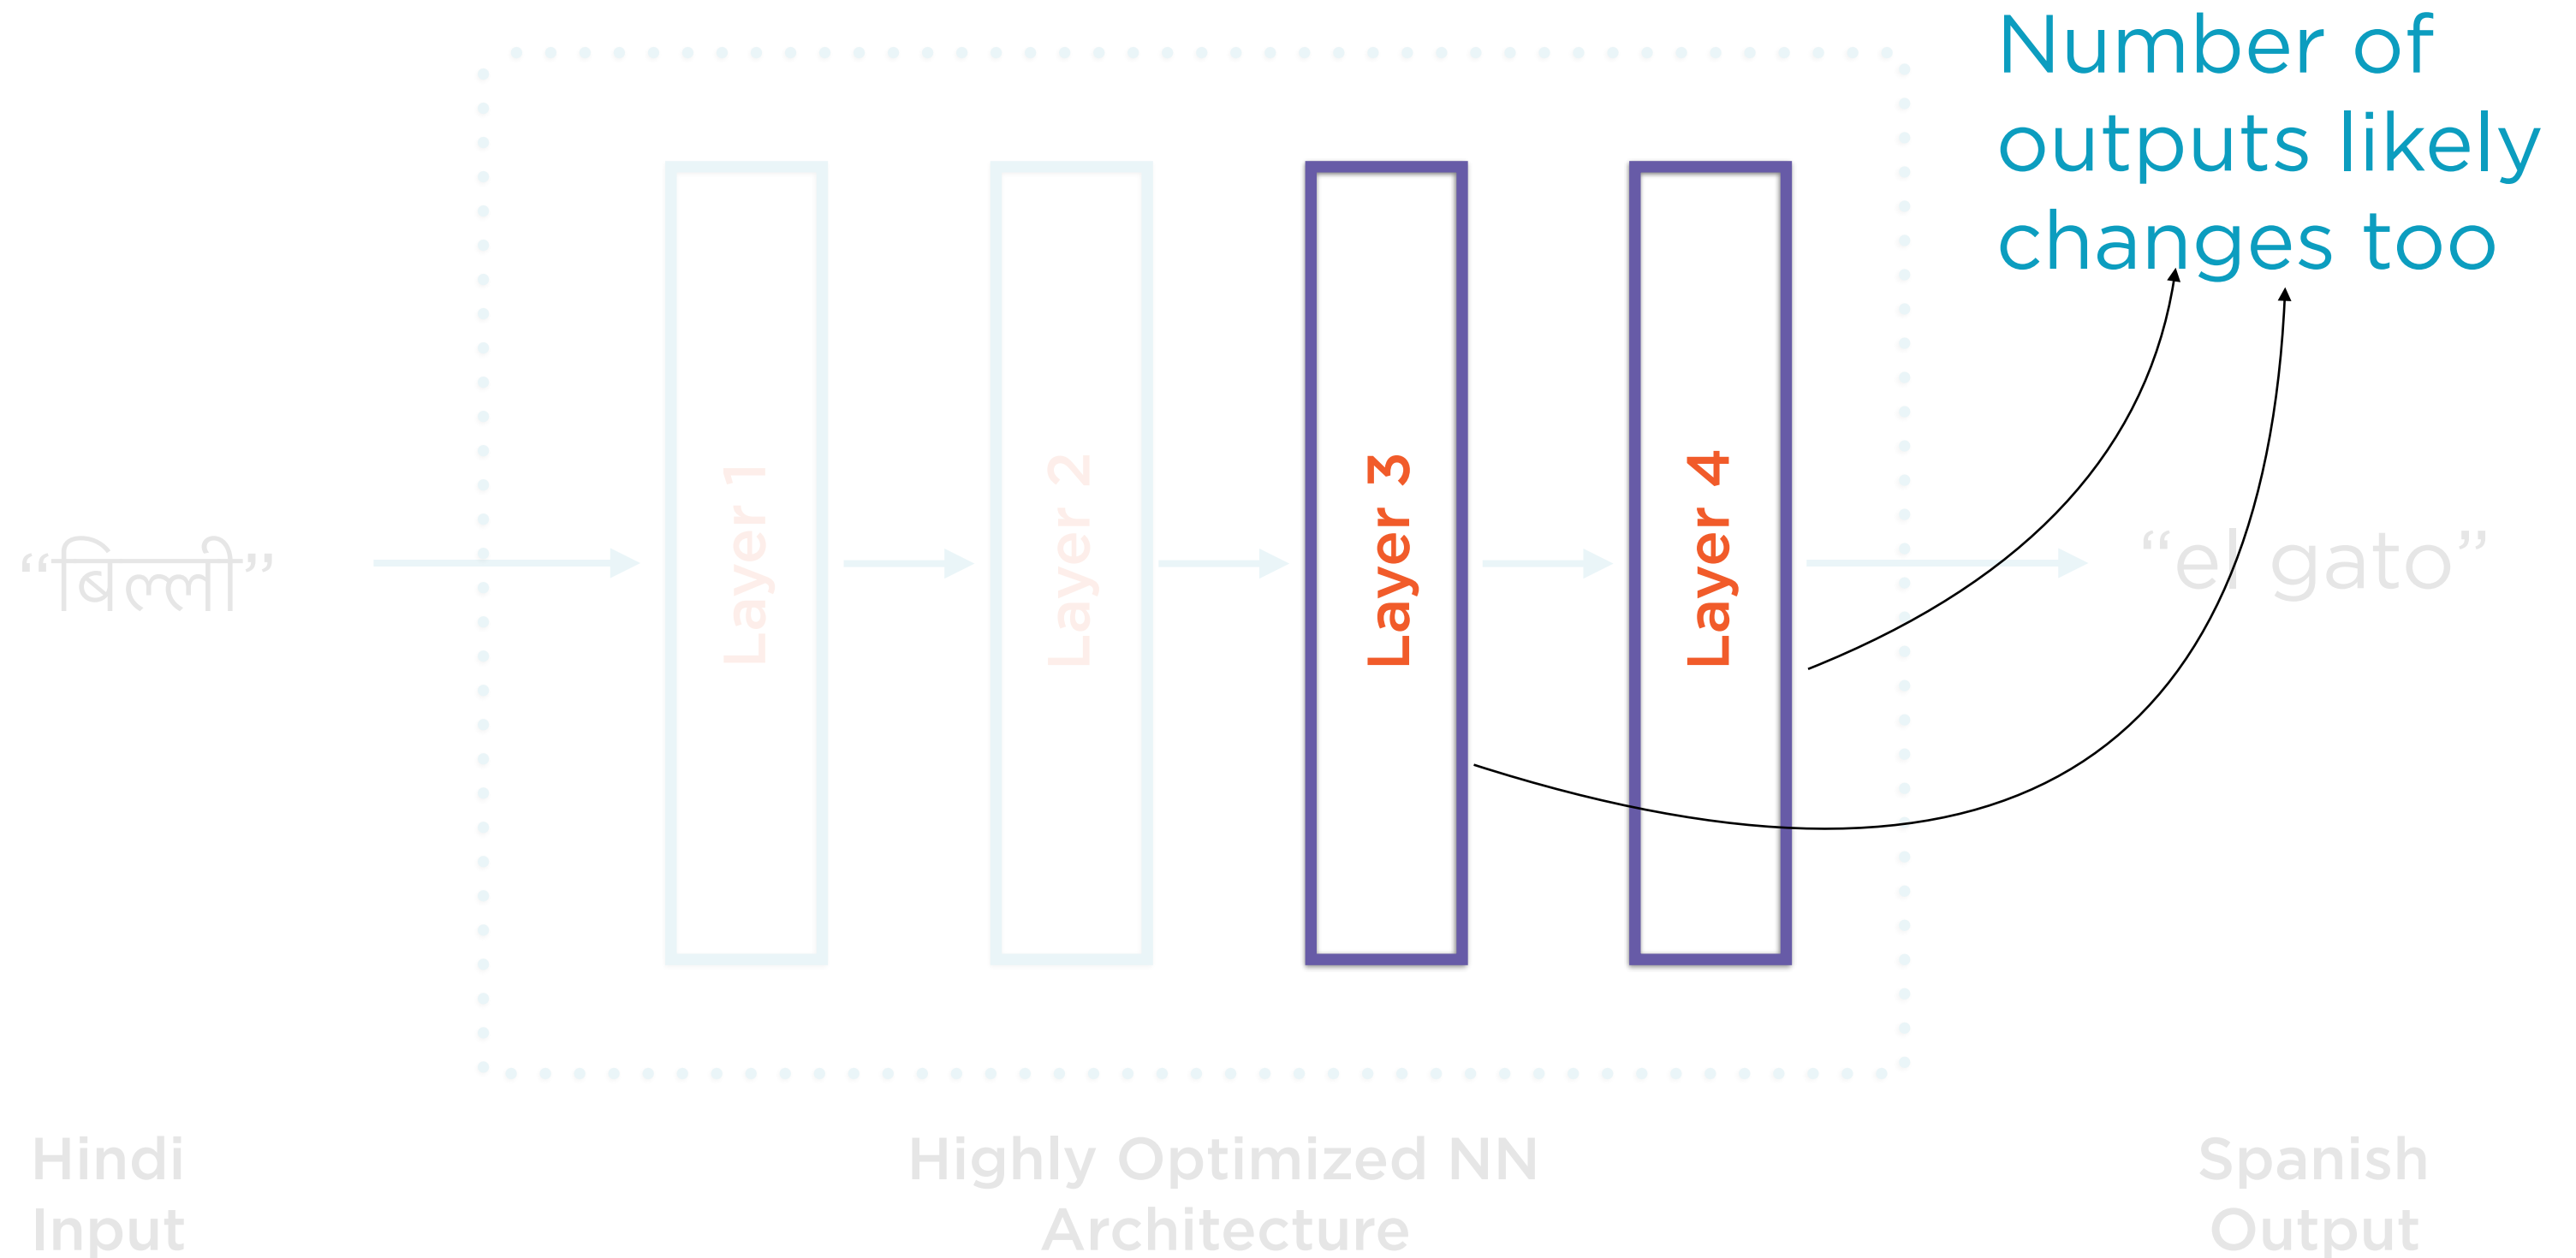
Original Model: English to French



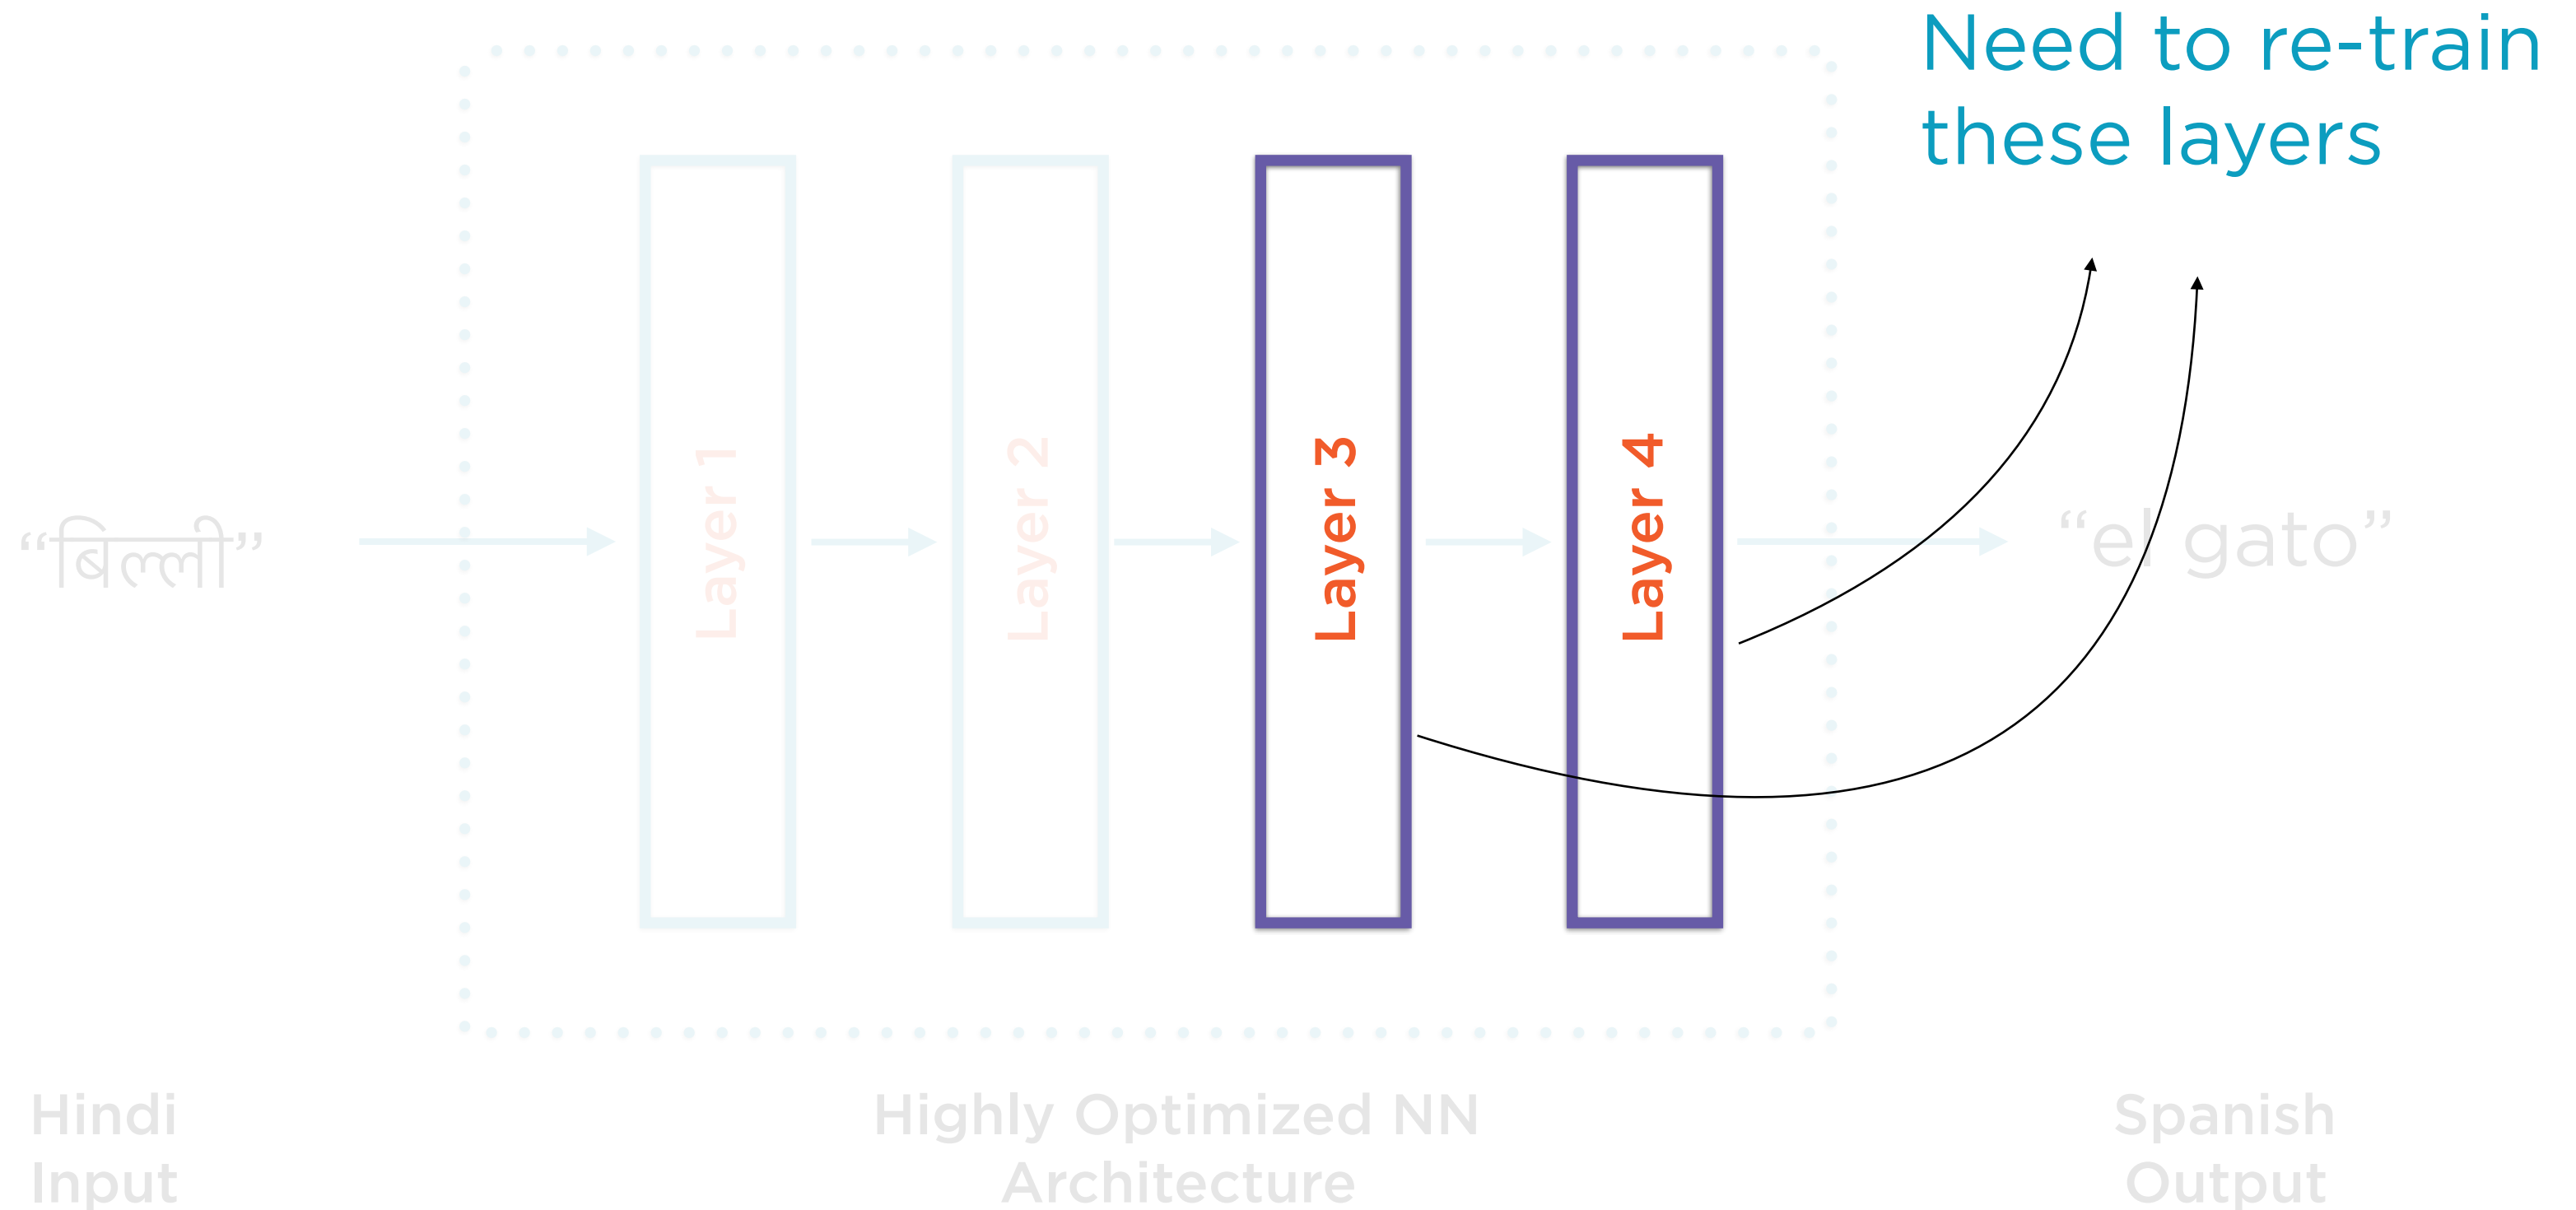
Transfer Learning: Hindi to Spanish



Transfer Learning: Hindi to Spanish



Transfer Learning: Hindi to Spanish



Transfer Learning

The practice of re-using a trained neural network that solves a problem similar to yours, freezing the lower layers and only re-training the higher layers

Benefits of Transfer Learning

“Ride on the shoulders of giants”

- NN architecture
- Choice of initialization
- Activation functions
- Number and density of layers

Benefits of Transfer Learning

“Do more with less”

Make do with less training data

- English to French: Lots of training data
- Hindi to Spanish: Little or no training data

Benefits of Transfer Learning

“Faster, cheaper”

Training process is far faster, easier

- Smaller training data
- Only higher layers to train
- In a cloud-enabled world, less time => less money

Transfer Learning in PyTorch

**Support for several famous NN
architectures**

torchvision.models

- Alexnet
- VGG
- ResNet
- Inception
- ...

Transfer Learning in PyTorch

Support for several famous NN
architectures

`torchvision.models`

- Alexnet
- VGG
- ResNet
- Inception
- ...

Demo

**Use the ResNet-18 pre-trained model to
classify flowers**

ResNet

Famous CNN architecture


Won famous challenge in 2015

Extremely deep

“Skip connections” aka shortcut connections

Shares many features with typical CNN architectures

ResNet



Big innovation - “skip connections”

Connect output of lower layers to far-ahead higher layers

Model is forced to focus on what is not learnt by intermediate layers

“Residual Learning”

Summary

Convolutional NNs are a deep learning technique easily implemented in PyTorch

ResNet is a famous CNN architecture

Transfer learning is a great way to re-use pre-trained models

PyTorch offers great support for transfer learning

Provides a ResNet model

Image classification using transfer learning and ResNet