

Building a Text Classification Model



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Recurrent Neural Networks (RNNs) are another common NN architecture

Great for working with sequences

Classification using RNNs

PyTorch has built-in support for RNNs

Recurrent Neural Networks (RNNs)

$$y = f(x)$$

Machine Learning

Machine learning algorithms seek to “learn” the function f that links the features and the labels

$$y_t = f(x_t, y_{t-1})$$

Learning the Past

Relationships where past values of the effect variable drive current values are called auto-regressive

$$y_t = f(x_t, y_{t-1})$$


Learning the Past

Relationships where past values of the effect variable drive current values are called auto-regressive


Feed-forward networks cannot
learn from the **past**

Recurrent neural networks can


Simplest Feed-forward Neuron



Simplest Recurrent Neuron



Recurrent Neuron




y_t = Output at time t


Depends upon

- y_{t-1} = Output at time $t - 1$
- x_t = New inputs available only at time t


Unrolling Through Time




Unrolling Through Time




Unrolling Through Time




Unrolling Through Time




Unrolling Through Time



Unrolling Through Time



Recurrent Neuron



Regular neuron: input is feature vector, output is scalar


$$Y = \text{activation}(Wx + b)$$

Recurrent neuron: **output is vector too**

Input: $[X_0, X_1, \dots, X_t]$

Output: $[Y_0, Y_1, \dots, Y_t]$

Memory and State




Recurrent neurons remember the past

They possess 'memory'

The stored state could be **more complex than simply y_{t-1}**

The internal state is represented by **h_t**


Recurrent Neuron



Now, each neuron has two weight vectors

W_x, W_y


Recurrent Neuron



Now, each neuron has two weight vectors

W_x , W_y


Recurrent Neuron



Now, each neuron has two weight vectors

W_x , W_y

Recurrent Neuron




Output of neuron as a whole is given as

$$y_t = \Phi(X_t W_x + y_{t-1} W_y + b)$$

(Φ is the activation function)


Training a Recurrent Neural Network

Layer of Recurrent Neurons



A layer of neurons forms an RNN cell - basic cell, LSTM cell, GRU cell (more on these later)

Layer of Recurrent Neurons




The cells unrolled through time form the layers of the
neural network


The actual training of a neural network happens via Gradient Descent Optimization

Gradient Descent

Converging on the “best”
value using an
optimization algorithm



Back Propagation Through Time



$$y_t = f(x_t, y_{t-1}, y_{t-2} \dots, y_{t-1000})$$

Learning the Distant Past


The unrolled RNN will be very, very deep - many layers to train, the gradient has to be propagated a long way

Recurrent neural networks may be
unrolled **very far back in time**

They're prone to the **vanishing** and
exploding gradients issue


Deep neural networks are prone to
overfitting

Overfitting



Low Training Error

Model does very well in training...



High Test Error

...but poorly with real data

Preventing Overfitting



Regularization - Penalize complex models




Cross-validation - Distinct training and validation phases




Dropout - Intentionally turn off some neurons during training

Long Memory RNN Cells


Simplest Recurrent Neuron




Simplest Recurrent Neuron



Long Memory Recurrent Neuron



Long Memory RNNs




Increase the amount of state in neuron

Effect is to increase memory of neuron

Could explicitly add:

- long-term state (c)
- short-term state (h)

Long Memory RNNs



Advantages in Training

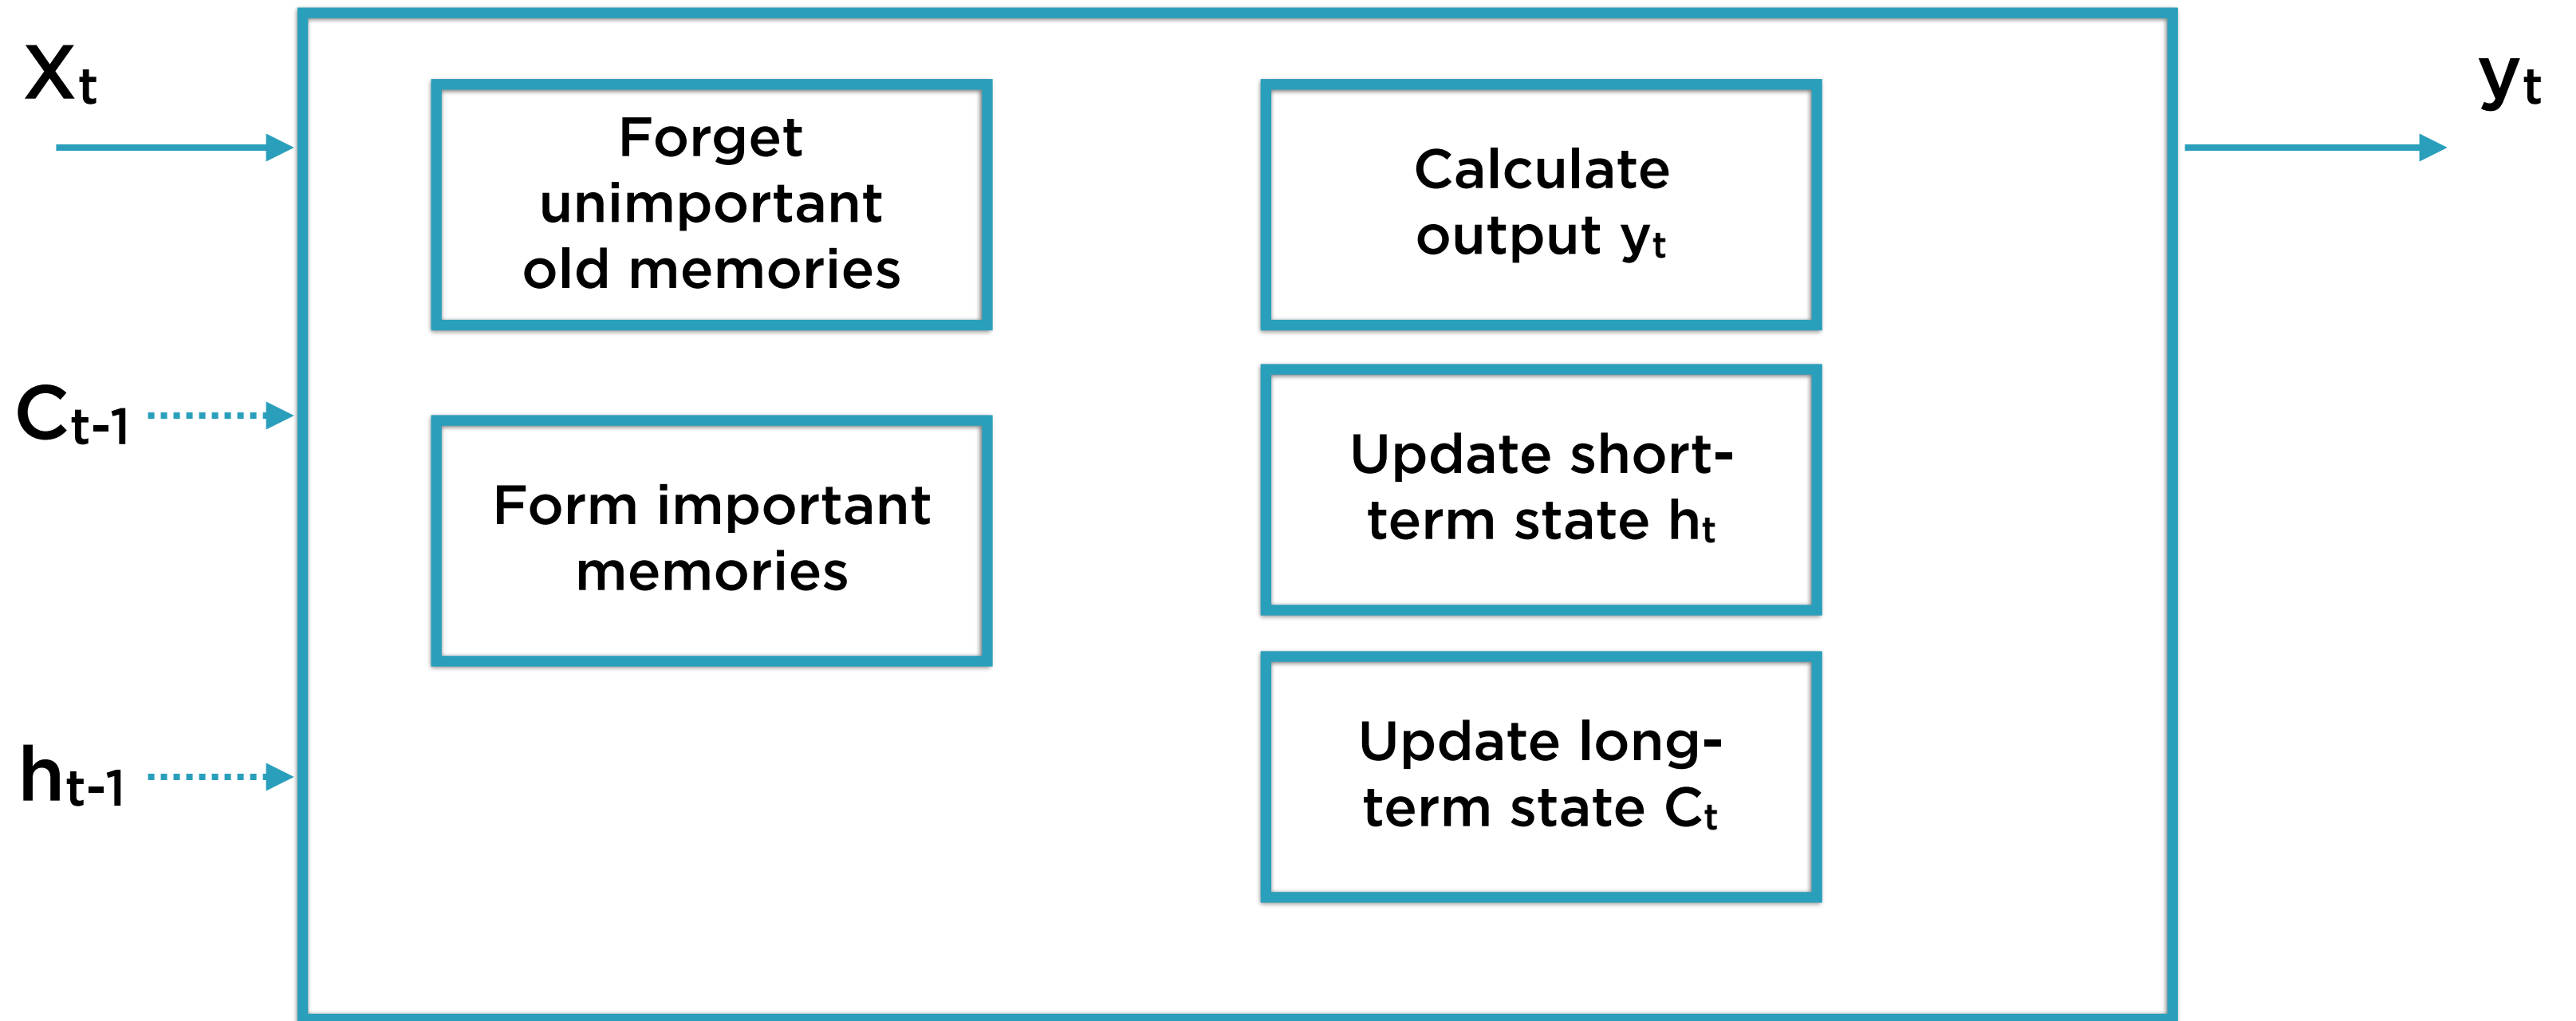
Faster training, nicer gradients




Advantages in Prediction

Works well with even very deep NNs

Long/Short-Term Memory Cell (LSTM)



LSTM Cells



Functionally like basic RNN cell

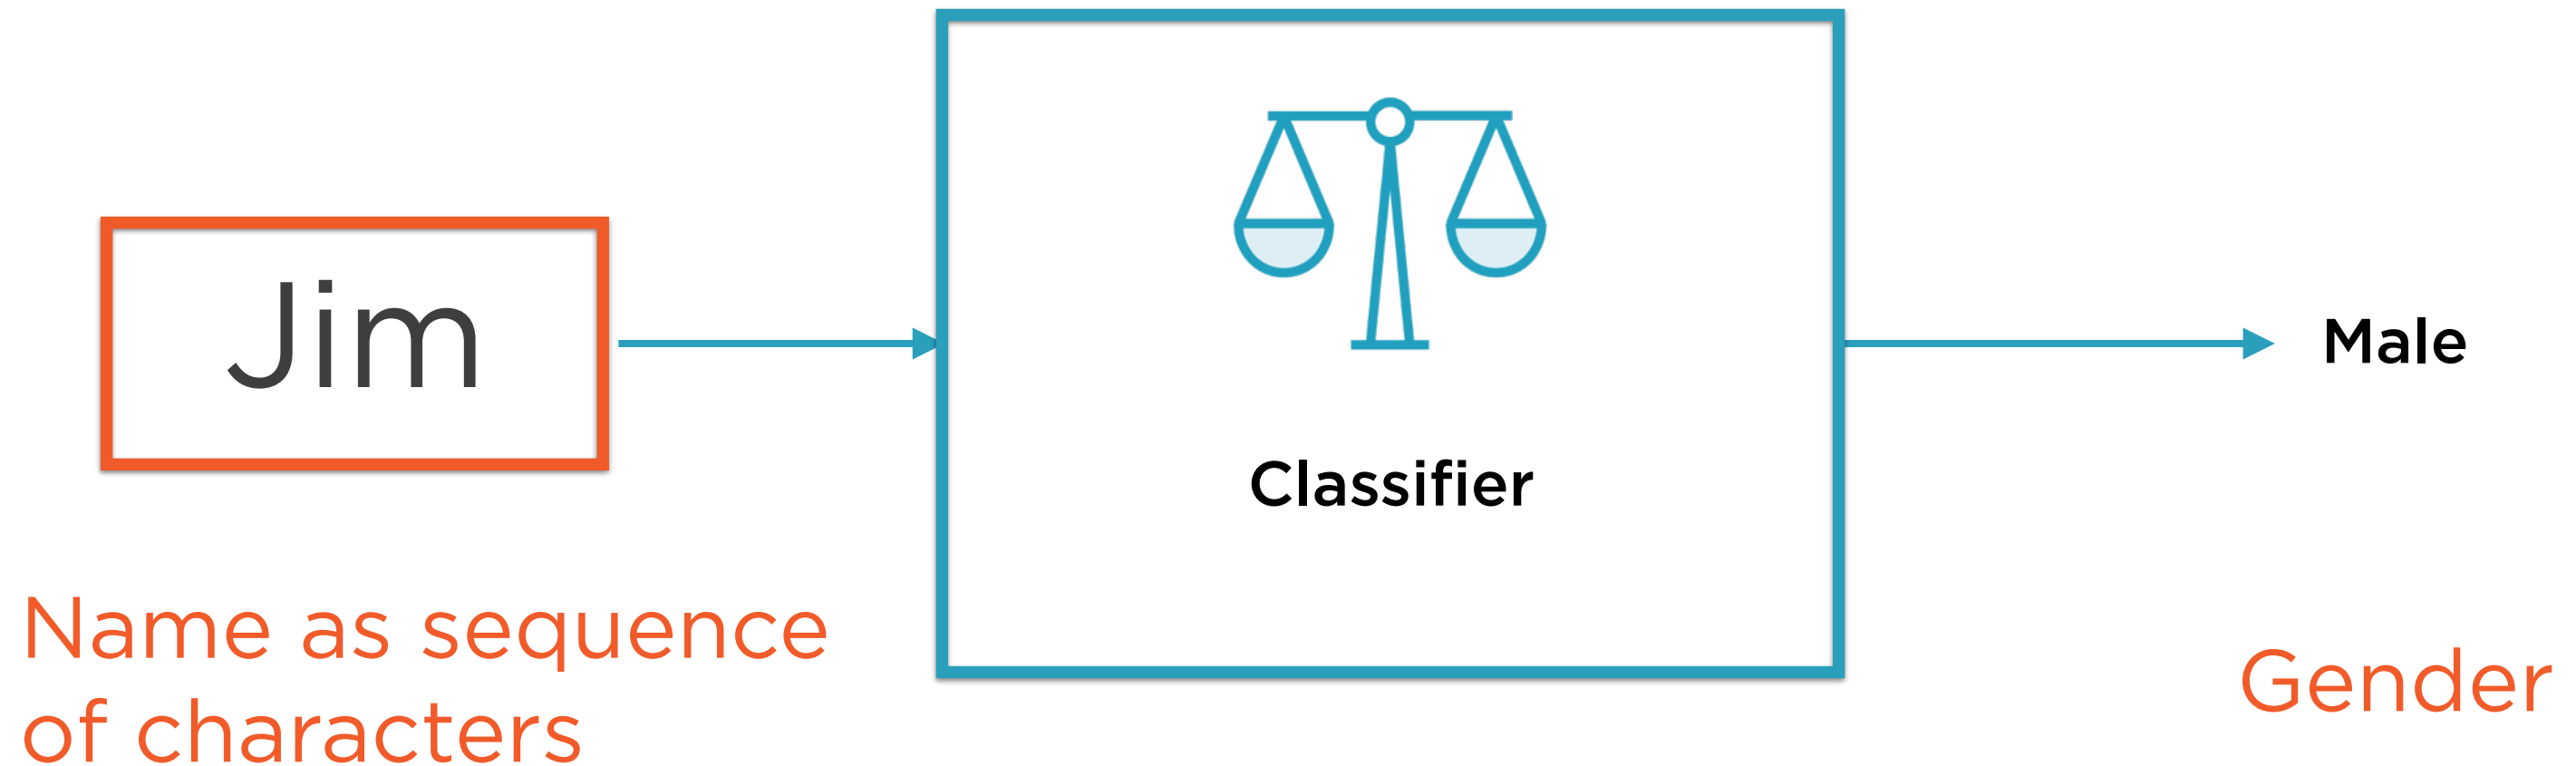
Performance far better

Amazing success at long-term patterns

Long text sequences, time series...

Gender Prediction Based on Names


Gender Prediction Based on Names



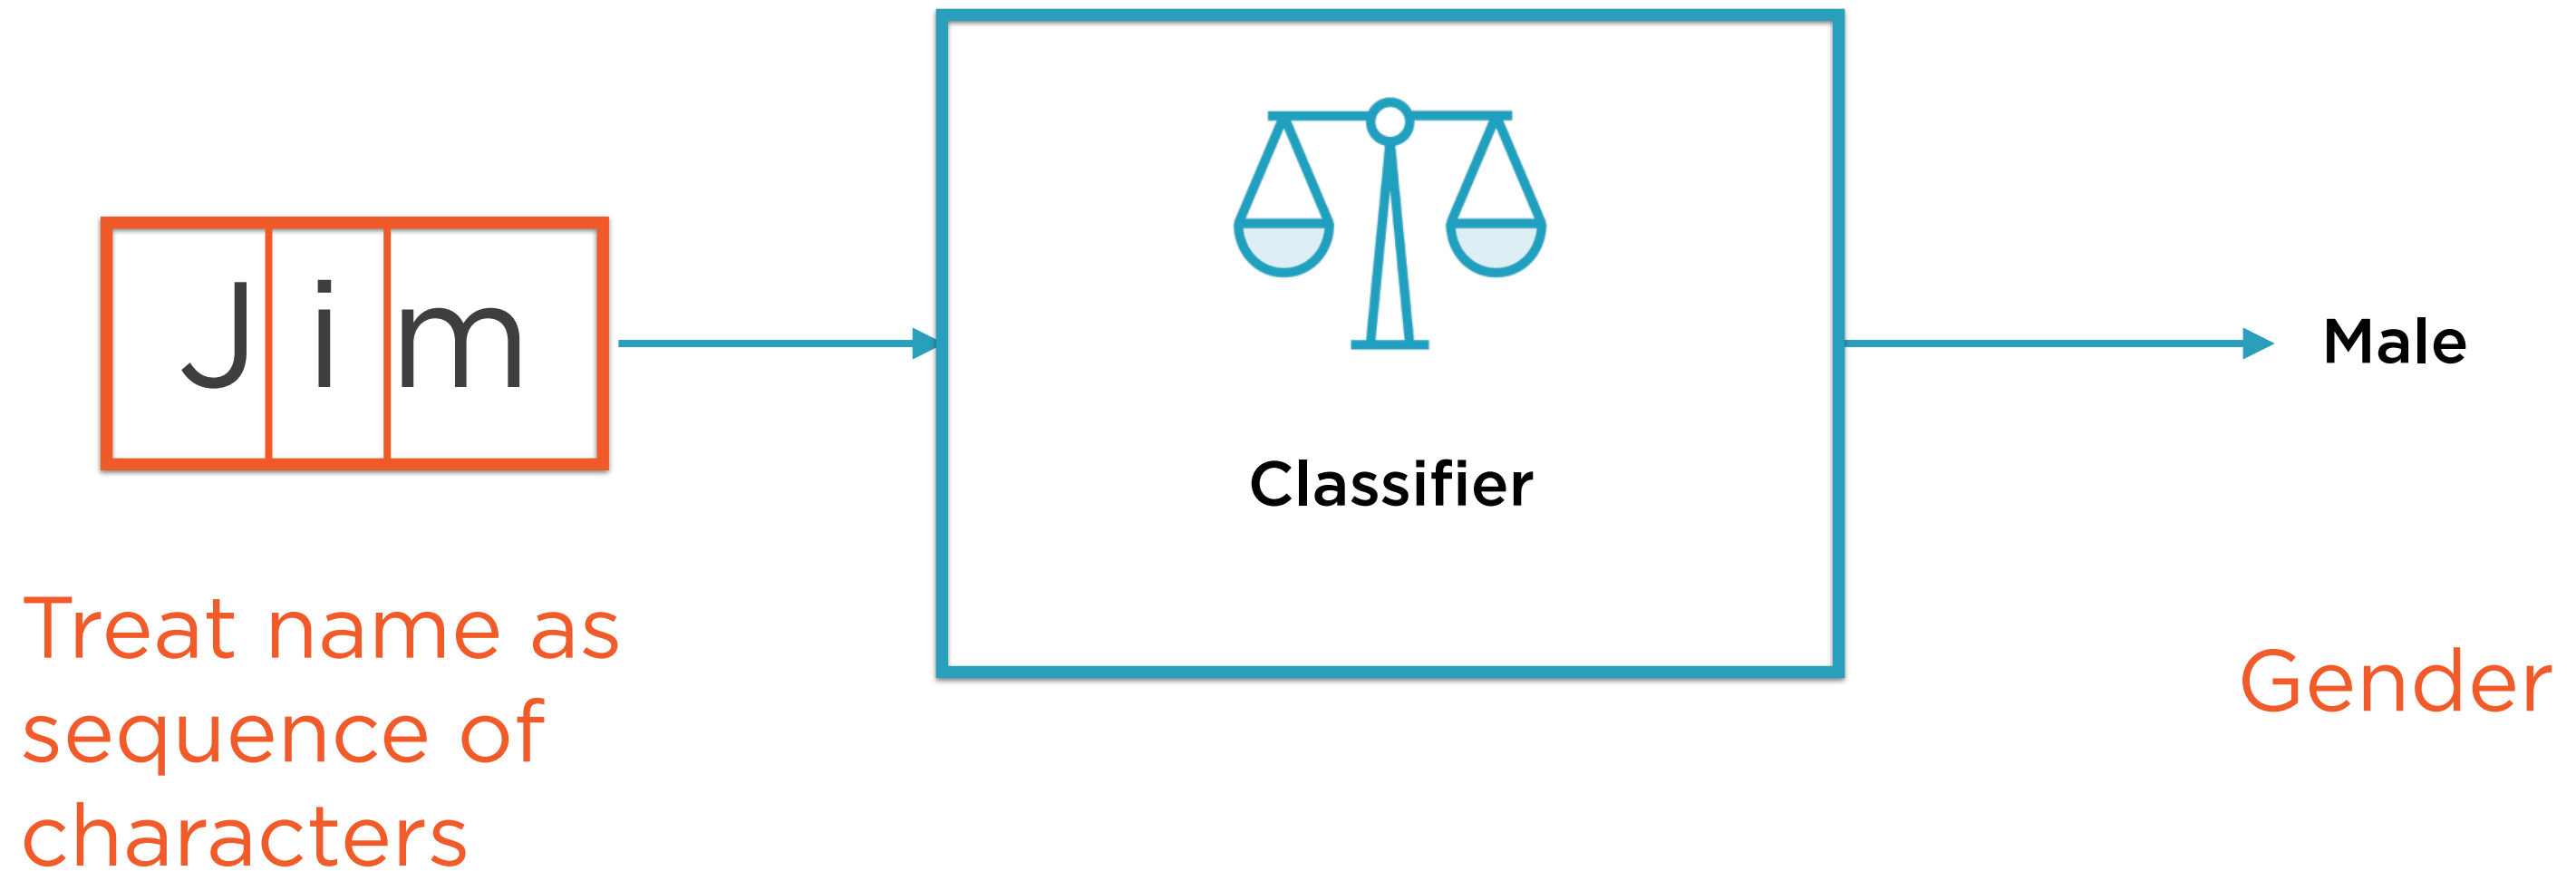
In PyTorch, because of dynamic computation graphs we do not need to pad names to be of the same length

This is a significant improvement over TensorFlow

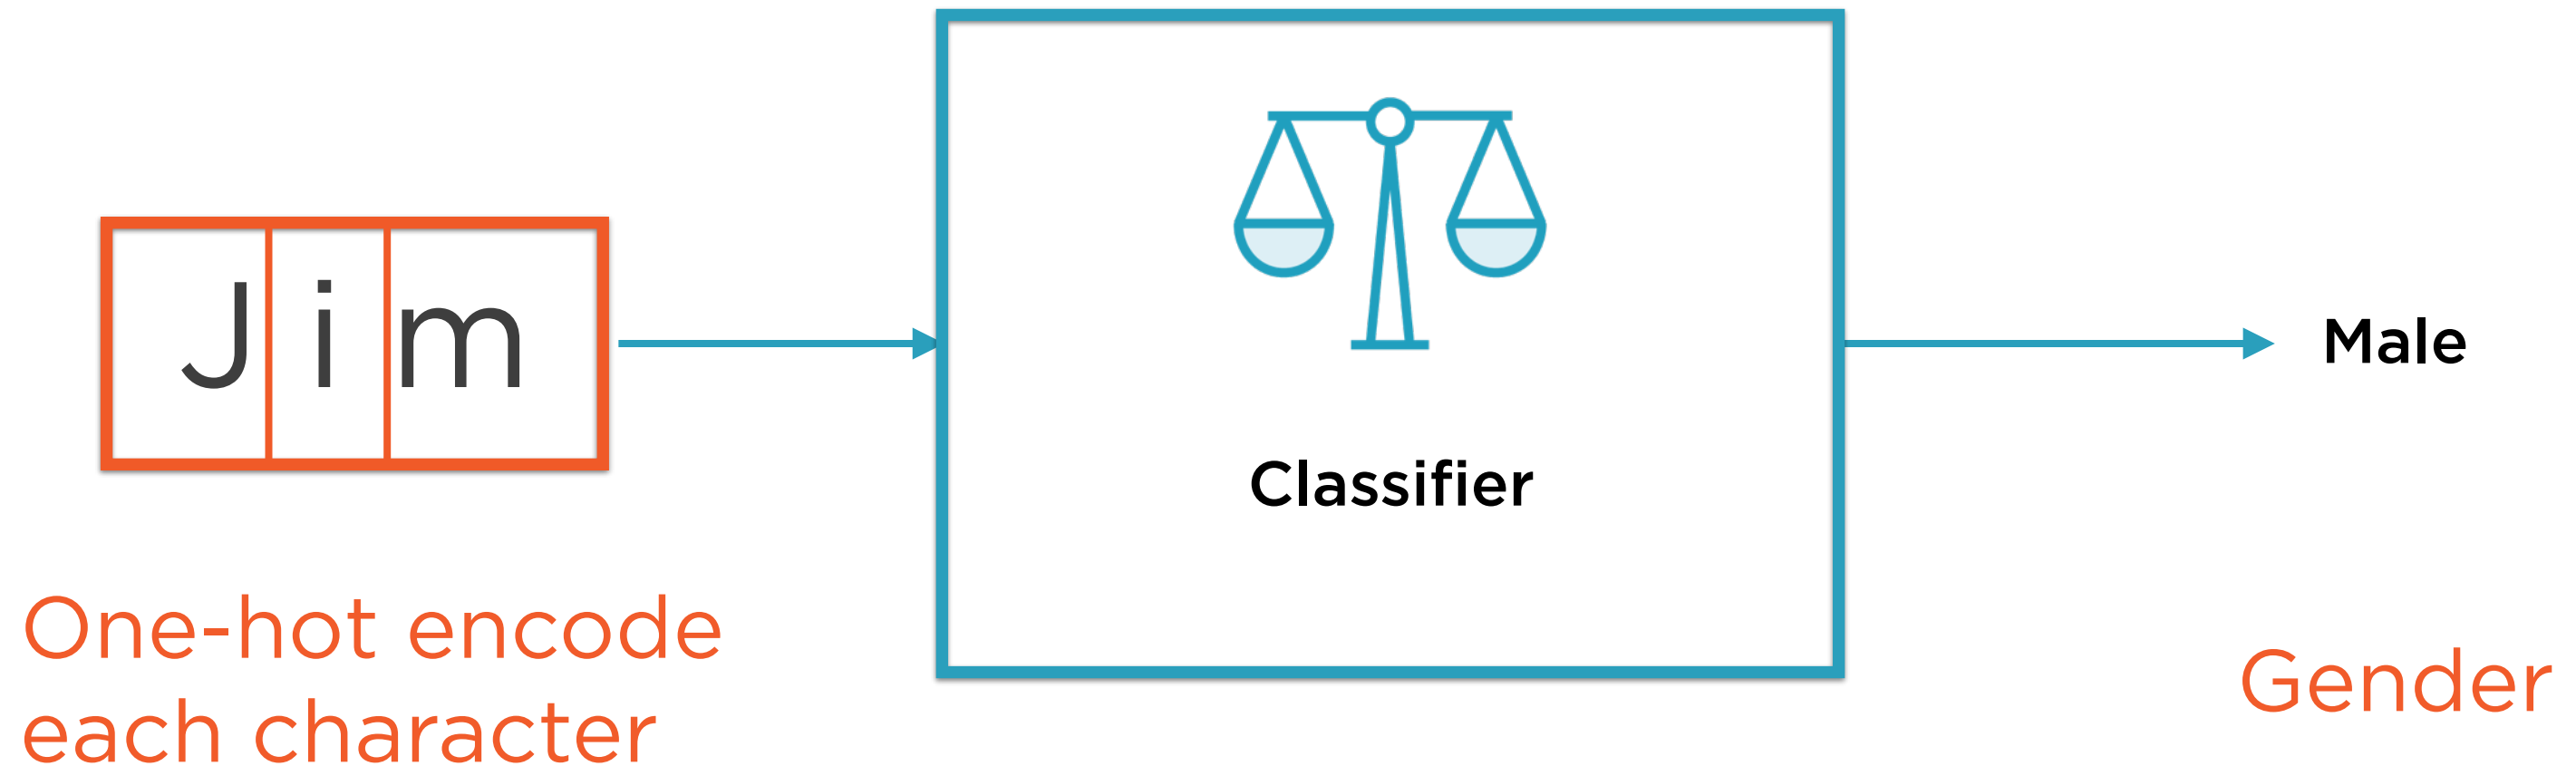
Gender Prediction Based on Names



Gender Prediction Based on Names



Gender Prediction Based on Names



One-hot Encoded Characters

Letter	a	...	j	...	i	...	m	...
J	0	0	1	0	0	0	0	0
i	0	0	0	0	1	0	0	0
m	0	0	0	0	0	0	1	0



26 elements

One-hot Encoded Characters

Letter	a ...	j	... i ... m ...
J	0	1	0 0 0 0 0 0
i	0	0	0 1 0 0 0 0
m	0	0	0 0 0 0 1 0

← 26 elements →

One-hot Encoded Characters

Letter	a ... j ...				i	... m ...		
j	0	0	1	0	0	0	0	0
i	0	0	0	0	1	0	0	0
m	0	0	0	0	0	0	1	0

← 26 elements →

One-hot Encoded Characters

Letter	a ... j ... i ...						m	...
j	0	0	1	0	0	0	0	0
i	0	0	0	0	1	0	0	0
m	0	0	0	0	0	0	1	0

← 26 elements →

Gender Prediction Using RNNs

Prepare Data

Download and extract

File manipulation - simple

Create Neural Network

Two linear layers

LogSoftMax output layer

Evaluate Model

Confusion Matrix

Similar to training, but no backprop

Encode Names as Tensors

No padding needed in PyTorch!

Sequence of one-hot encoded characters

Train Model

Backprop with `loss.backward()`

NLL loss to go with LogSoftMax

Use in Prediction

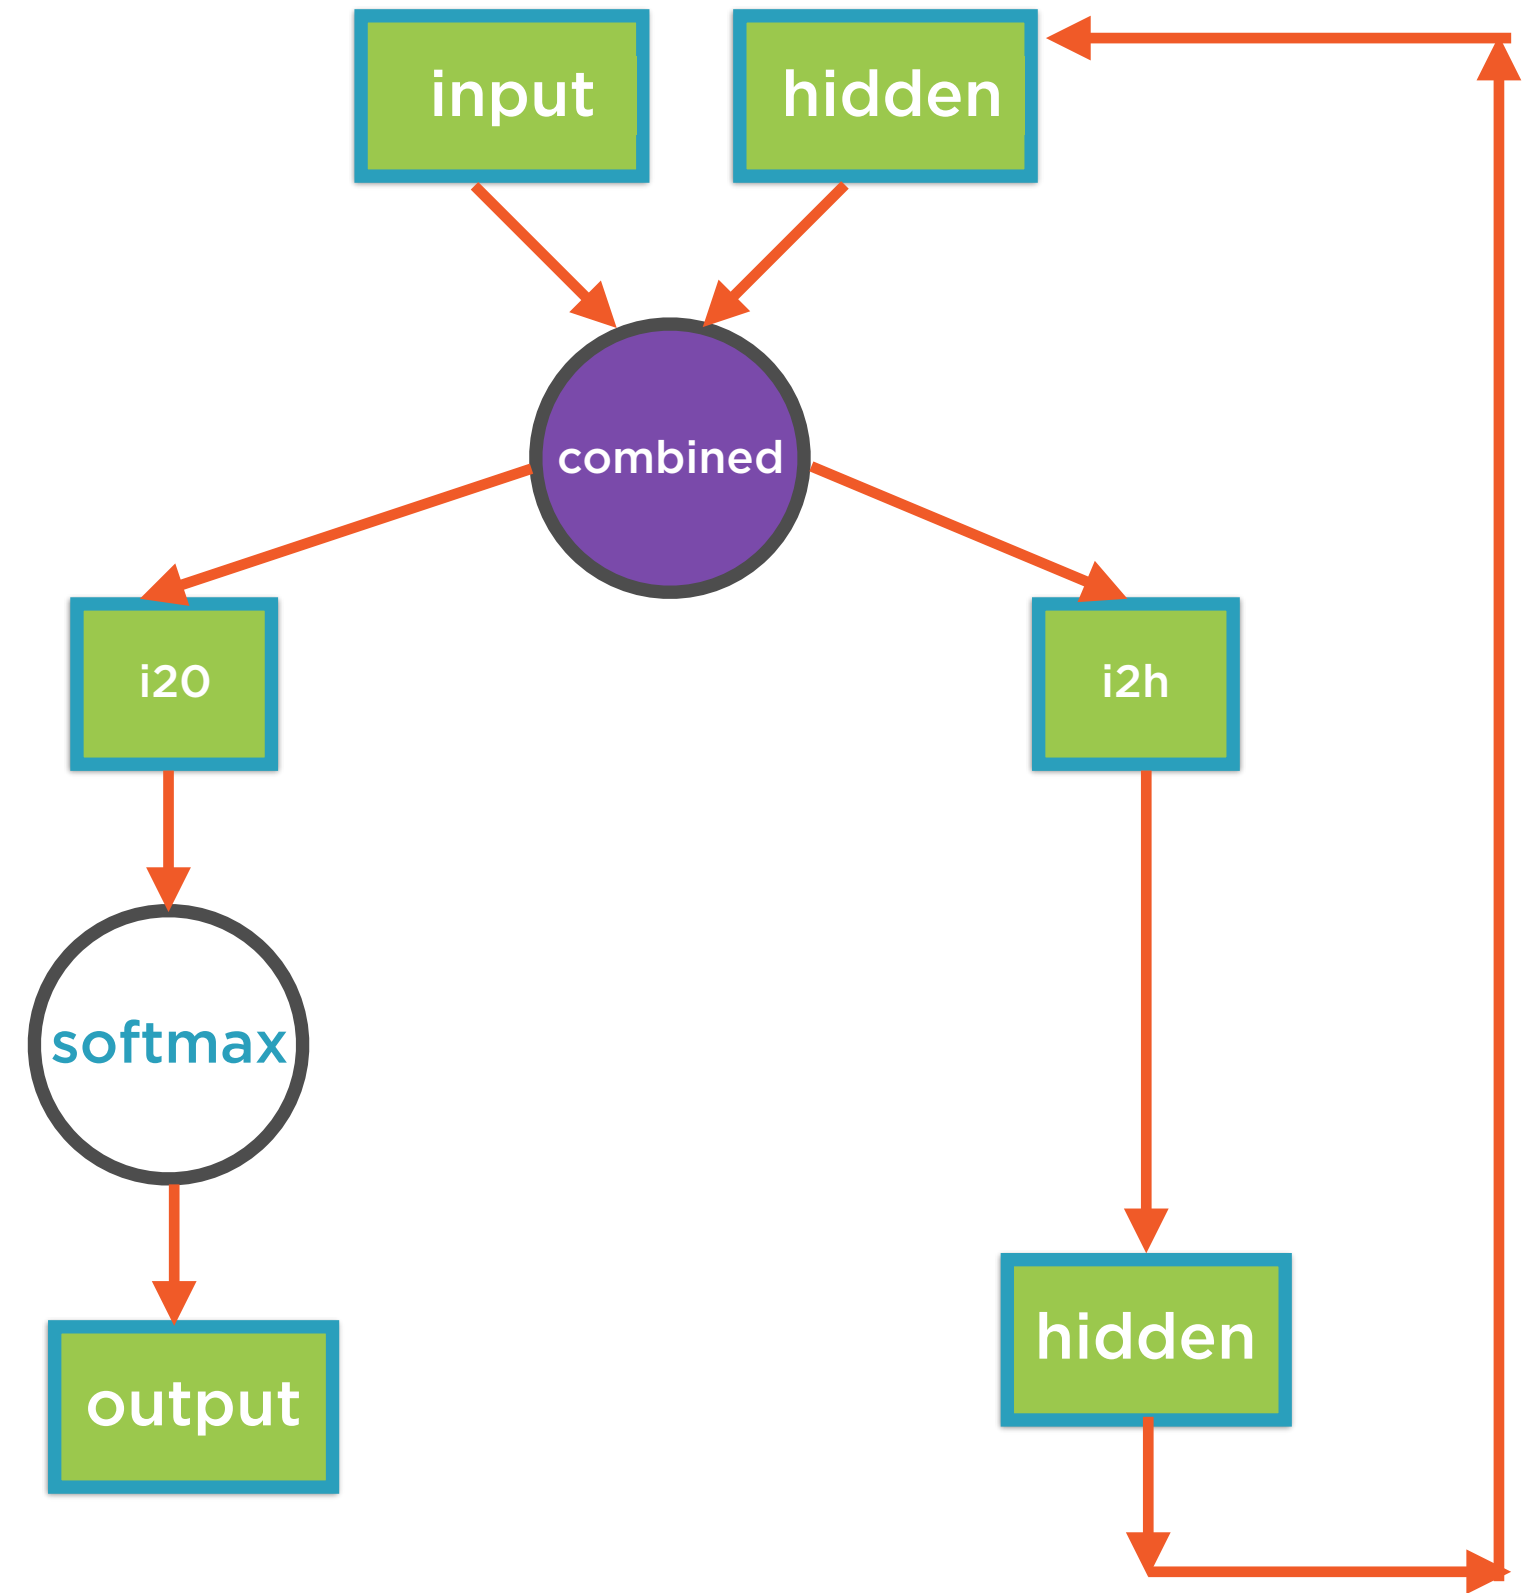
Ready to go!

Could add REST API server

```

class RNN(nn.Module):
# In constructor
    self.i2h = nn.Linear()
    self.i2o = nn.Linear()
    self.softmax = nn.LogSoftmax()
# Forward method
    def forward(self, input, hidden):
        combined = torch.cat((input, hidden), 1)
        hidden = self.i2h(combined)
        output = self.i2o(combined)
        output = self.softmax(output)
        return output, hidden

```



Demo

**Gender prediction/classification from
names using RNNs**

Confusion Matrix in Evaluating Classifiers

Confusion Matrix

Predicted Labels



Cancer

No
Cancer

Actual Label



Cancer

10 instances

4 instances

No
Cancer

5 instances

1000 instances

	Cancer	No Cancer
Cancer	10 instances	4 instances
No Cancer	5 instances	1000 instances

Confusion Matrix

Predicted Labels



Cancer

No
Cancer

Actual Label



Cancer

No
Cancer

	Cancer	No Cancer
Cancer	10 instances	4 instances
No Cancer	5 instances	1000 instances

Confusion Matrix

Predicted Labels

Actual Label

	Cancer	No Cancer
Cancer	10	4
No Cancer	5	1000

True Positive

Predicted Labels

Cancer

No
Cancer


Actual Label

Cancer

No
Cancer

10 TP	4
5	1000

Actual Label = Predicted Label



False Positive

Predicted Labels

Cancer

No
Cancer

Actual Label

Cancer

10

4

No
Cancer

5

FP

1000

Actual Label \neq Predicted Label

	Cancer	No Cancer
Cancer	10	4
No Cancer	5	1000

True Negative

Predicted Labels

Cancer

No
Cancer

Actual Label

Cancer

10

4

No
Cancer

5

1000

TN

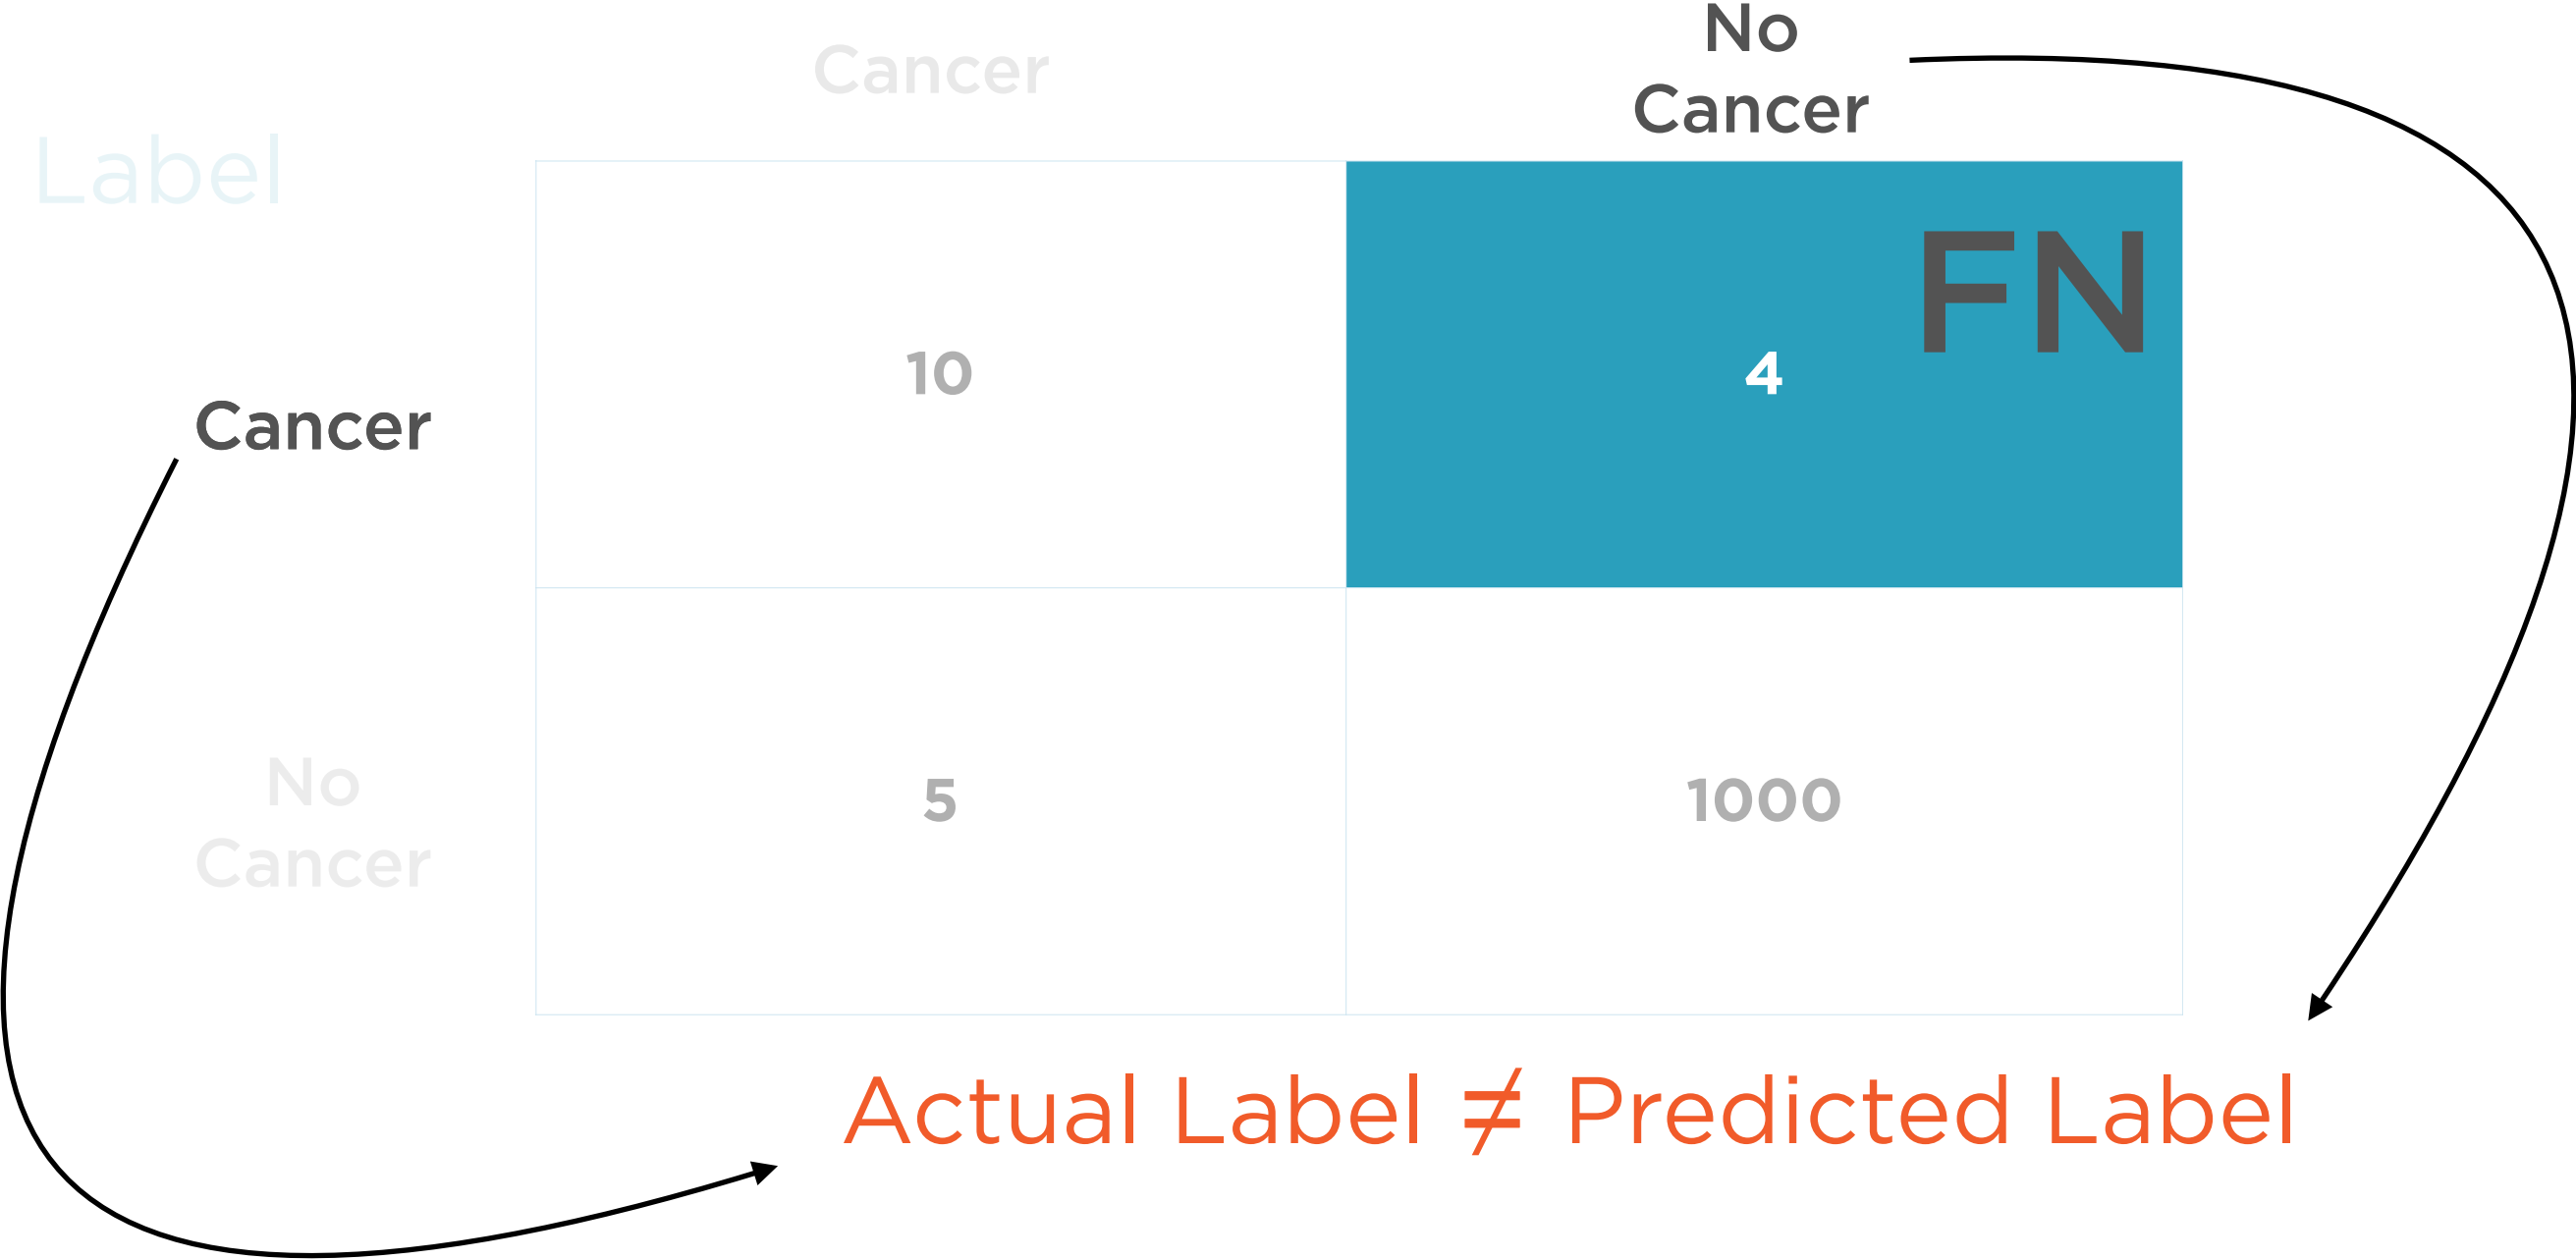
Actual Label = Predicted Label

	Cancer	No Cancer
Cancer	10	4
No Cancer	5	1000 TN

False Negative

Predicted Labels

Actual Label



Accuracy

Predicted Labels

Cancer

No
Cancer

Actual Label

Cancer

No
Cancer

	Cancer	No Cancer
Cancer	TP 10	FN 4
No Cancer	FP 5	TN 1000

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{Num Instances}} = \frac{1010}{1019} = 99.12\%$$

Precision

Predicted Labels

Actual Label

	Cancer	No Cancer
Cancer	10 TP	4 FN
No Cancer	5 FP	1000 TN

Precision = Accuracy when classifier flags cancer

Recall

Predicted Labels

Cancer

No
Cancer

Actual Label

Cancer

10

TP

4

FN

No
Cancer

5

FP

1000

TN

Recall = Accuracy when cancer actually present


Summary

Recurrent Neural Networks (RNNs) are another common NN architecture

Great for working with sequences

Classification using RNNs

PyTorch has built-in support for RNNs



Related Courses

Building Classification Models with TensorFlow

Building Unsupervised Learning Models with TensorFlow

Deep Learning Using TensorFlow and Apache MXNet on AWS SageMaker

Building and Deploying Keras in a Multi-cloud Environment